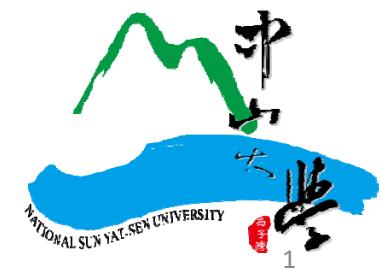


Basic Machine Learning (1)

Linear-regression

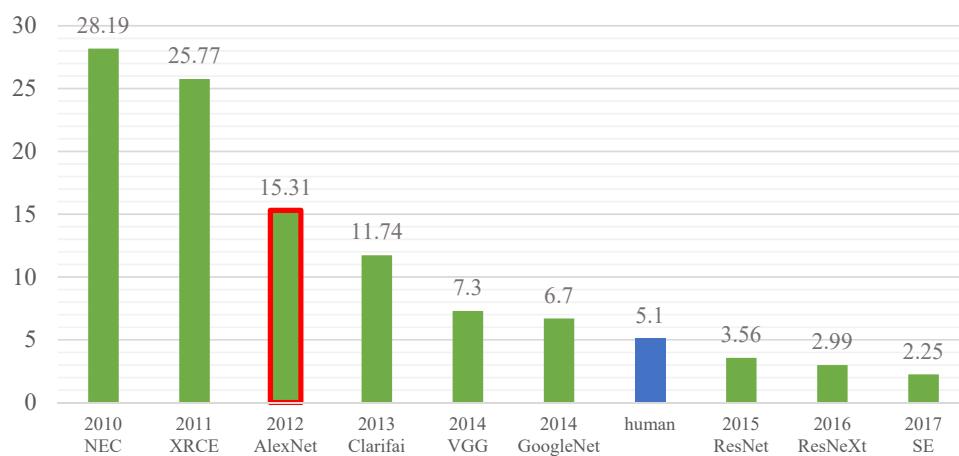
國立中山大學
資訊工程系
張雲南



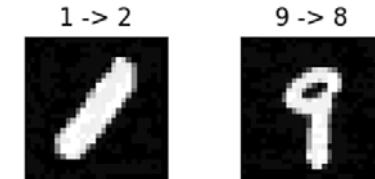
Outlines

- ❖ Introduction to machines learning
- ❖ Linear Regression
- ❖ Data preprocessing
- ❖ Using scikit-learn for machine learning

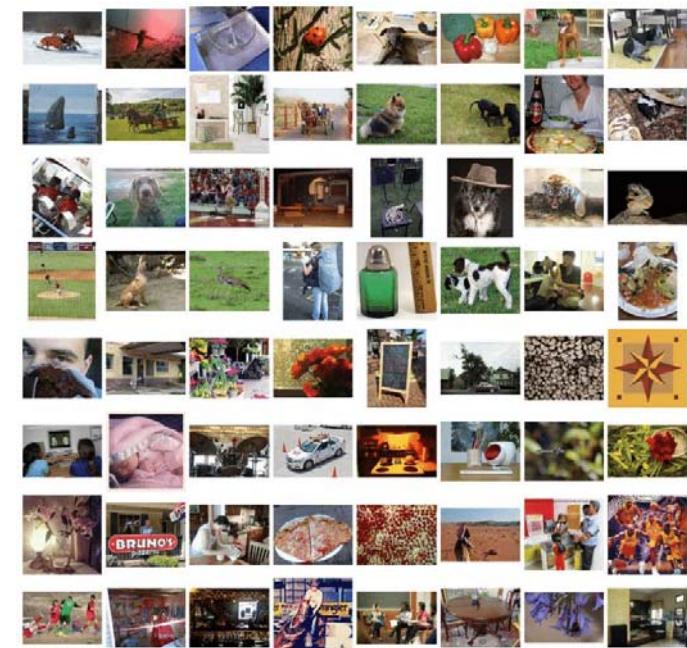
Right perception of current AI tech



ImageNet Large Scale Visual Recognition Competition



14 million images, 20000 categories



Three driving forces for AI

❖ 數據 Data

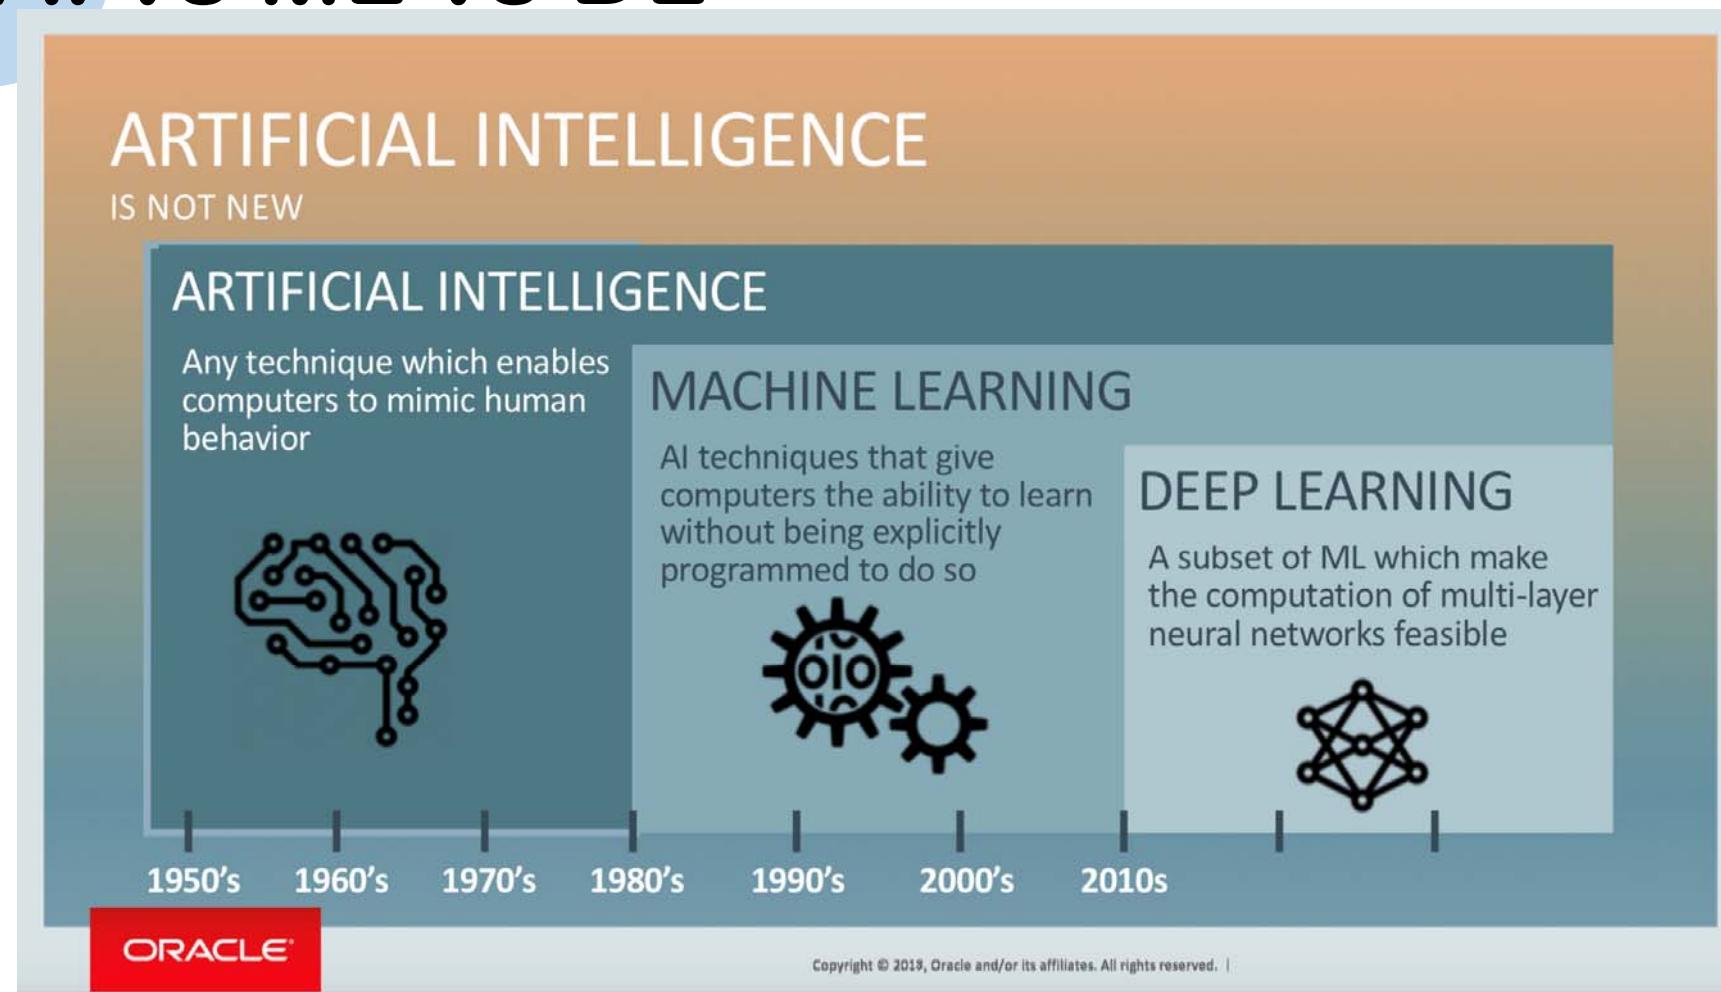
❖ 演算法 Algorithm ★

- ◆ Machine learning model (Neural Network architecture)
- ◆ Training schemes and methods

❖ 計算能力 Computing Power

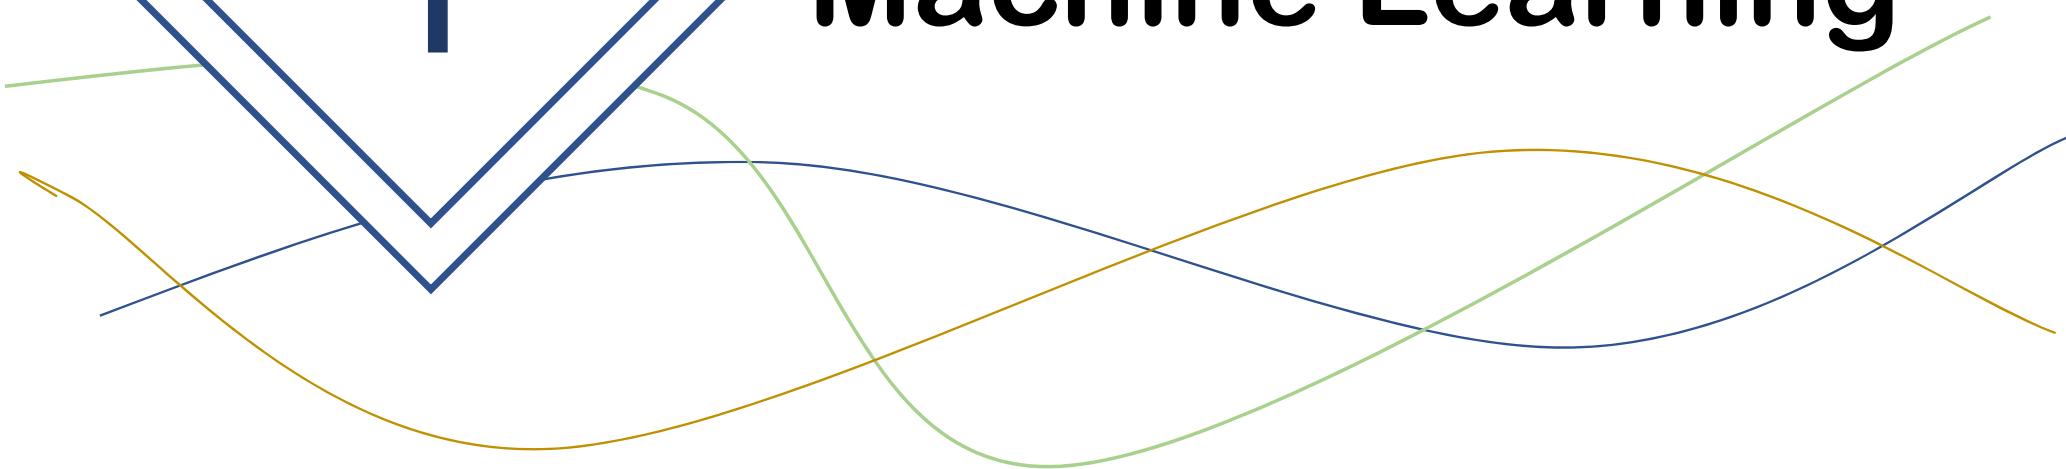
- ◆ Tesla M40 GPU: 14 days, nVidia DGX (8 x P100 GPU): 21hr
 - ❖ Training Data : ImageNet 1K, ~1M image, 1000 class
 - ❖ ResNet-50, 90 epochs, Top-1 Accuracy : 72%

AI vs ML vs DL





Introduction to Machine Learning



Machine Learning

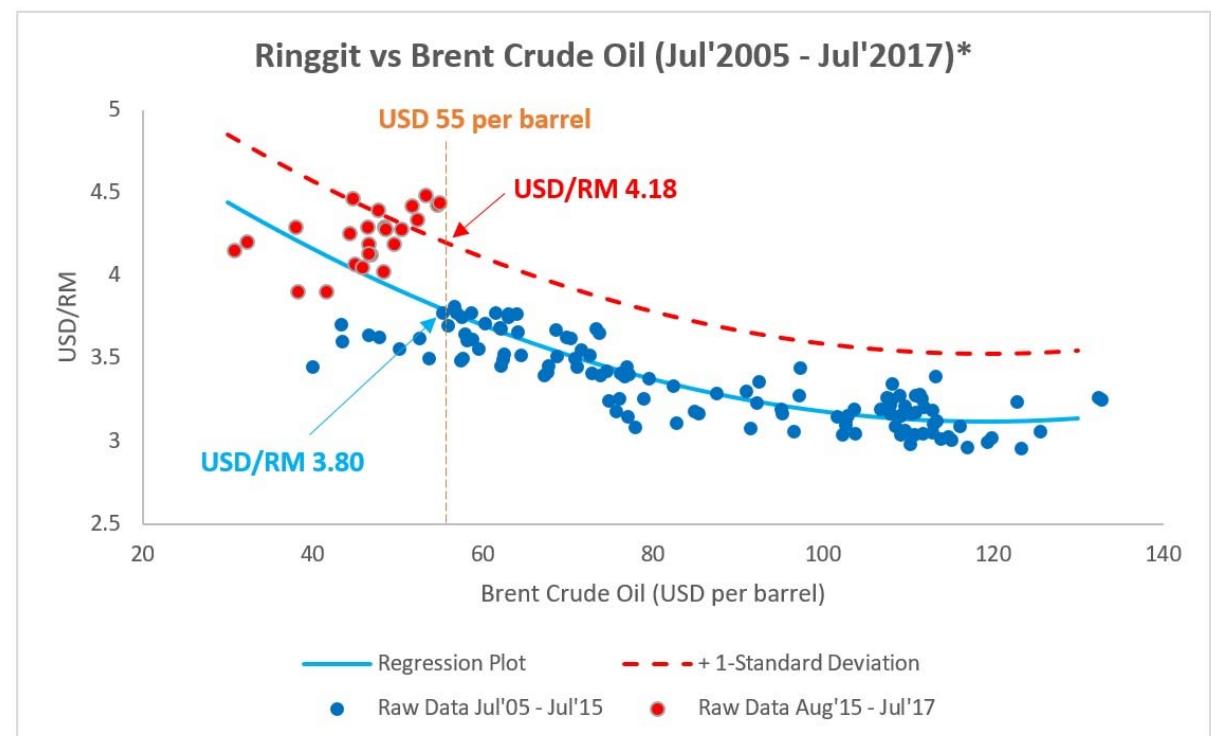
- ◆ **Train** a *model* by *data* to learn the mapping of *features (X)* to the *results (Y)*.

$$f: X \rightarrow Y$$

- ◆ The *model* is used to realize the mapping function, and will be implemented by the software.
- ◆ Depending on the relationship between *X* and *Y*, different types of *models* may be chosen.
 - ◆ f : 身高 → 體重 \Rightarrow can choose *linear regression model*
 - ◆ f : 客戶資料 → 信用核可 \Rightarrow *decision tree model* may be better
- ◆ The dimension of *X* and *Y* may be low or high, but they all have to be encoded by numbers.

Regression

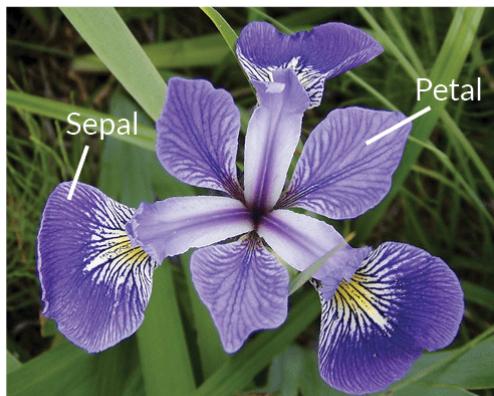
- ❖ $y = f(x_1, x_2, \dots, x_n)$
- ❖ Target y is a float.
- ❖ For example:
 - ❖ Oil price forecast
 - ❖ House price estimation
- ◆ Use *data* to learn *hypothesis functions* based on different machine-learning *models*.



* The data range was chosen from July 2005 onwards because the Ringgit was unpegged after July 2005.

Classification

- ❖ $y = f(x_1, x_2, \dots, x_n)$
- ❖ Target y is a binary vector or discrete number, which corresponds to some categories.
- ❖ For example:
 - ❖ Iris flowers classification
 - ❖ Yes/NO, Live/Dead, Positive/negative,.....



Iris Versicolor



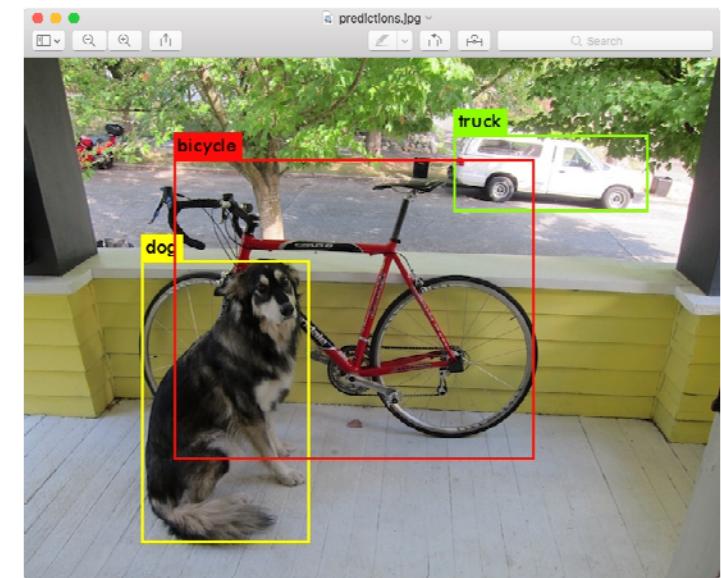
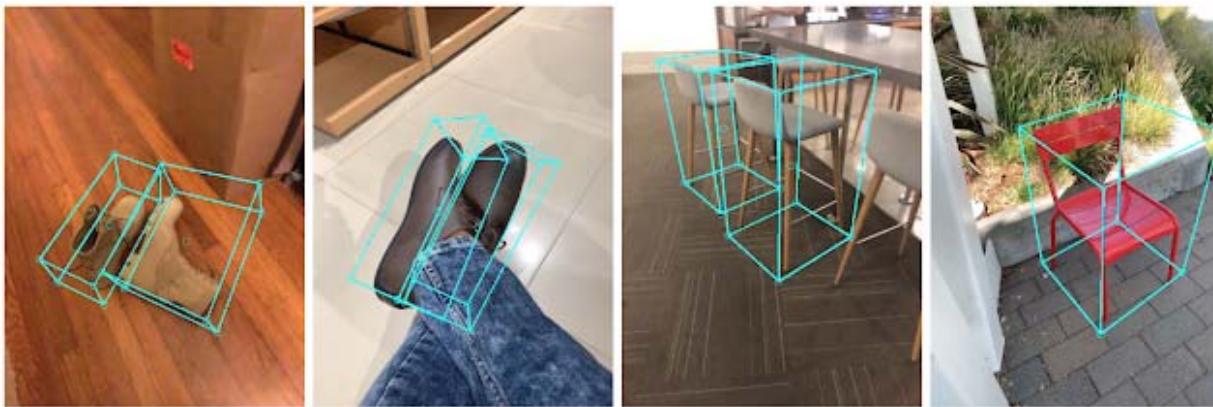
Iris Setosa



Iris Virginica

Pattern recognition

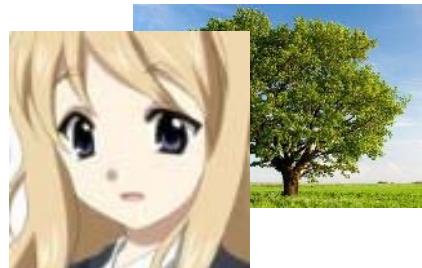
- ❖ $y = f(x_1, x_2, \dots, x_n)$
- ❖ Target y is a list of triplets (*class, w, h*)
- ❖ For example:
 - ❖ Object Detection



Data generation

- ◊ $y = f(x_1, x_2, \dots, x_n)$
- ◊ Target y is high-dimensional data.

(1.2,0.8)



Hello



I'm fine

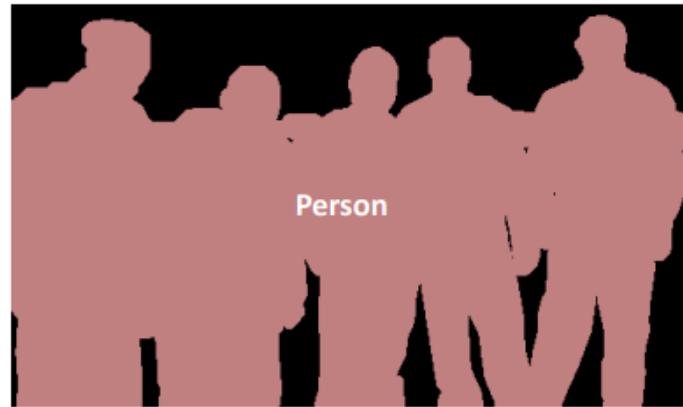


this bird is red with
white and has a very
short beak

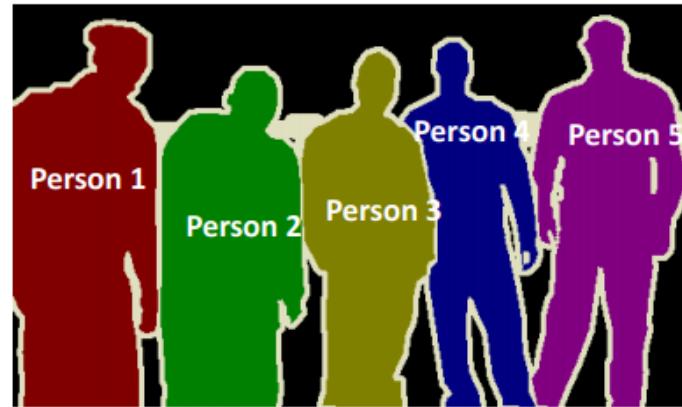


Segmentation

- ◊ $y = f(x_1, x_2, \dots, x_n)$
- ◊ Target y is a high-dimensional vector consisting of discrete numbers.



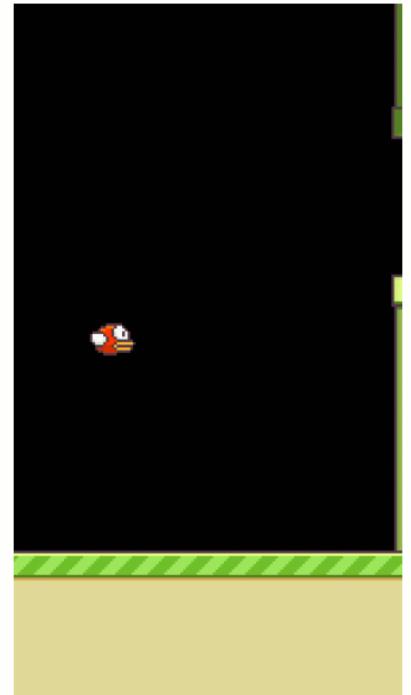
Semantic Segmentation



Instance Segmentation

Action choice

- ❖ $y = f(x_1, x_2, \dots, x_n)$
- ❖ Target y corresponds to the actions to be taken.
- ❖ For example:
 - ◆ AlphaGo
 - ◆ Play video game



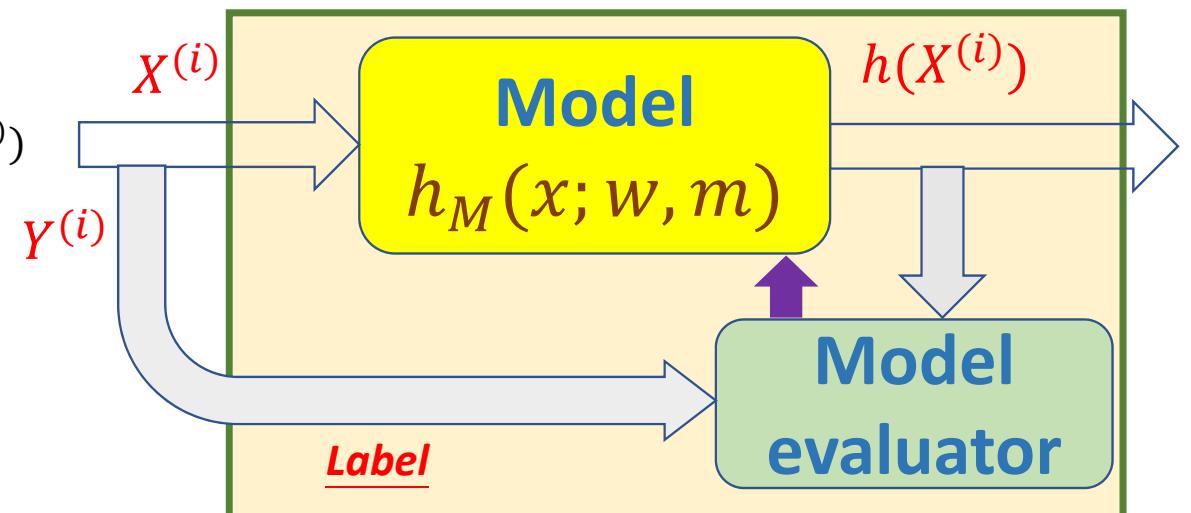
Supervised learning

◊ Training stage

dataset

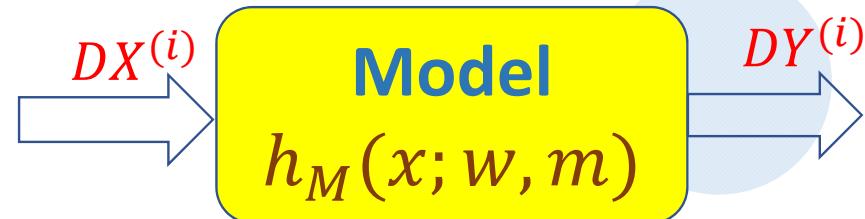
$$f: X \rightarrow Y$$

$(X^{(N)}, Y^{(N)}) \dots (X^{(2)}, Y^{(2)}) \quad (X^{(1)}, Y^{(1)})$



◊ Inference stage

* Depend on whether the object function includes “label”.



Supervised learning

❖ Training stage



X_2
 Y_2 : Monkey



X_1
 Y_1 : Monkey



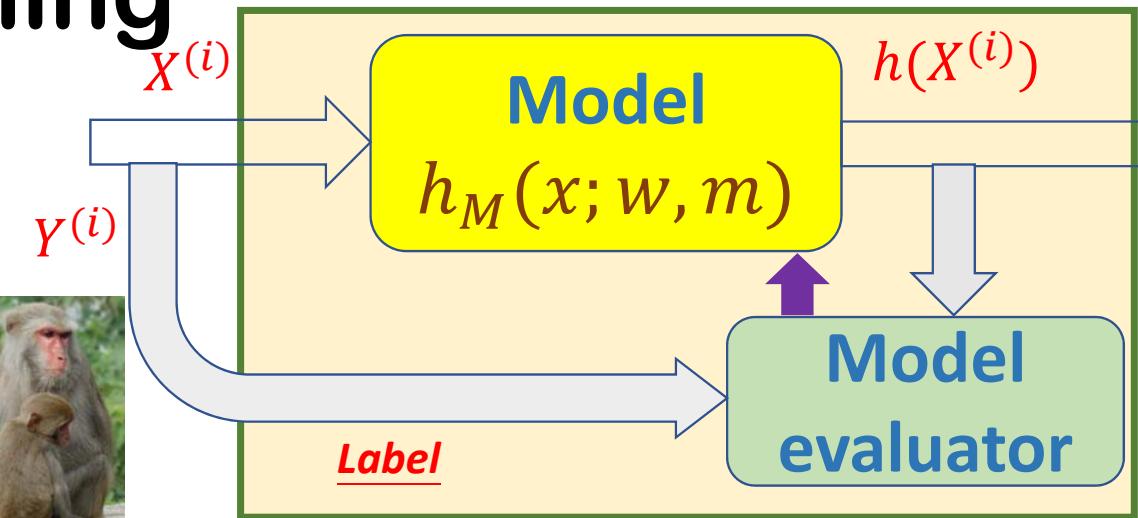
X_0
 Y_0 : Monkey



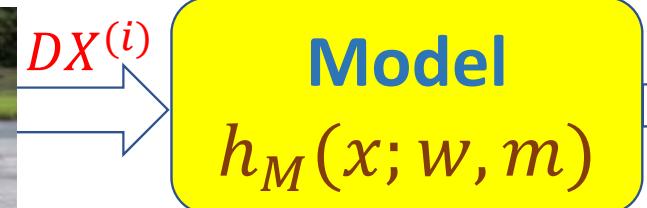
X_4
 Y_4 : Dog



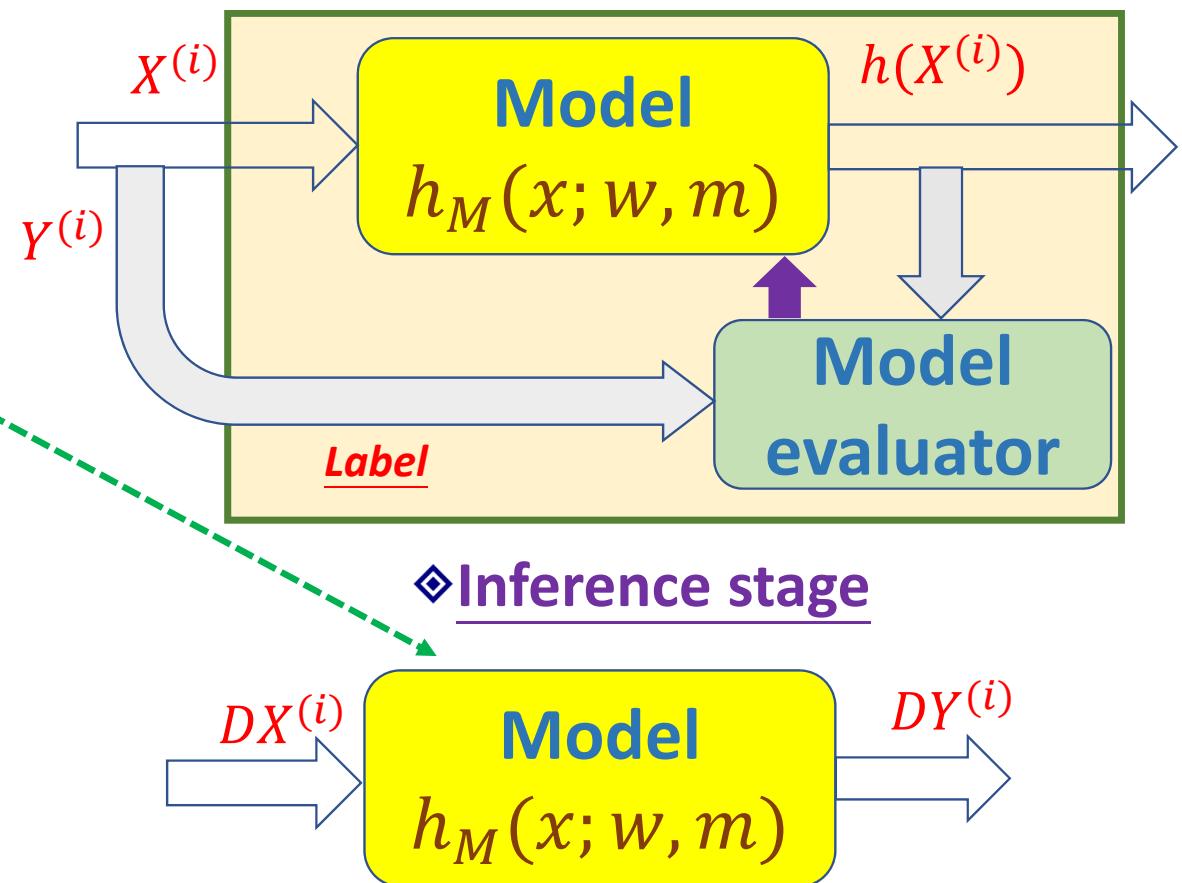
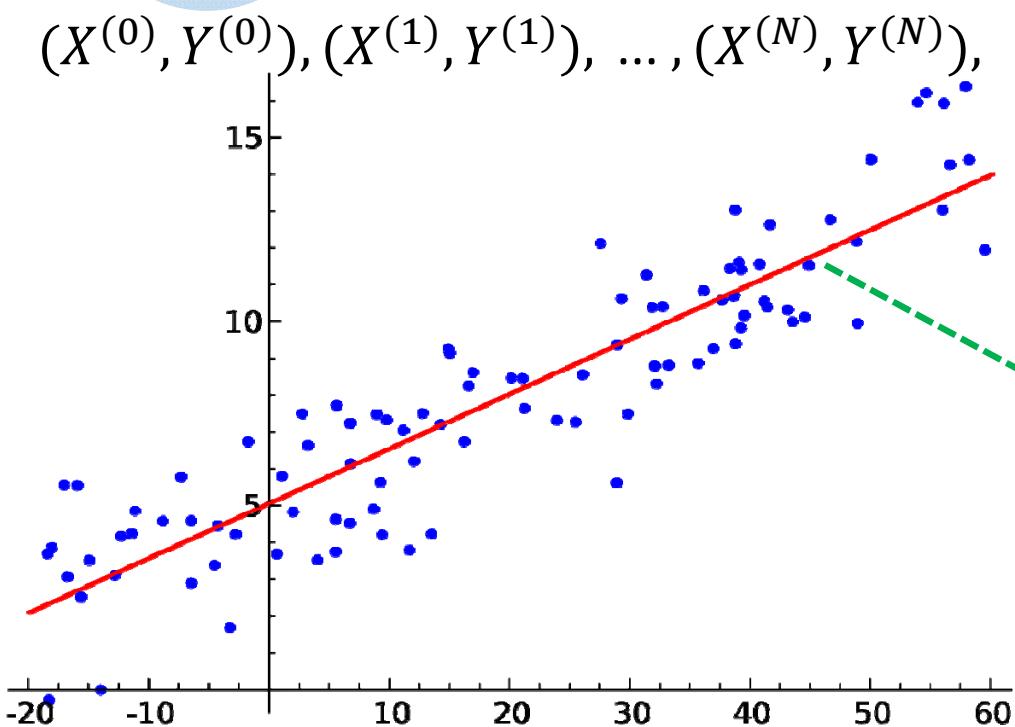
X_3
 Y_3 : Dog



❖ Inference stage

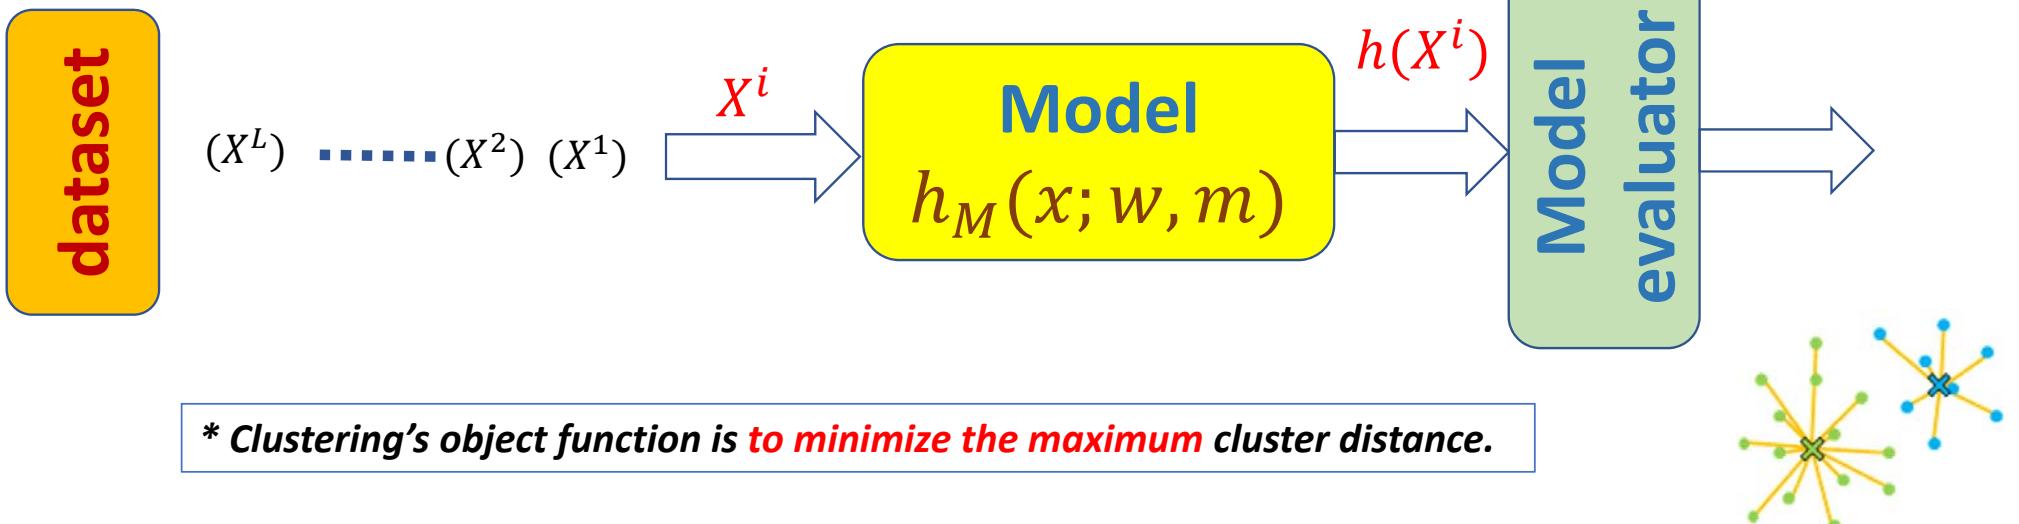


Supervised learning for regression

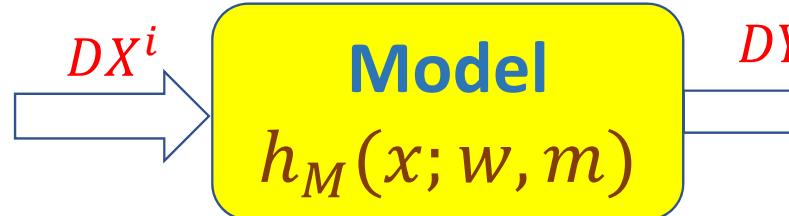


Unsupervised learning

◊ Training stage

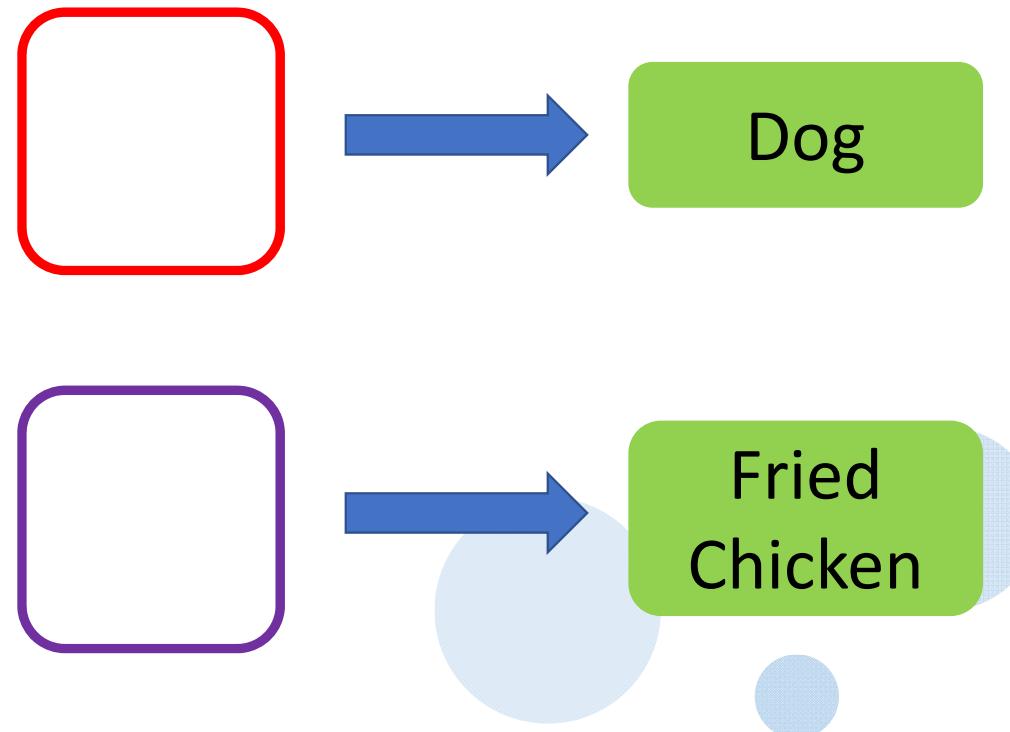
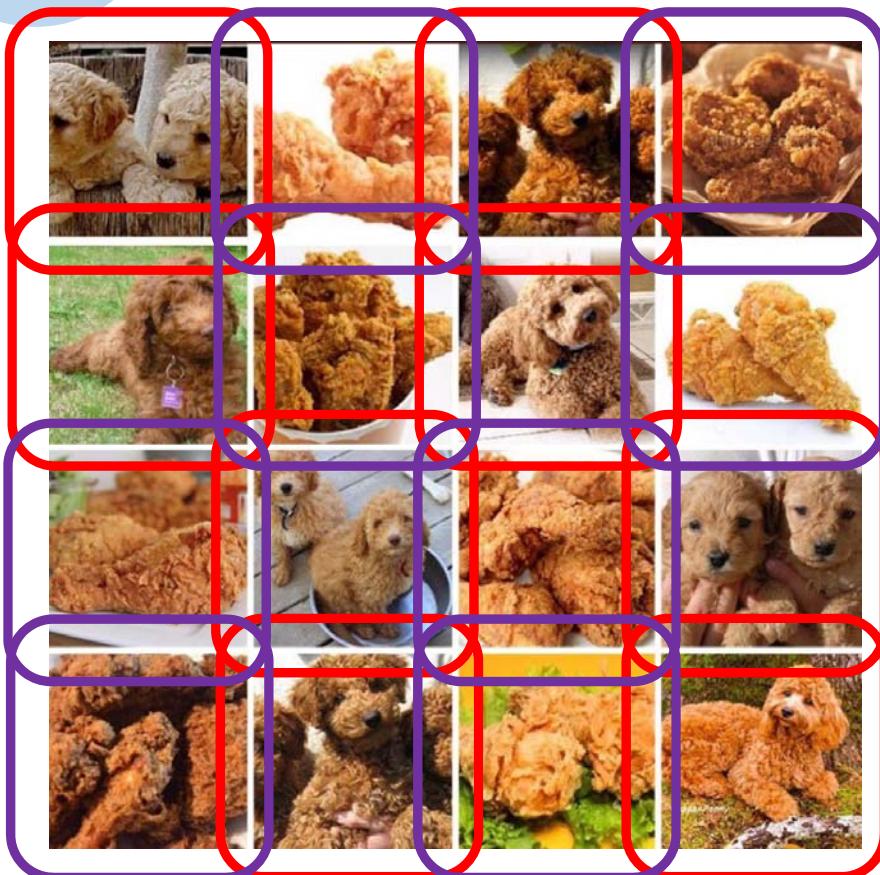


◊ Inference stage



Unsupervised learning

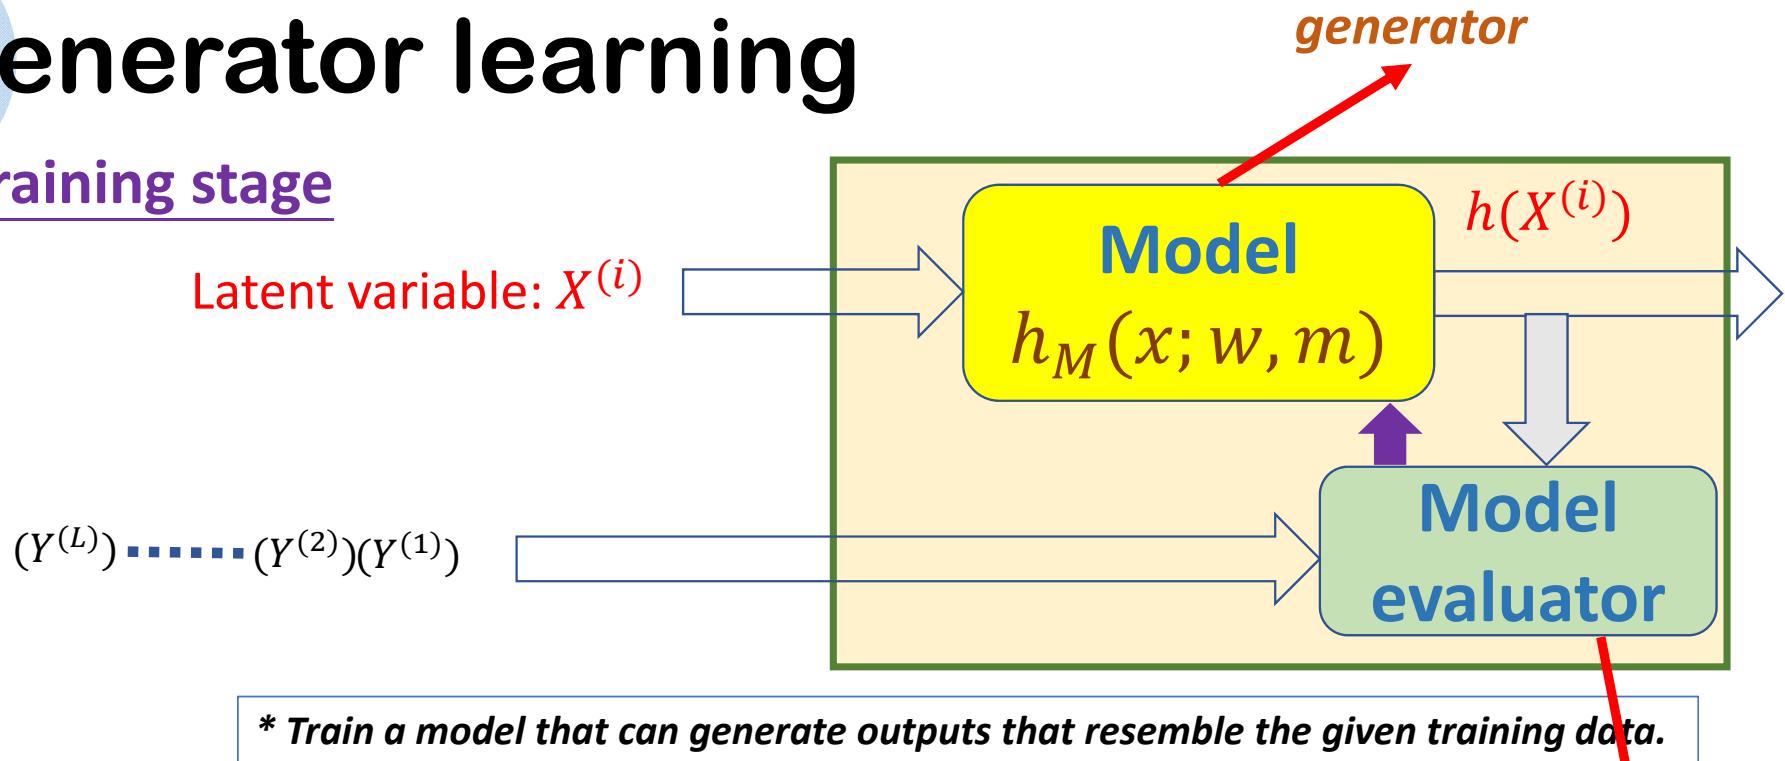
Clustering



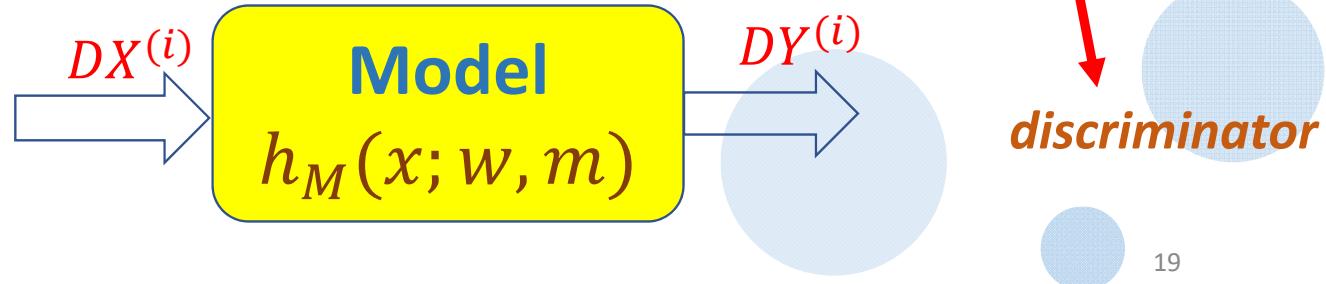
Generator learning

◊ Training stage

dataset



◊ Inference stage

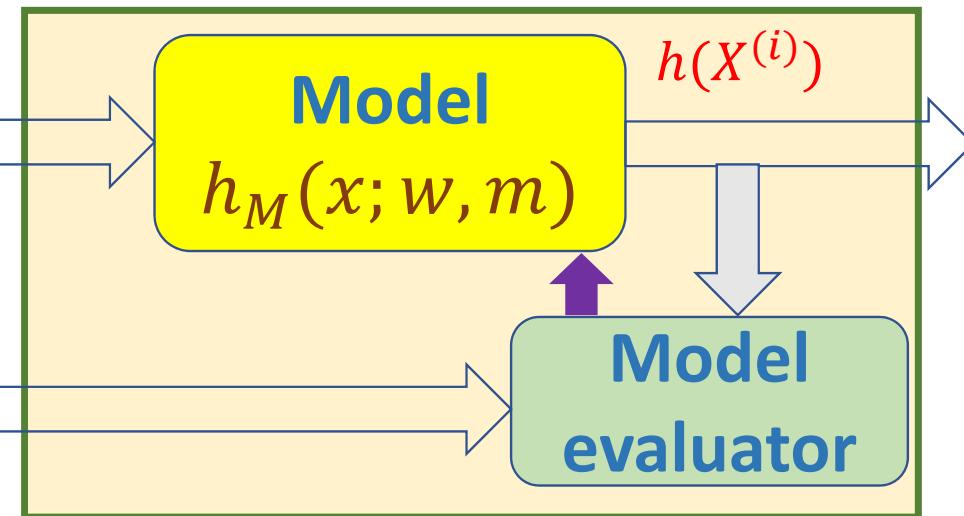


Generator learning

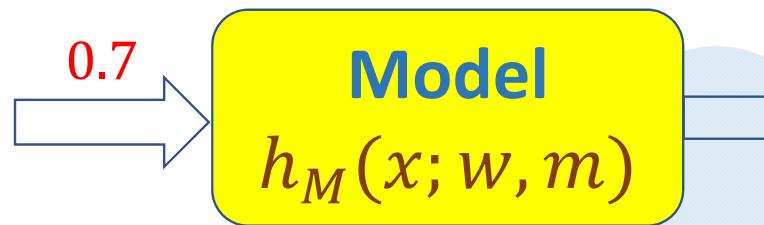
◊ Training stage



Latent variable: $X^{(i)}$



◊ Inference stage

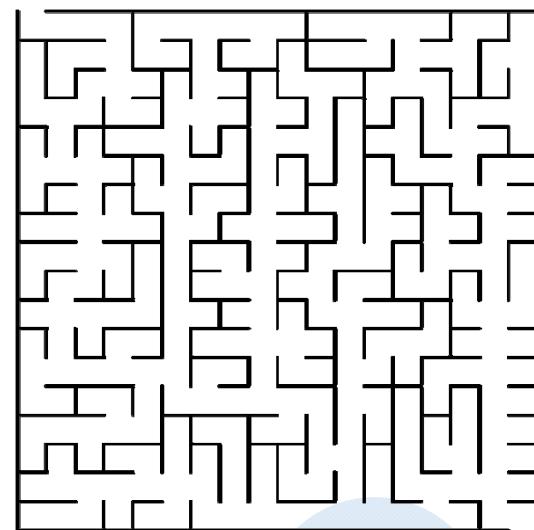
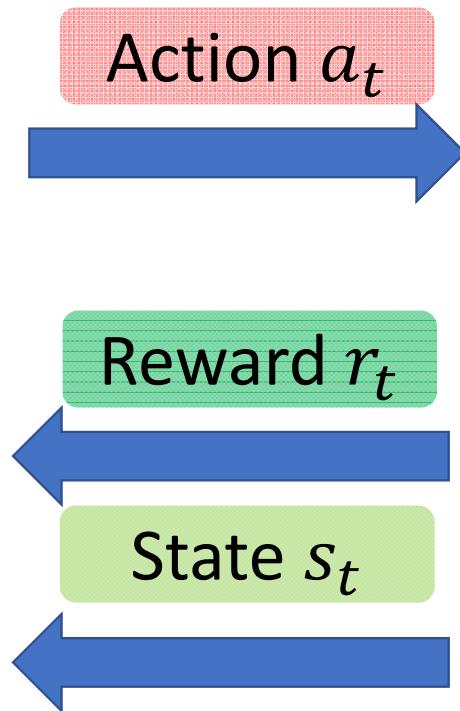


Reinforcement

- ◆ Train an agent function to decide actions.



Agent



Environment

Reinforcement Learning

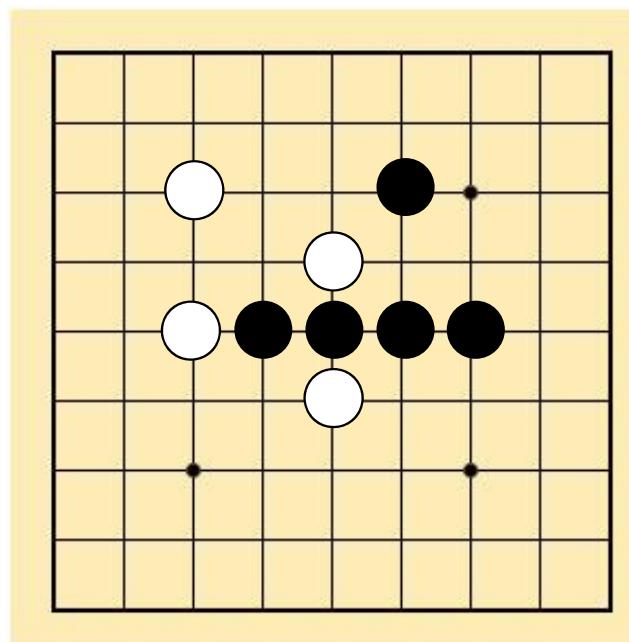
- ❖ In Gomoku, one who can form an unbroken chain of five is the winner

State s_t

Board

Environment

Opponent



Action a_t

Where to place the stone

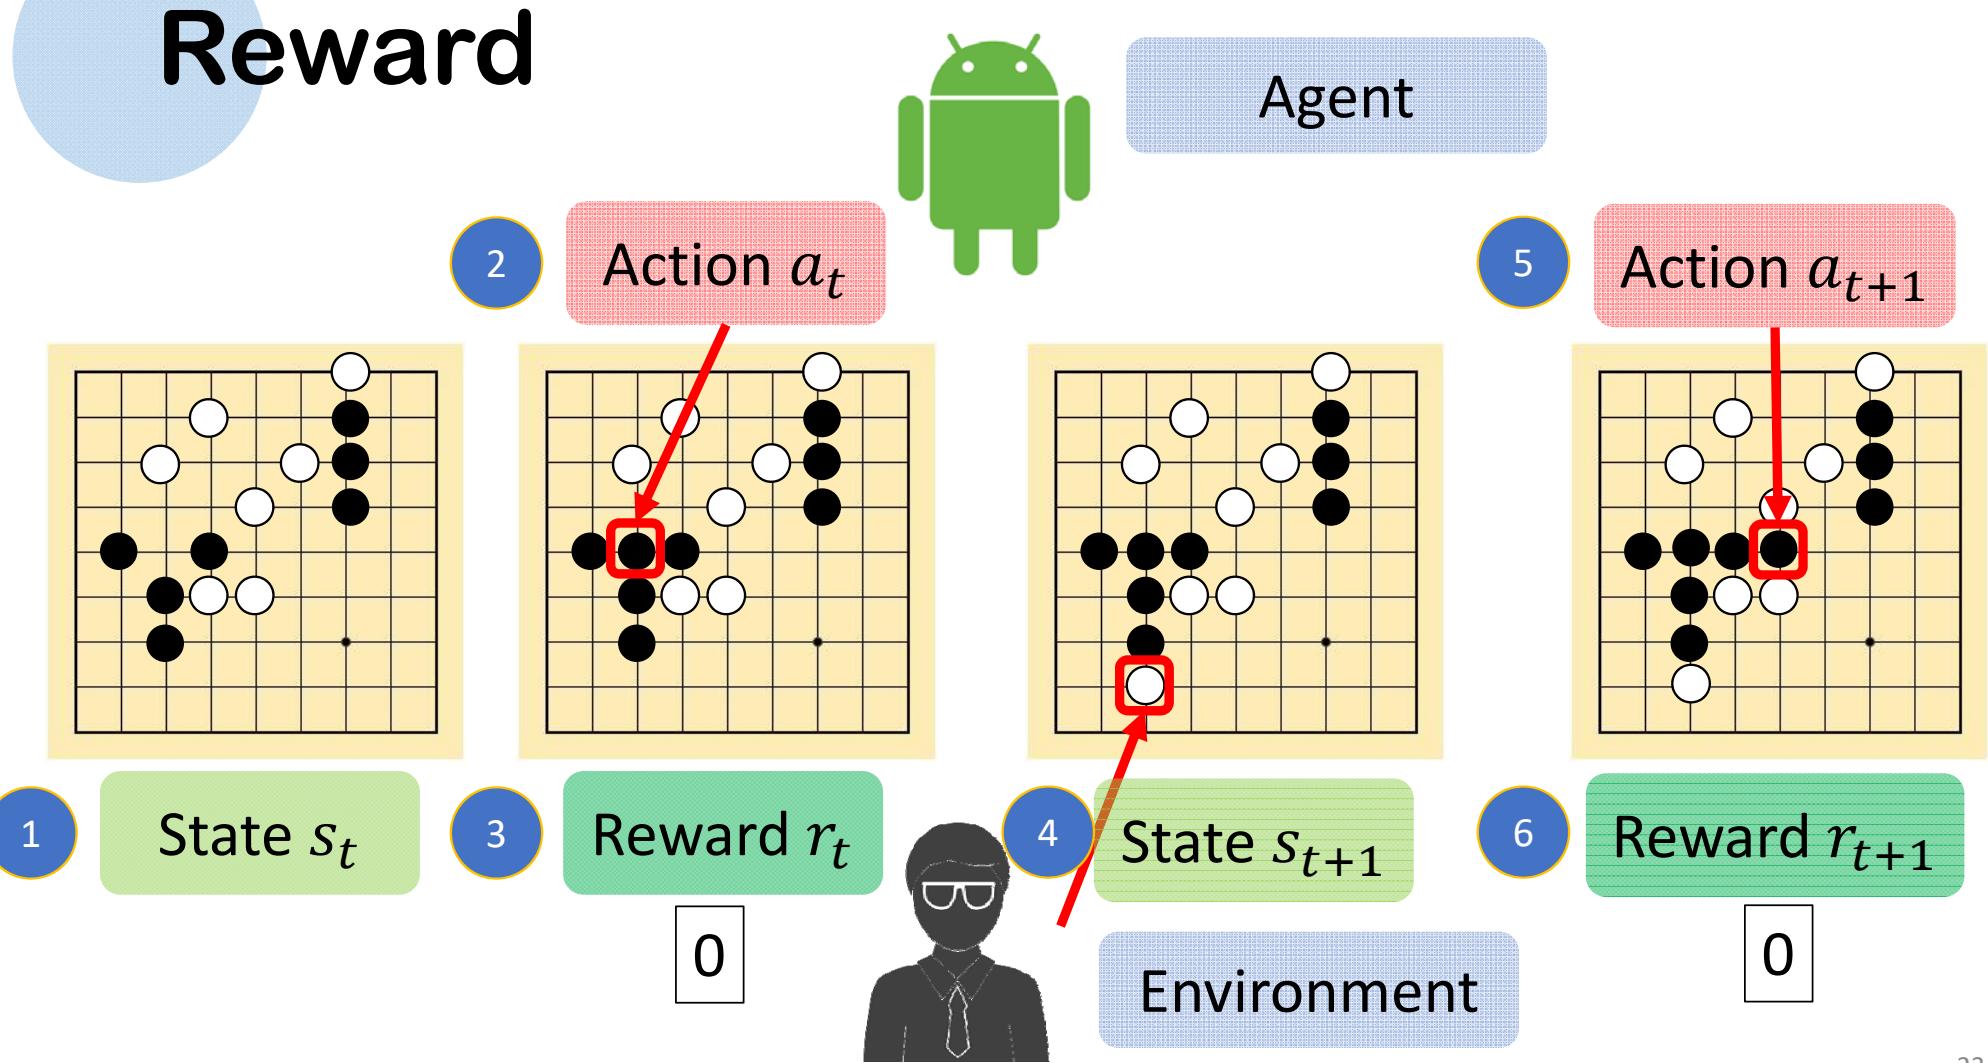
Reward r_t

Win : get 1 point

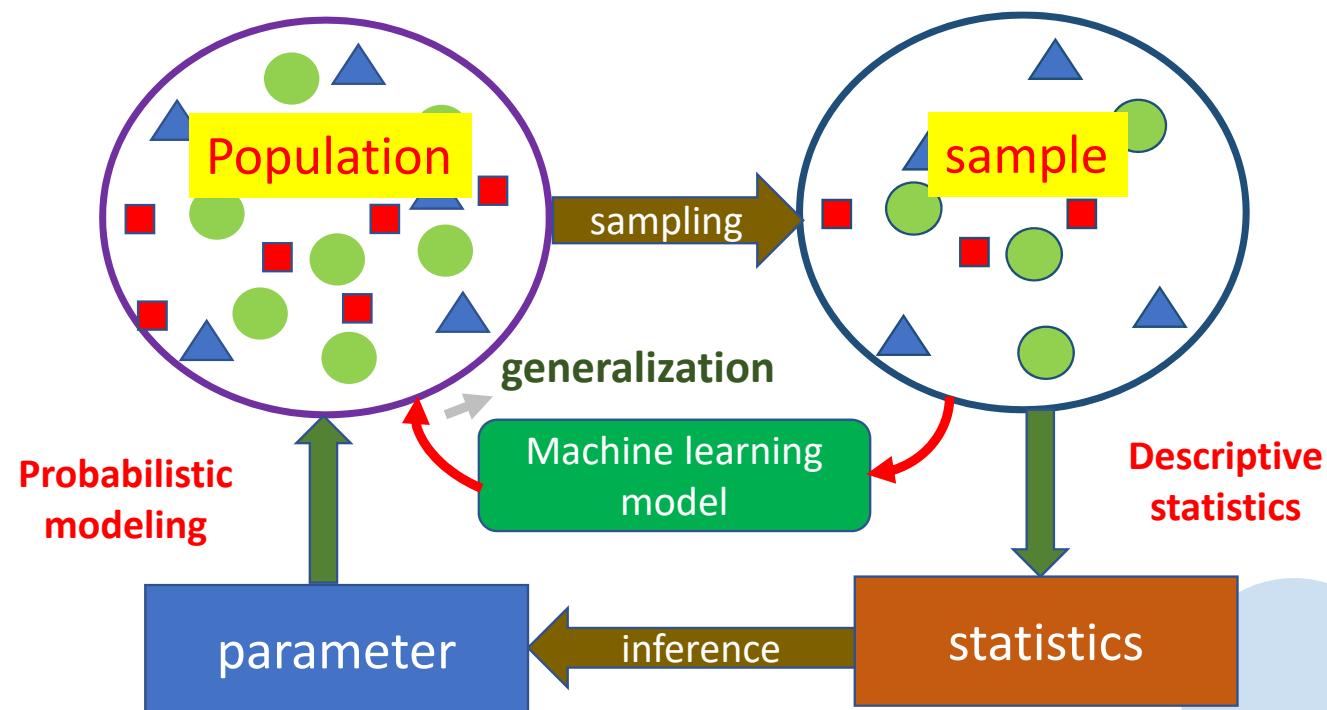
Not-Win : get 0

Loss: get -1 point

Reward

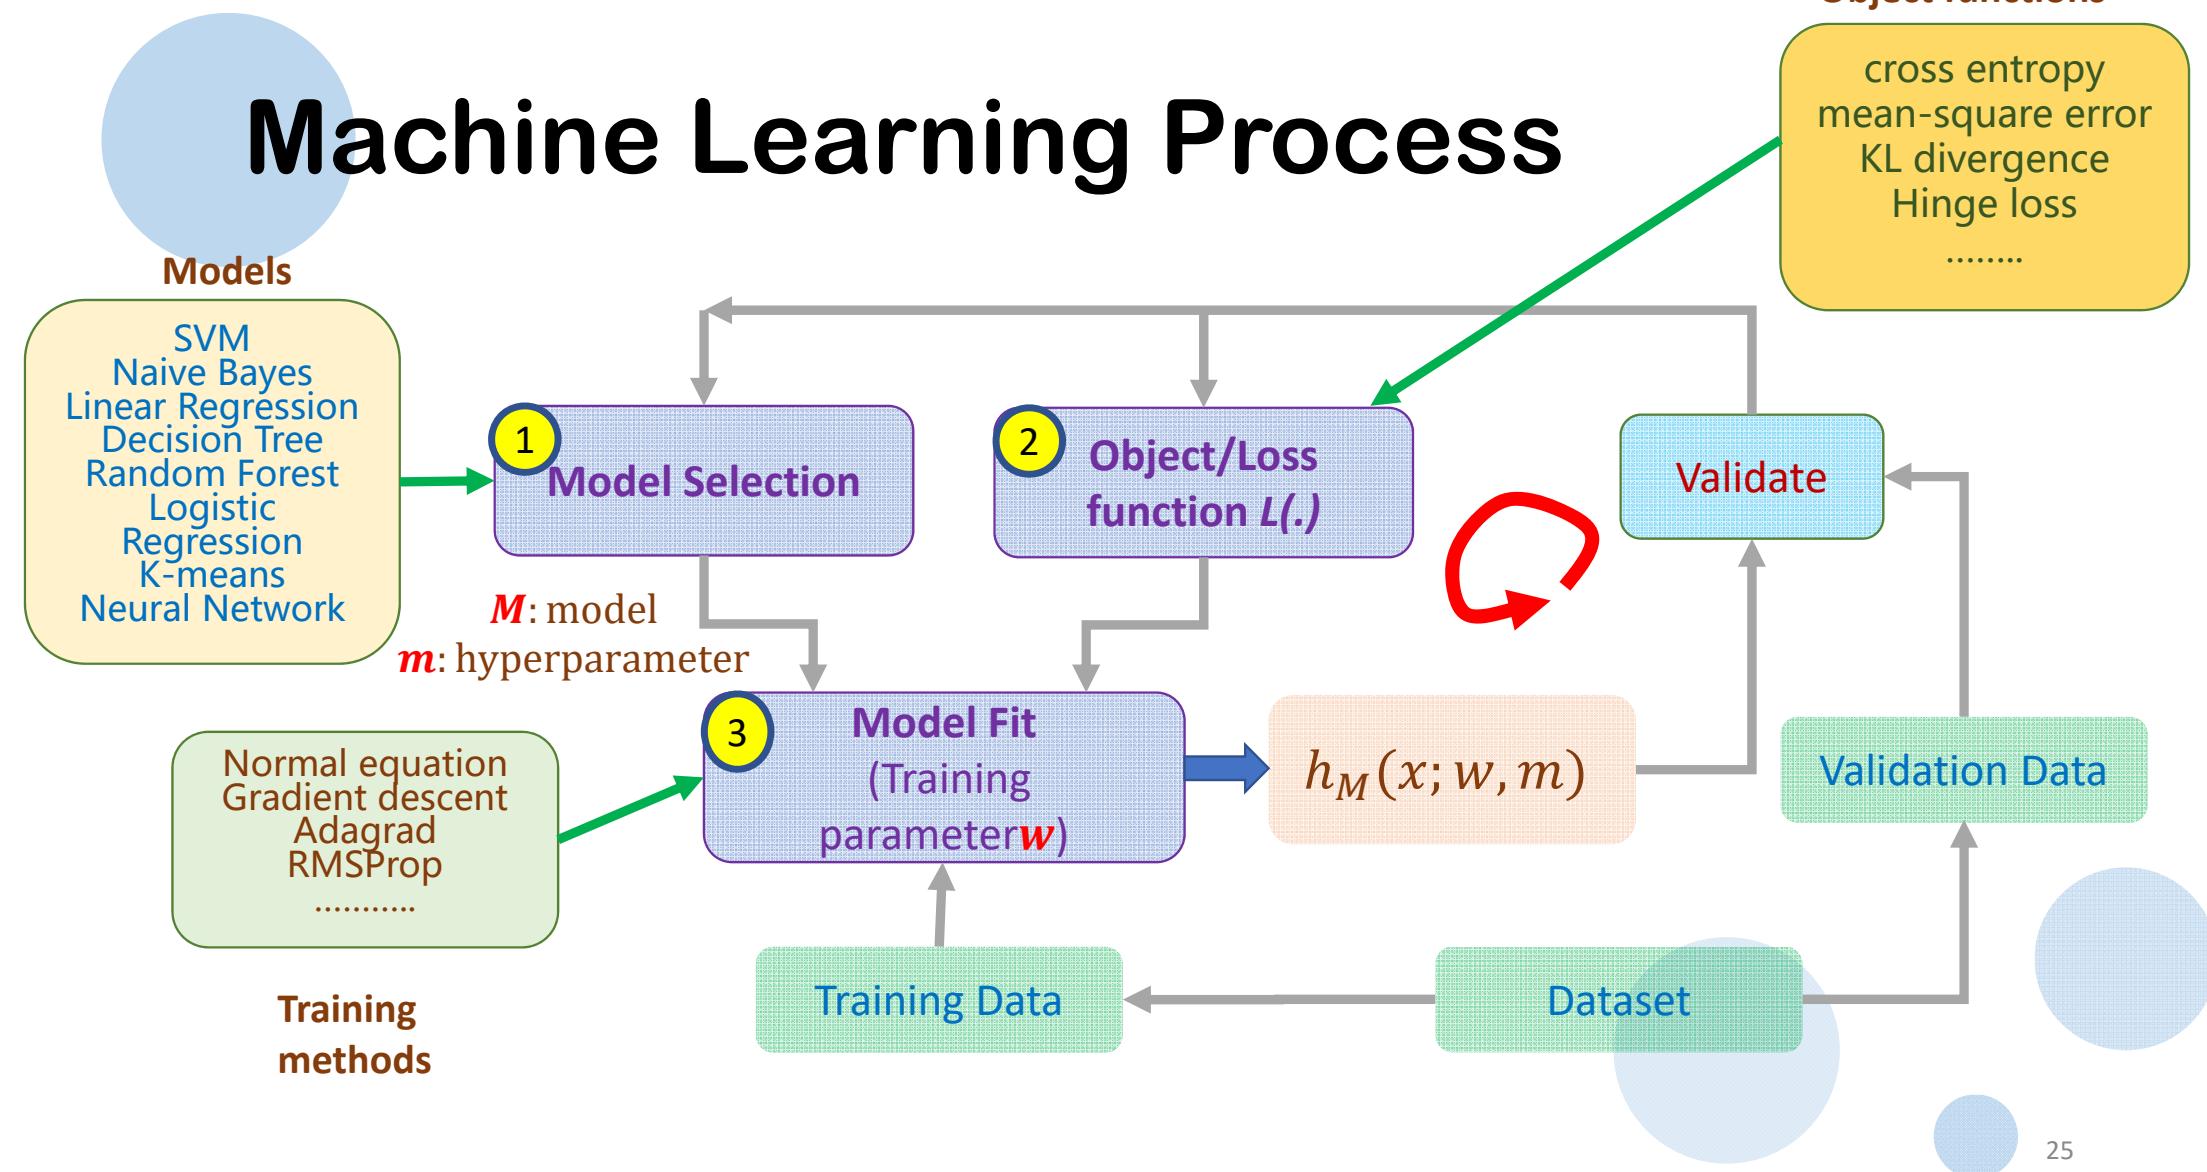


Sampling & Learning

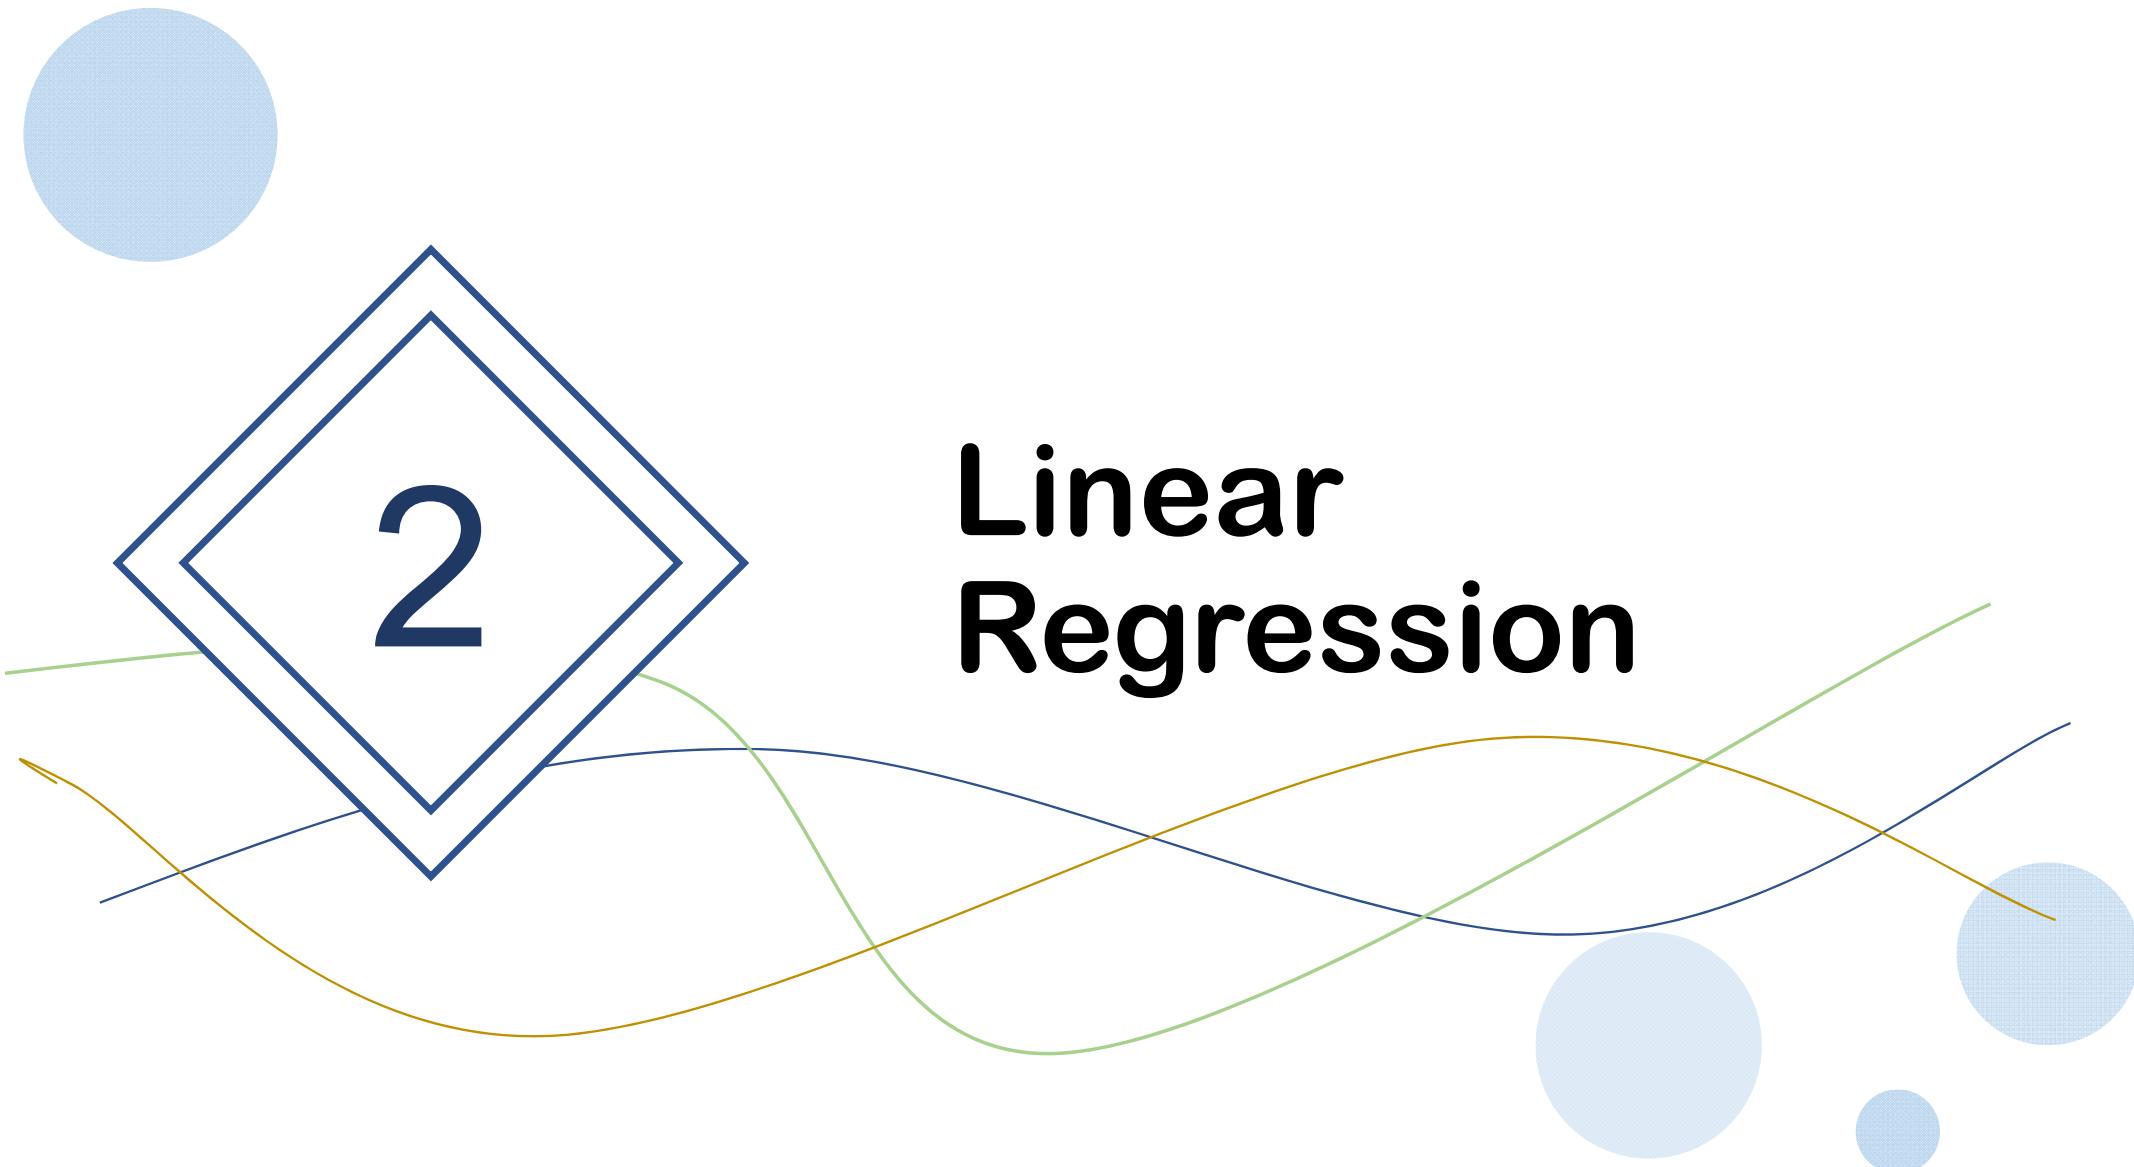


Overfitting vs Underfitting

Machine Learning Process



Linear Regression



What's linear regression

- ❖ A *linear approach* to modelling the relationship between a scalar response (or dependent variable) and one or more explanatory variables (or independent variables).

- ◆ *Summation of weighted results*

$$y(x_1, x_2 \dots x_n) = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$$

Model parameter, unknown in advance.

- ❖ Some related regression terminology

- ◆ Multiple linear regression

- ❖ More than one independent variable

- ◆ Multivariate linear regression

- ❖ More than one dependent variable.

- ◆ Polynomial regression

- ❖ The independent variable has more than one degree in the model of linear equation

Normal equation

- ❖ A linear system of $n+1$ variables:

$$y(x_1, x_2 \dots x_n) = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n$$

- ❖ Suppose our dataset contains m data

$$y^{(1)} = w_0 + w_1 x_1^{(1)} + w_2 x_2^{(1)} + \dots + w_n x_n^{(1)}$$

$$y^{(2)} = w_0 + w_1 x_1^{(2)} + w_2 x_2^{(2)} + \dots + w_n x_n^{(2)}$$

.....

$$y^{(m)} = w_0 + w_1 x_1^{(m)} + w_2 x_2^{(m)} + \dots + w_n x_n^{(m)}$$

- ❖ Can be rewritten as $Xw = b$

- ❖ w can be solved using *normal equation*:

$$X^T X w = X^T b \Rightarrow w = (X^T X)^{-1} X^T b$$

Polynomial Curve Fitting (regression)

❖ Model :

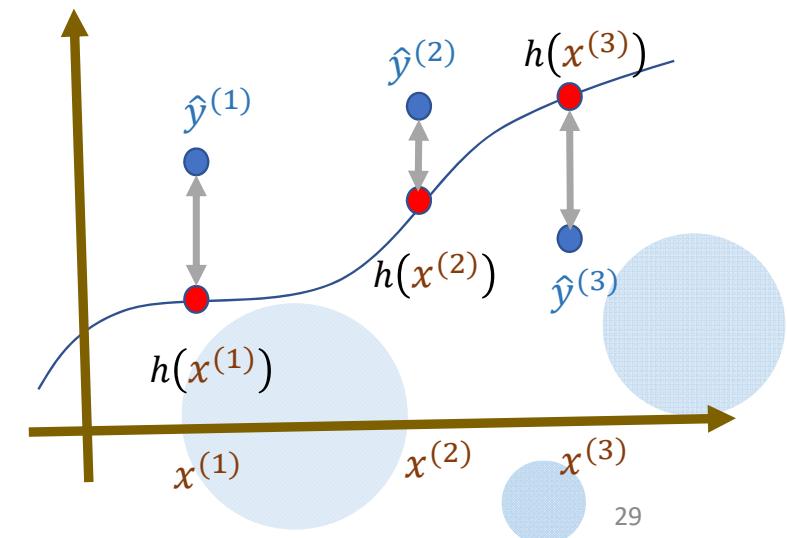
$$h(x; w, m) = w_0 + w_1 x + w_2 x^2 + \cdots + w_m x^m$$

Model parameter, unknown in advance.

Model hyper-parameter, set in advance.

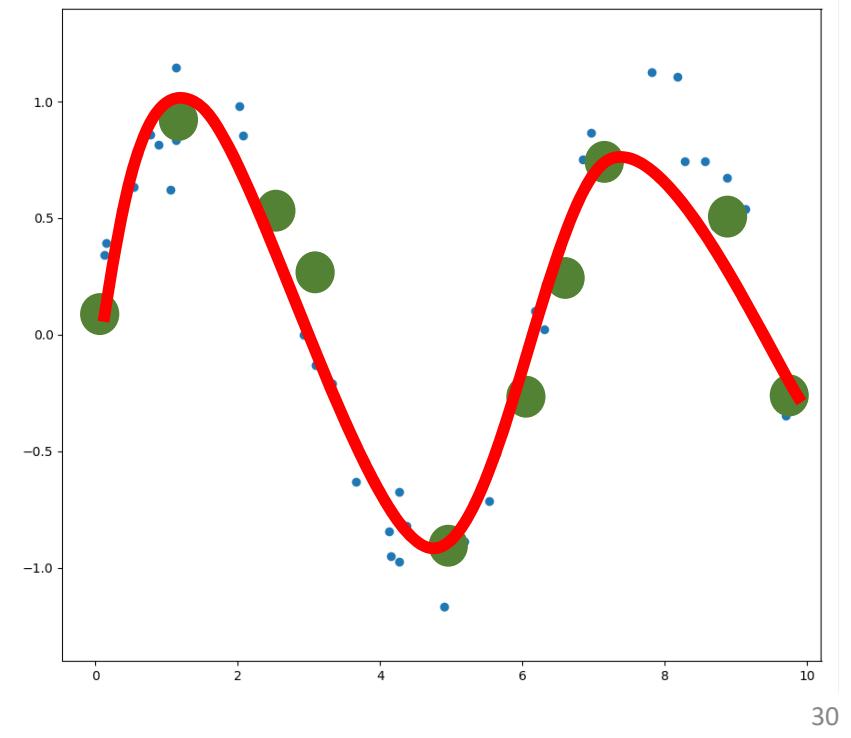
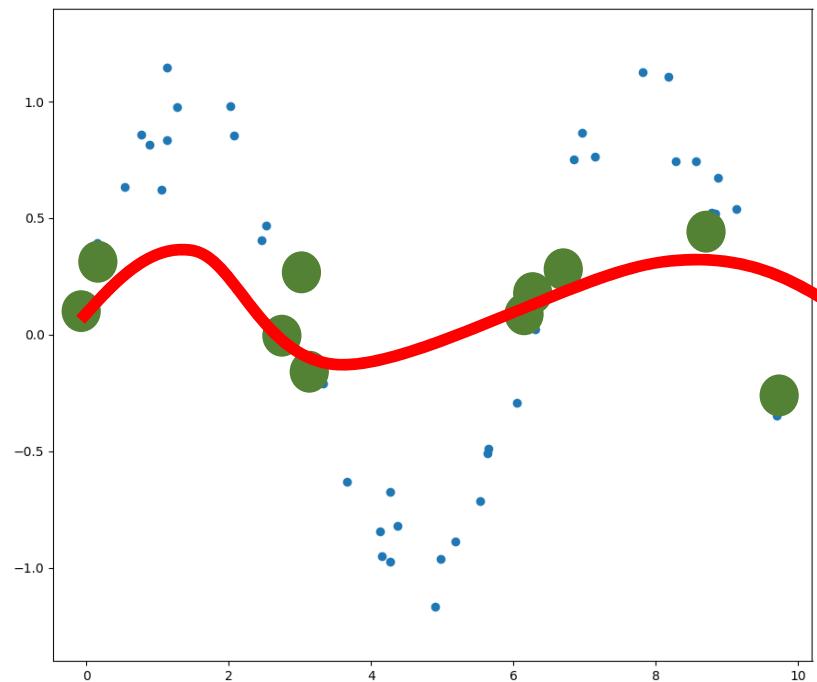
❖ Loss function: Mean-square-error (MSE)

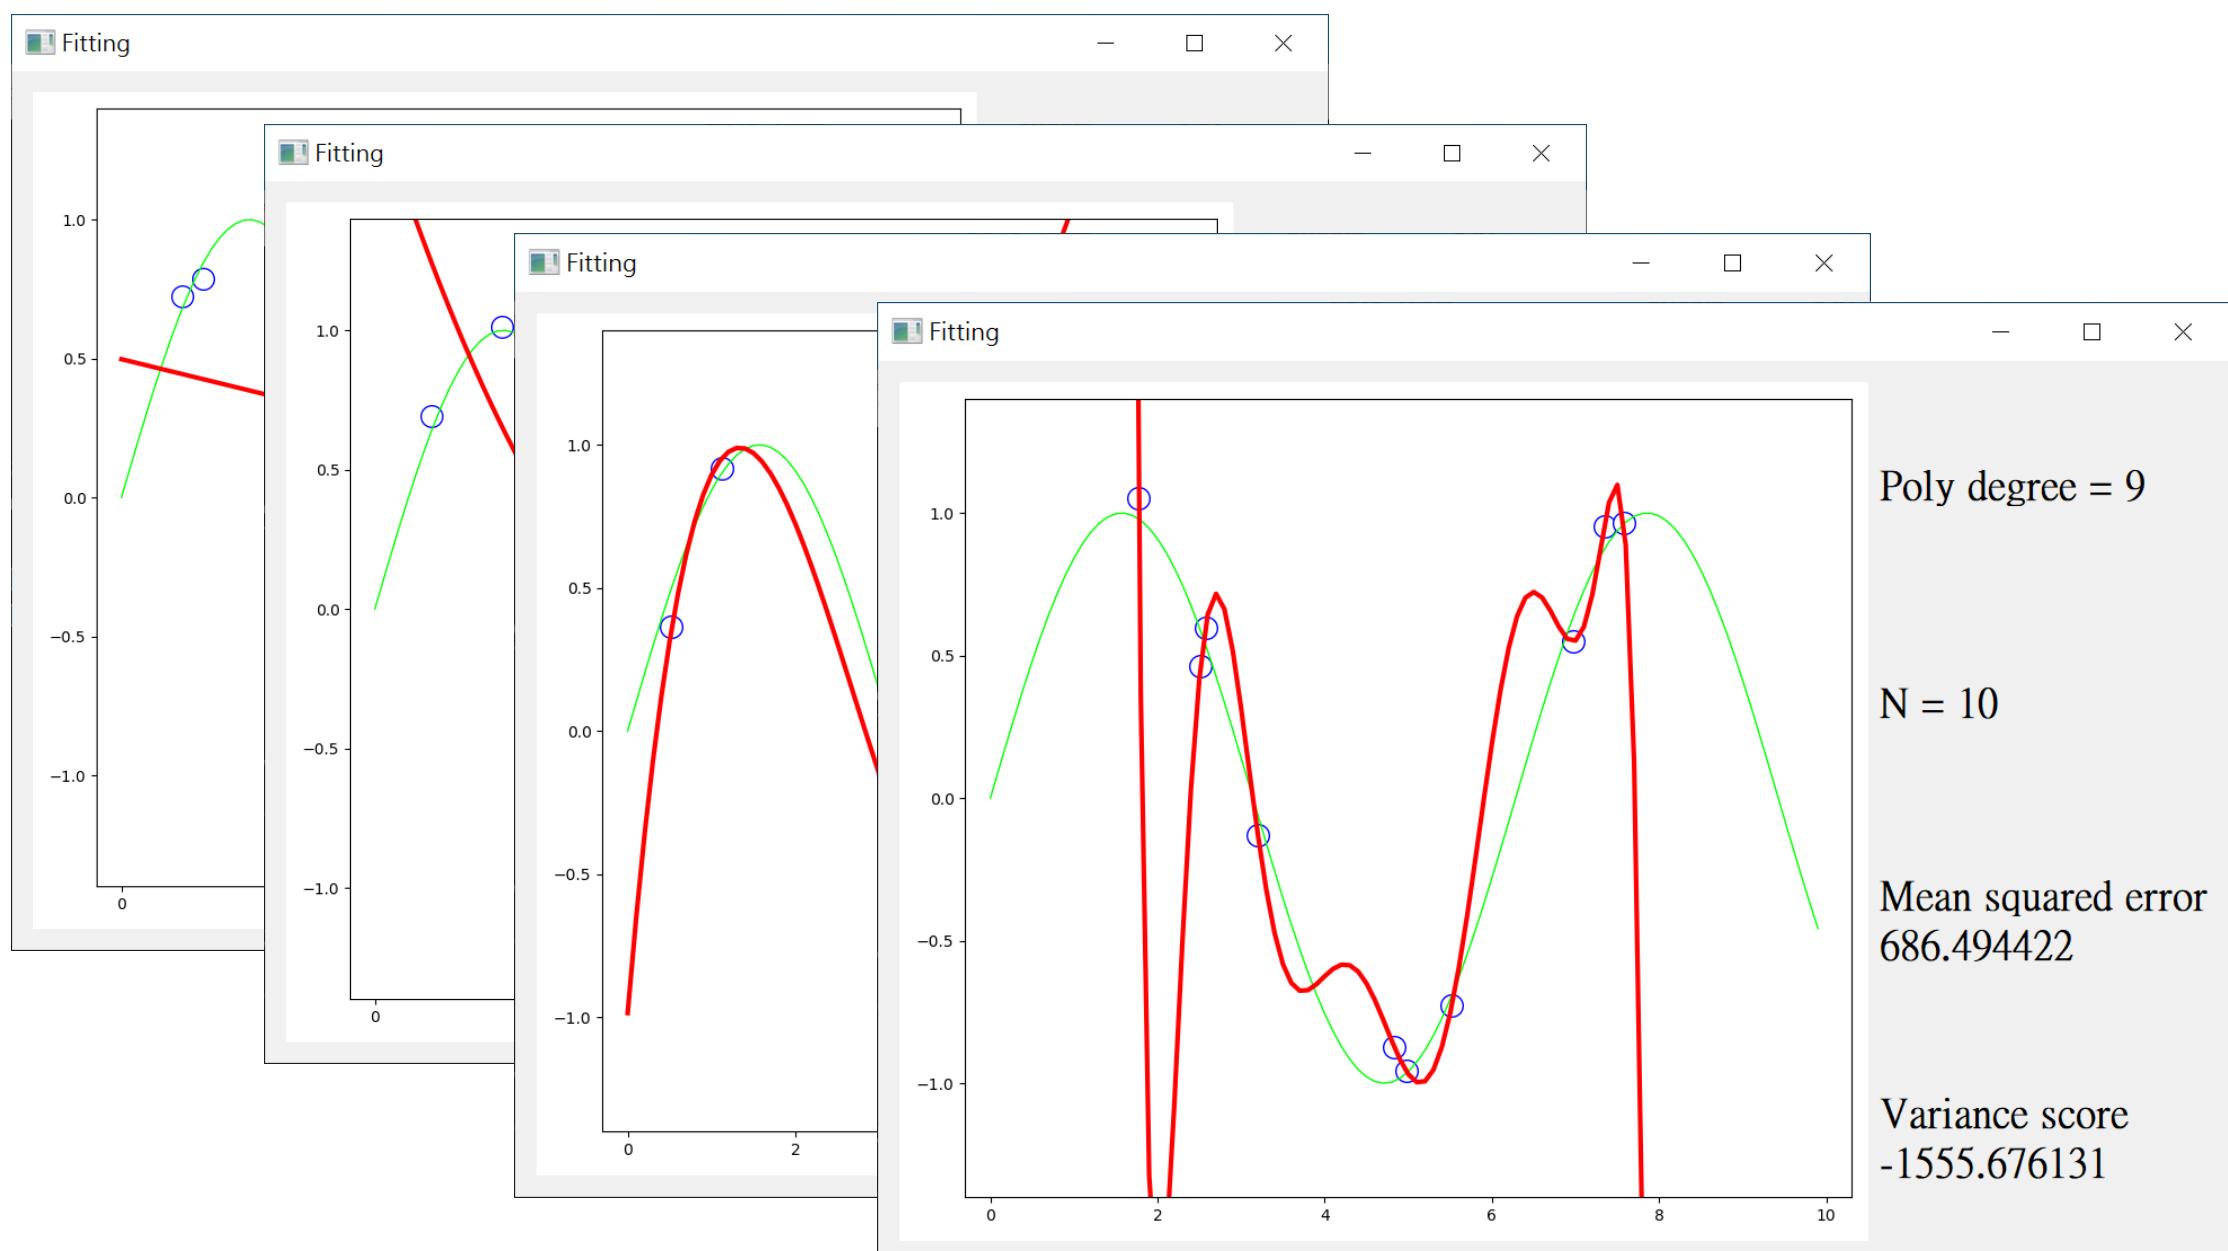
$$L(w) = \frac{1}{N} \sum_{i=1}^N (h(x^{(i)}; w, m) - \hat{y}^{(i)})^2$$

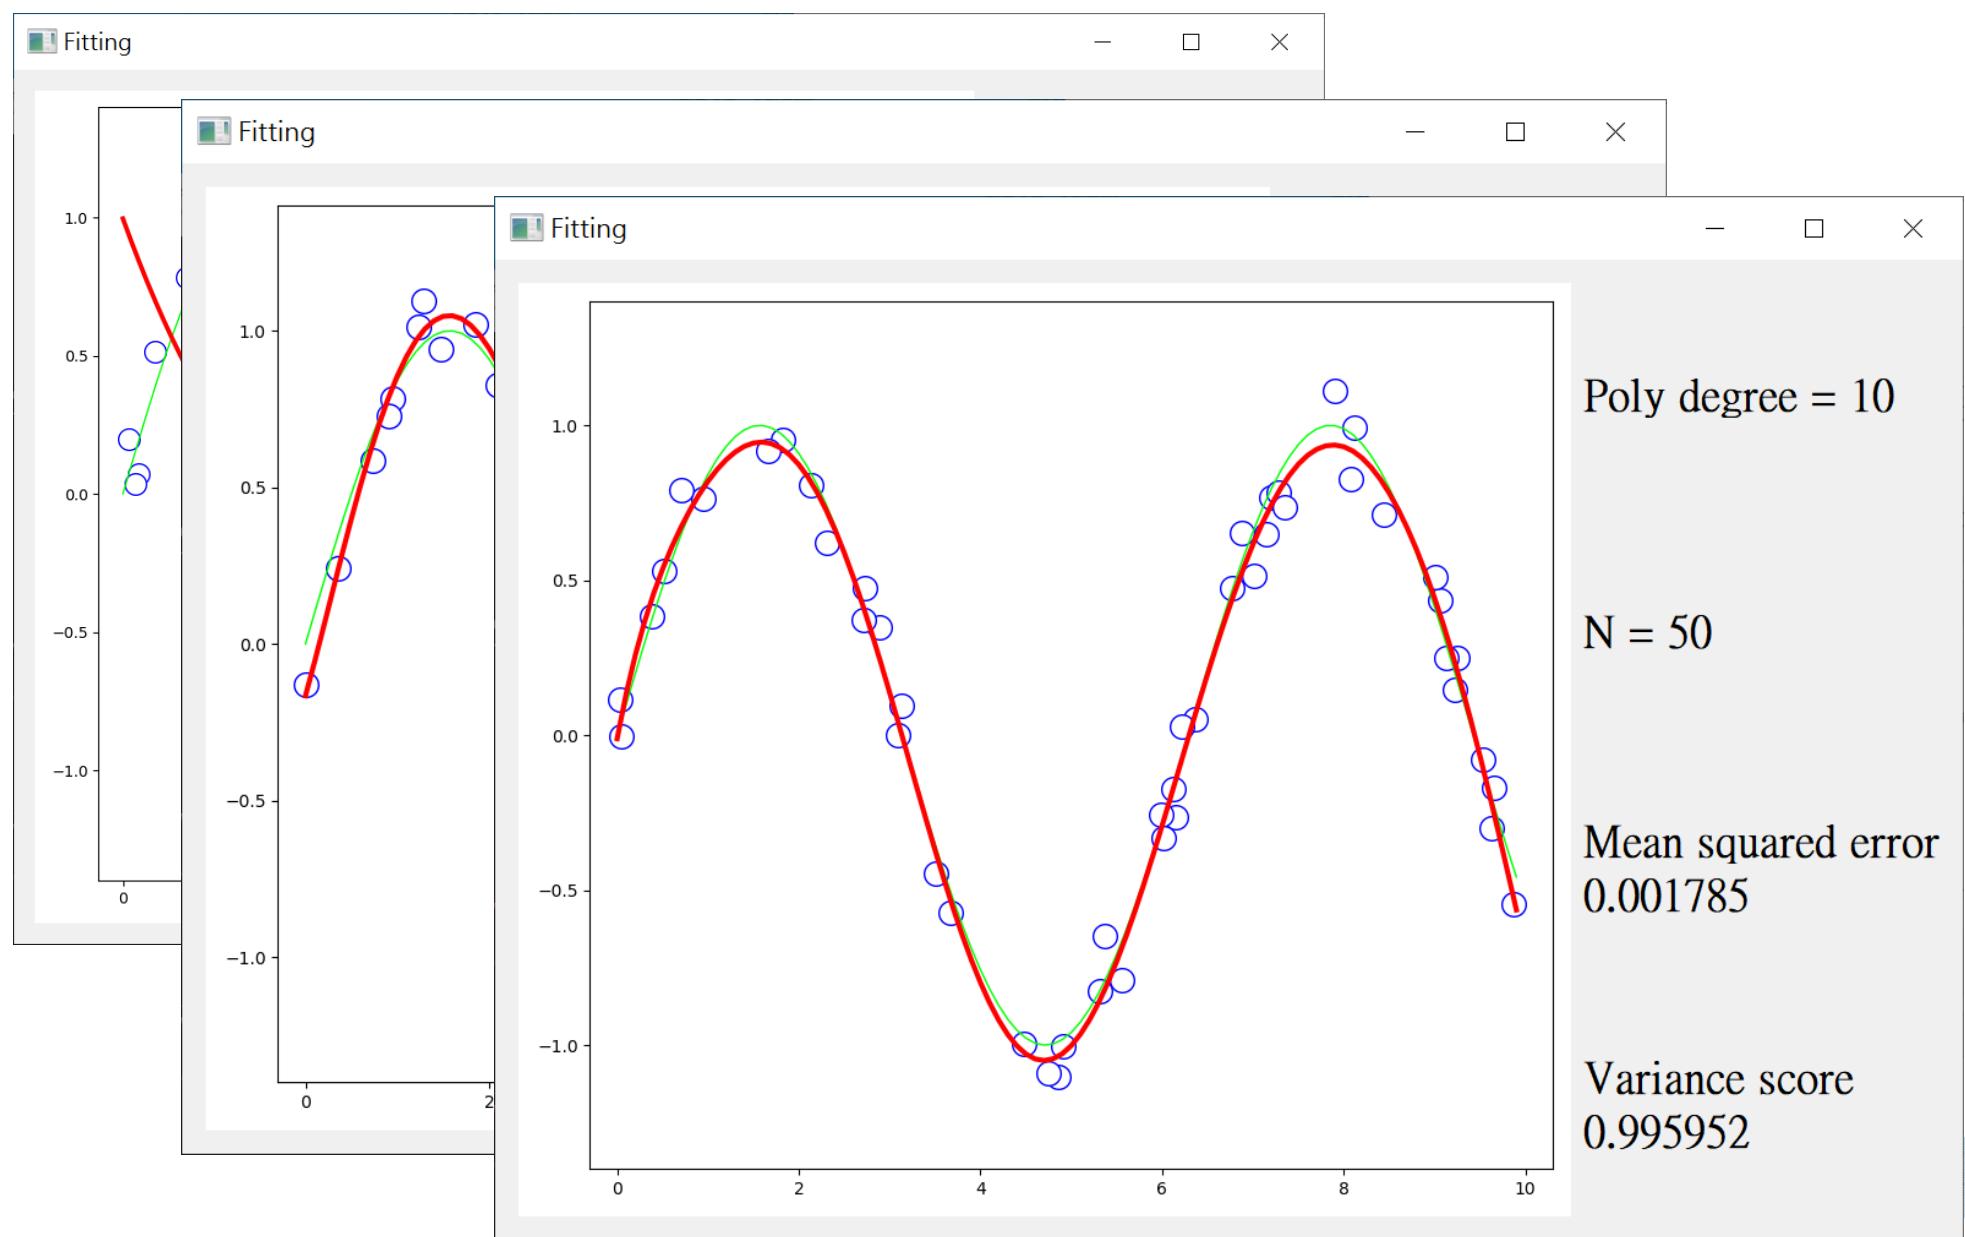


Sample size vs variance

- ◊ If the size of sample data is small, the model may vary a lot for different samples.

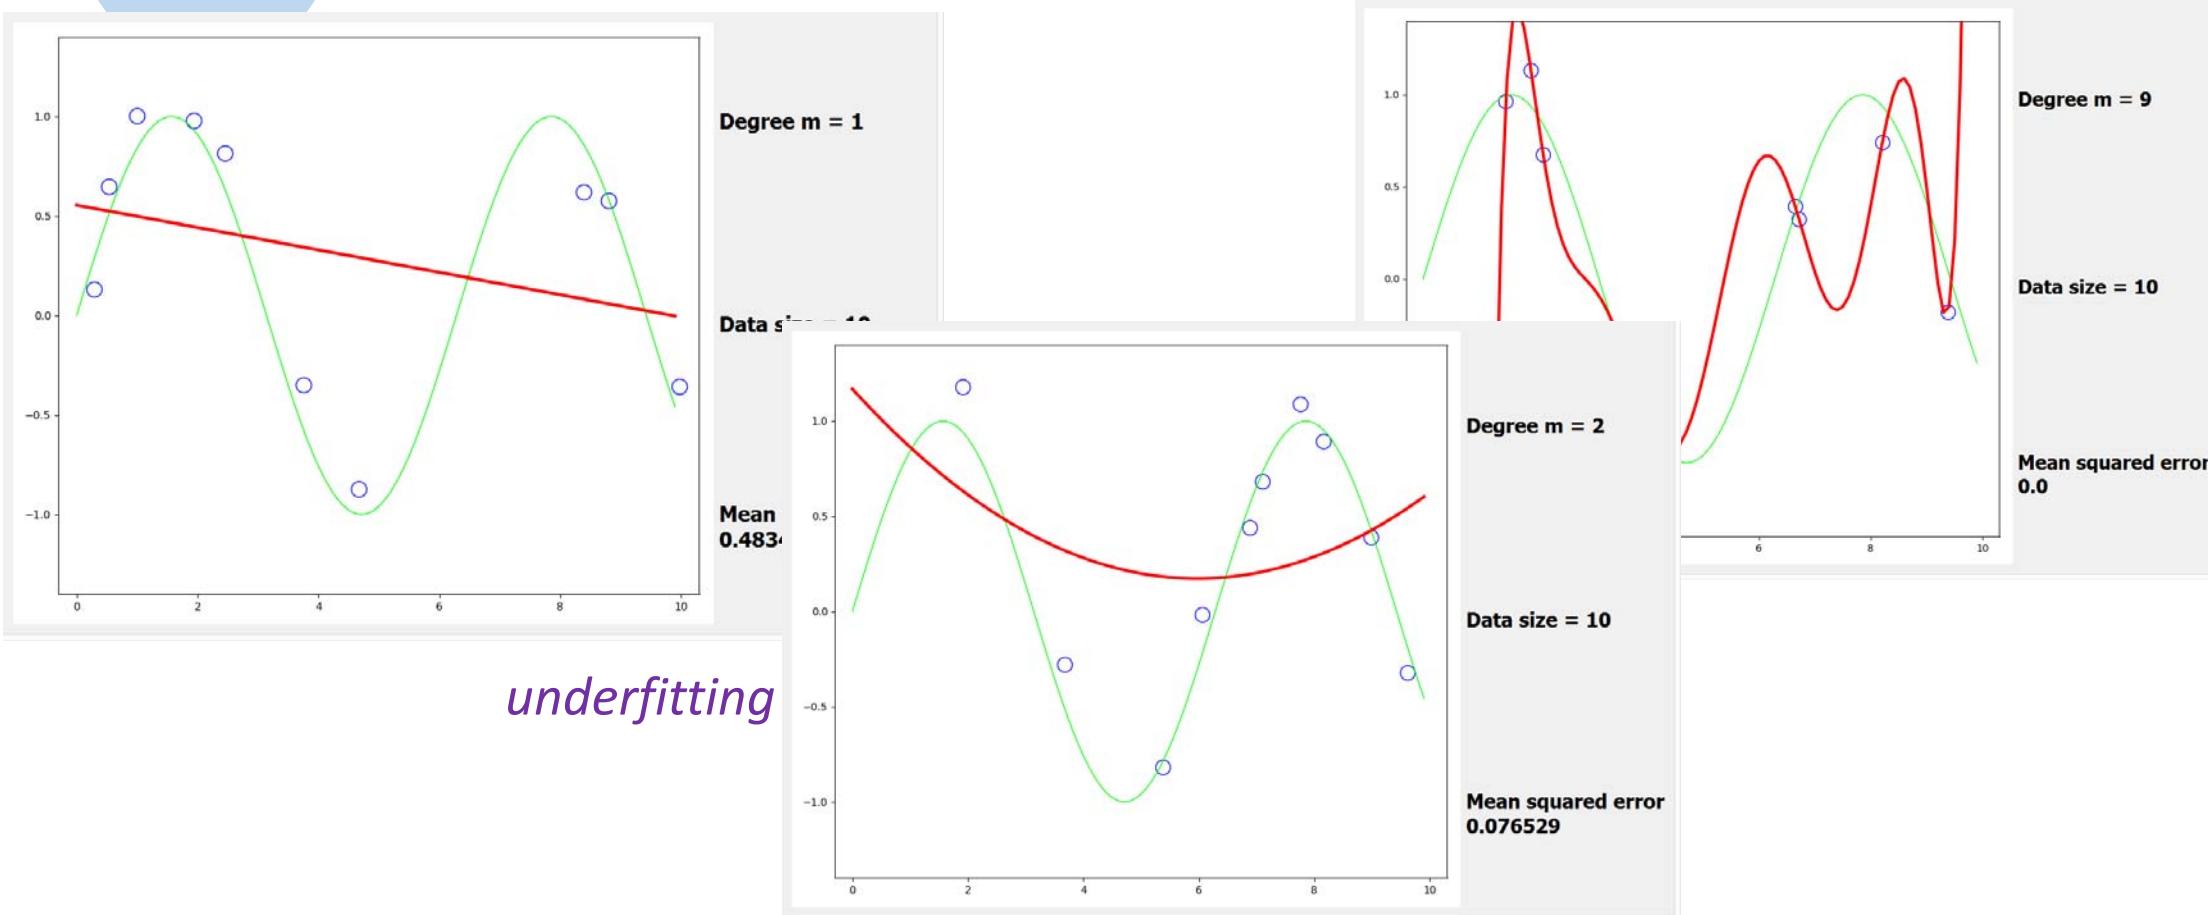






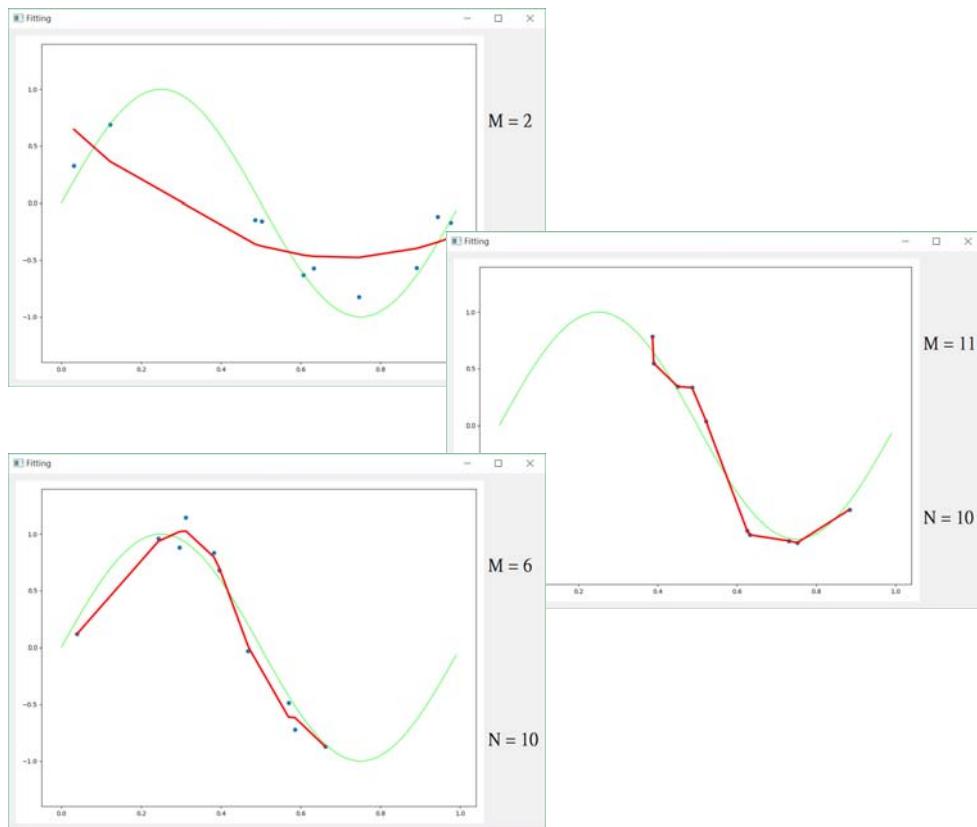
Polynomial Curve Fitting

overfitting



- **Lesson:** Complex models should only be used for large data samples.

Training results of weight parameters



Weight	M = 2	M = 6	M = 11
w_0	0	0	2.67850417
w_1	-3.48815	100.8412	-776755.1085
w_2	2.46723	-1046.57	4502951.55
w_3		5108.034	-14066940.4
w_4		-12573.3	24327608.55
w_5		14976.32	-18752447.23
w_6		-6864.09	-6136446.715
w_7			20332026.69
w_8			-2160559.429
w_9			-20347291.69
w_{10}			17932831.65
w_{11}			-4913022.544

Regularization for the control of overfitting

- ◊ The coefficient α governs the relative importance of the regularization term.
- ◊ Ridge regression, L2 regularization

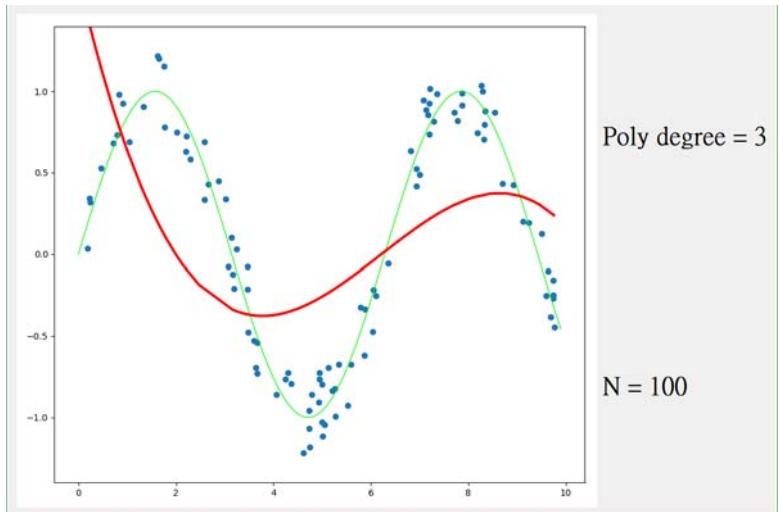
$$E(w) = \frac{1}{2} \sum_{i=1}^N (h(\mathbf{x}^{(i)}; w, m) - \hat{y}^{(i)})^2 + \alpha \|w\|^2$$

$$\|w\|^2 \equiv w^T w = w_0^2 + w_1^2 + \dots + w_M^2$$

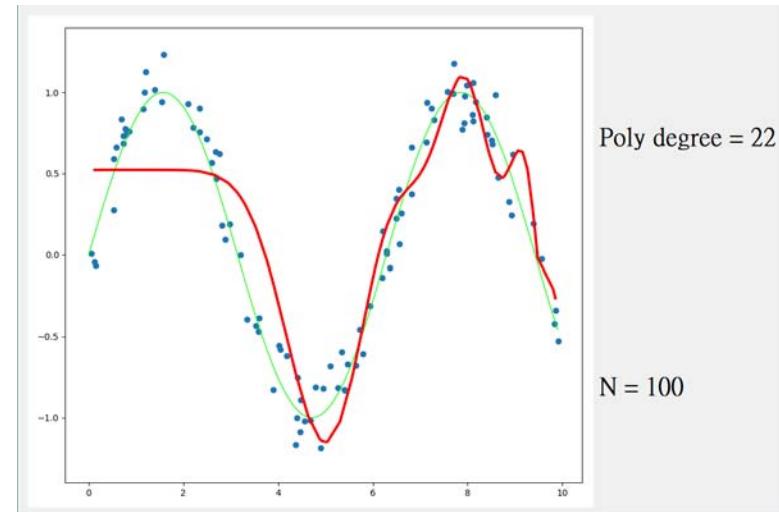
- ◊ Lasso regression, L1 regularization

$$E(w) = \frac{1}{2} \sum_{i=1}^N \{h(\mathbf{x}^{(i)}; w, m) - \hat{y}^{(i)}\}^2 + \alpha \|w\|^1$$

Over-fitting and Under-fitting



Small degree (simple model)
-> underfitting
-> Less variance



Large degree (complex model)
-> overfitting
-> Less bias

Model complexity

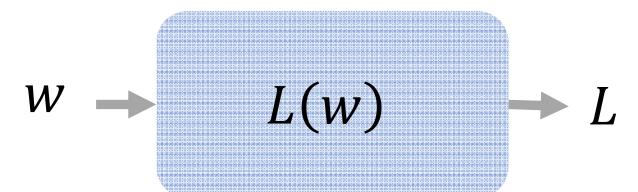
- ❖ We should choose the model with suitable complexity for the size of the given dataset.
 - ❖ High complexity -> more freedom -> low bias with potential high variance
 - ❖ Low complexity -> less freedom -> low variance with potential high bias
- ❖ Should we use all the features in the given dataset?

How to solve function's minimum/maximum

❖ Suppose we want to find the parameter w to minimize $L(w)$

- ◆ For example: $L(w) = \frac{1}{N} \sum_{i=1}^N \{h(x^{(i)}; w, m) - \hat{y}^{(i)}\}^2$
- ◆ The most direct naïve approach is to find many w candidates.
- ◆ What if we have multiple parameters to solve?

$$L(w_1, w_2, w_3, \dots, w_{100})$$



- ❖ the number of candidates will be very huge. Ex: 10^{100}
- ◆ The actual weights to train could be up to millions.

Gradient Descent Optimization Method

for k in range (max_iteration):

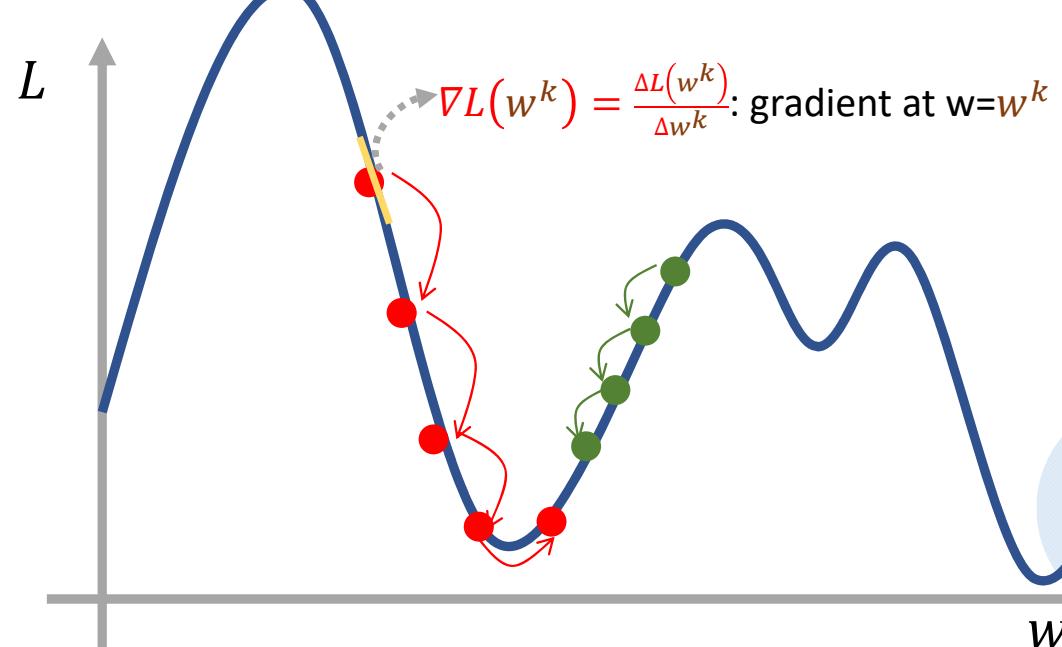
$$w^{k+1} = w^k - \eta \nabla L(w^k) \quad \text{\# update weight}$$

$$\text{if } \|w^{k+1} - w^k\| < \epsilon \quad \text{\# stopping criteria}$$

break

η : adjustment (learning) rate

❖ If the increase of W can reduce the loss, W will be increased.



Gradient Descent

Training data: $(x^{(i)}, \hat{y}^{(i)})$

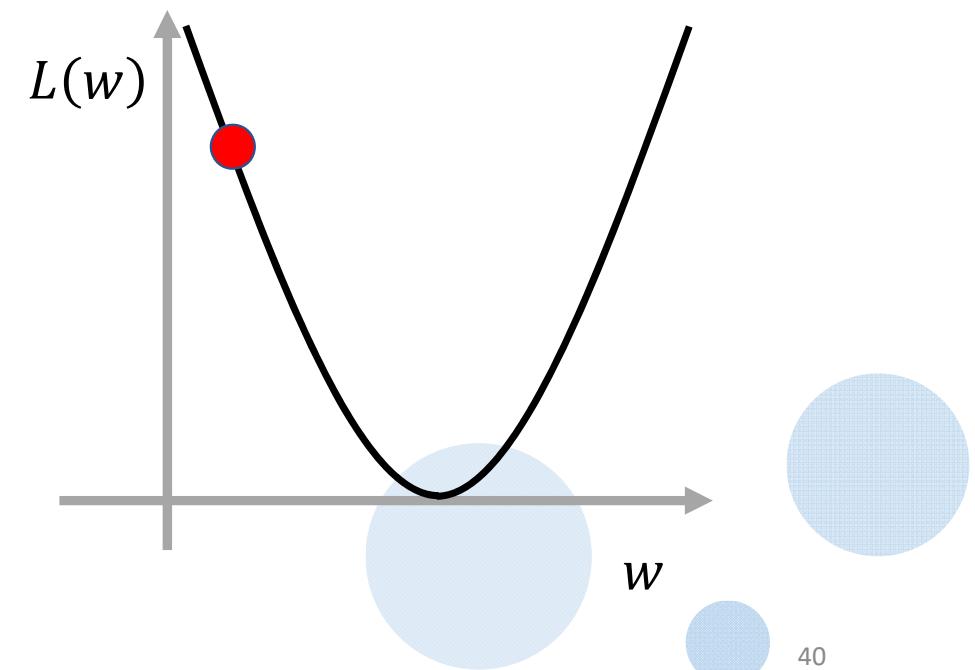
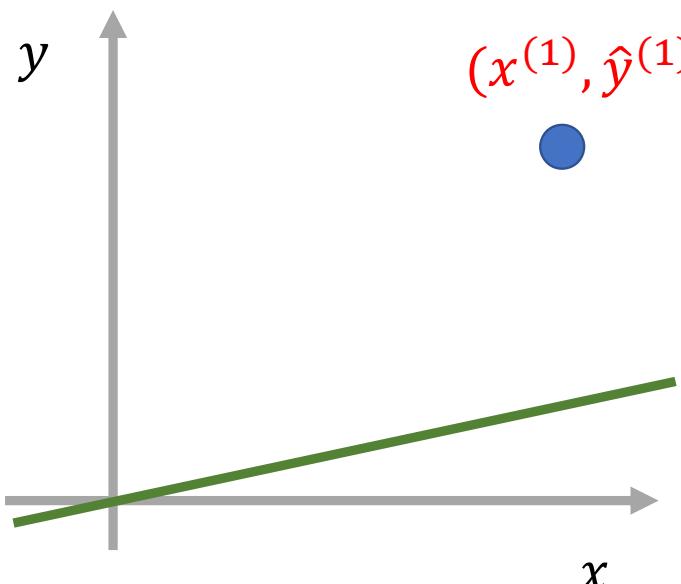
$$h(x; w, m) = x \times w$$

$$L(w) = (x^{(1)} \times w - \hat{y}^{(1)})^2$$

Model

Epoch = 0

Loss

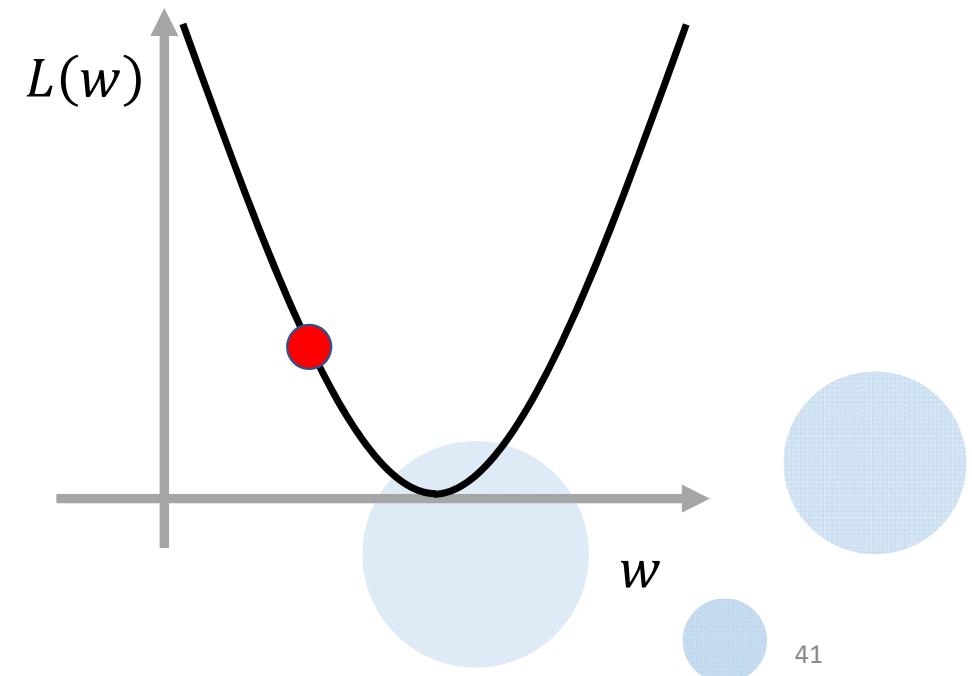
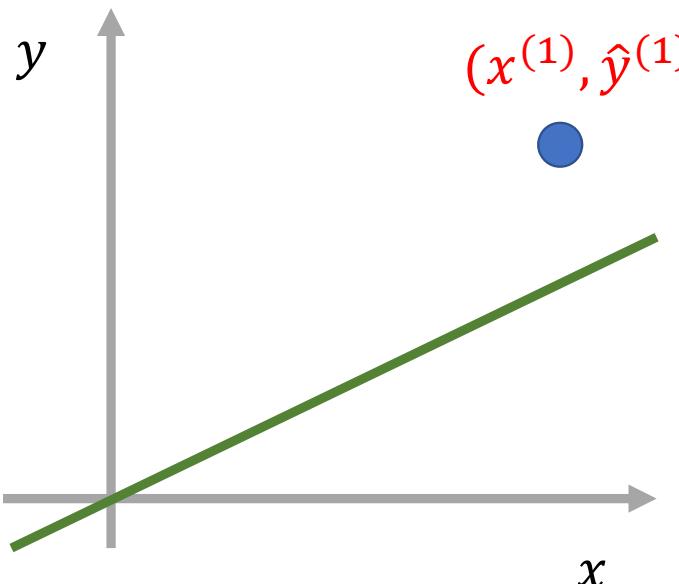


Gradient Descent

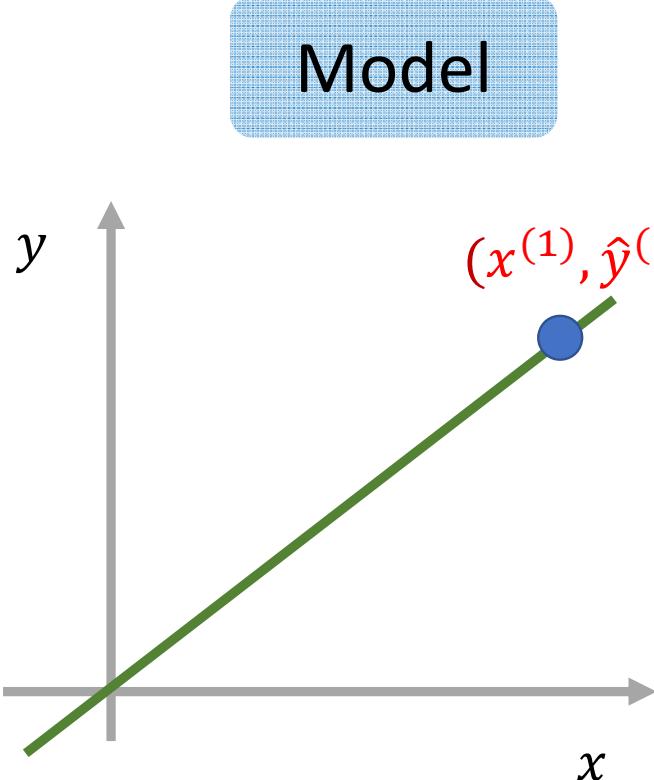
Model

Epoch = 10

Loss



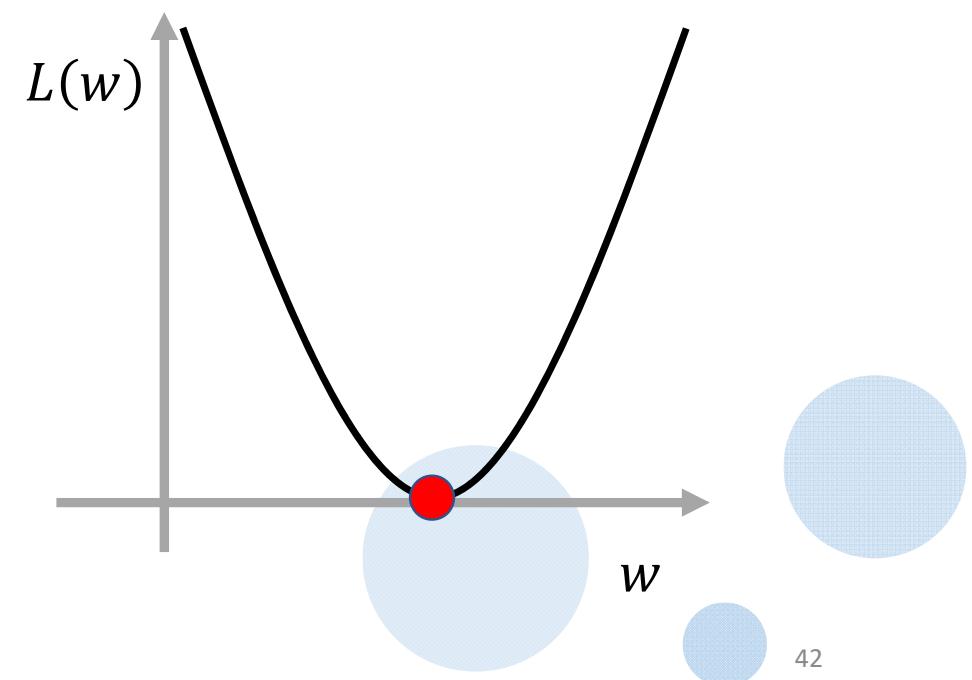
Gradient Descent



Model

Epoch = 100

Loss



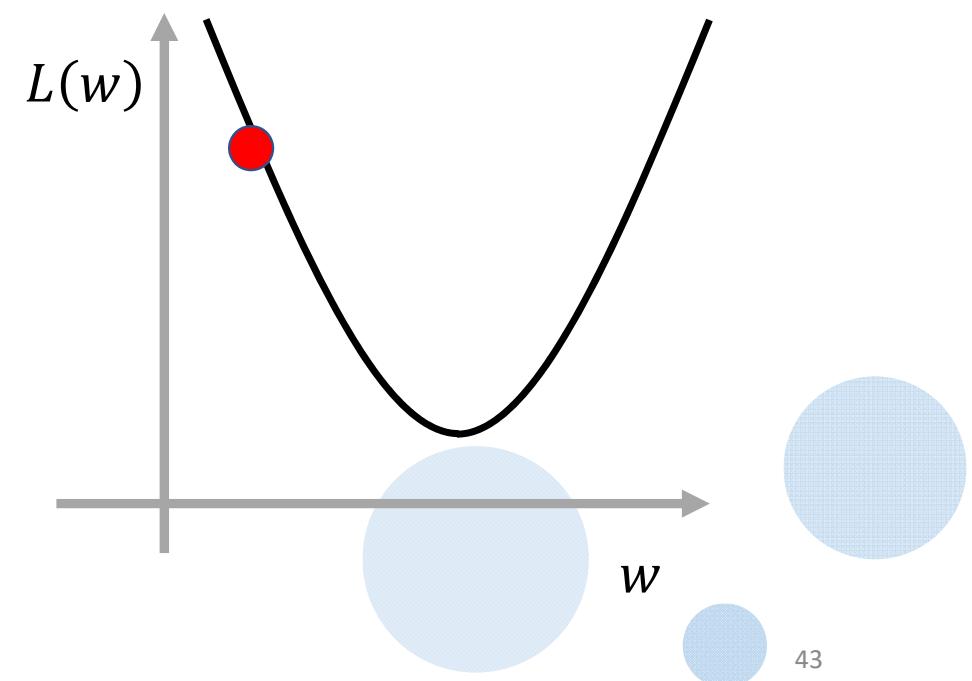
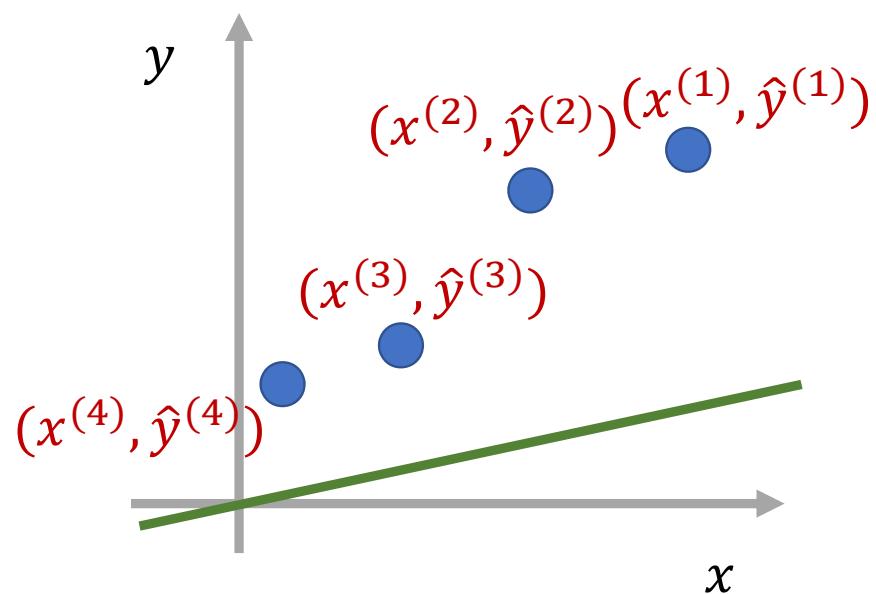
Gradient Descent

$$L(w) = \sum_{i=1}^4 (x^{(i)} \times w - \hat{y}^{(i)})^2$$

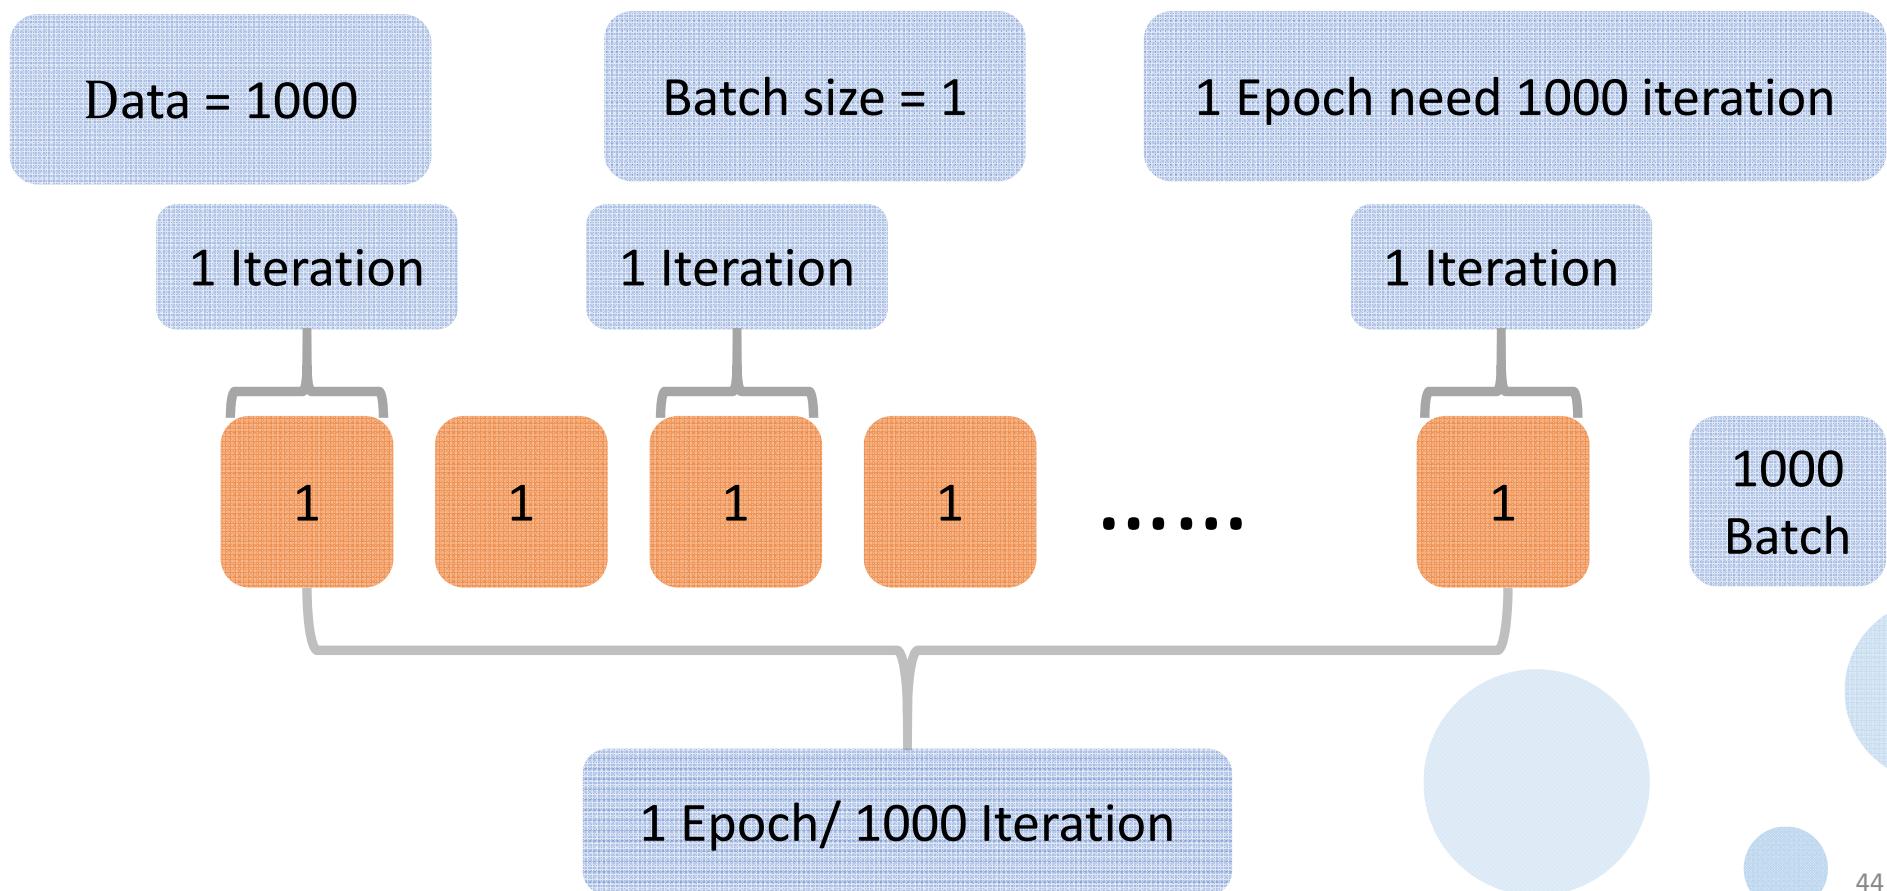
Model

Epoch = 0

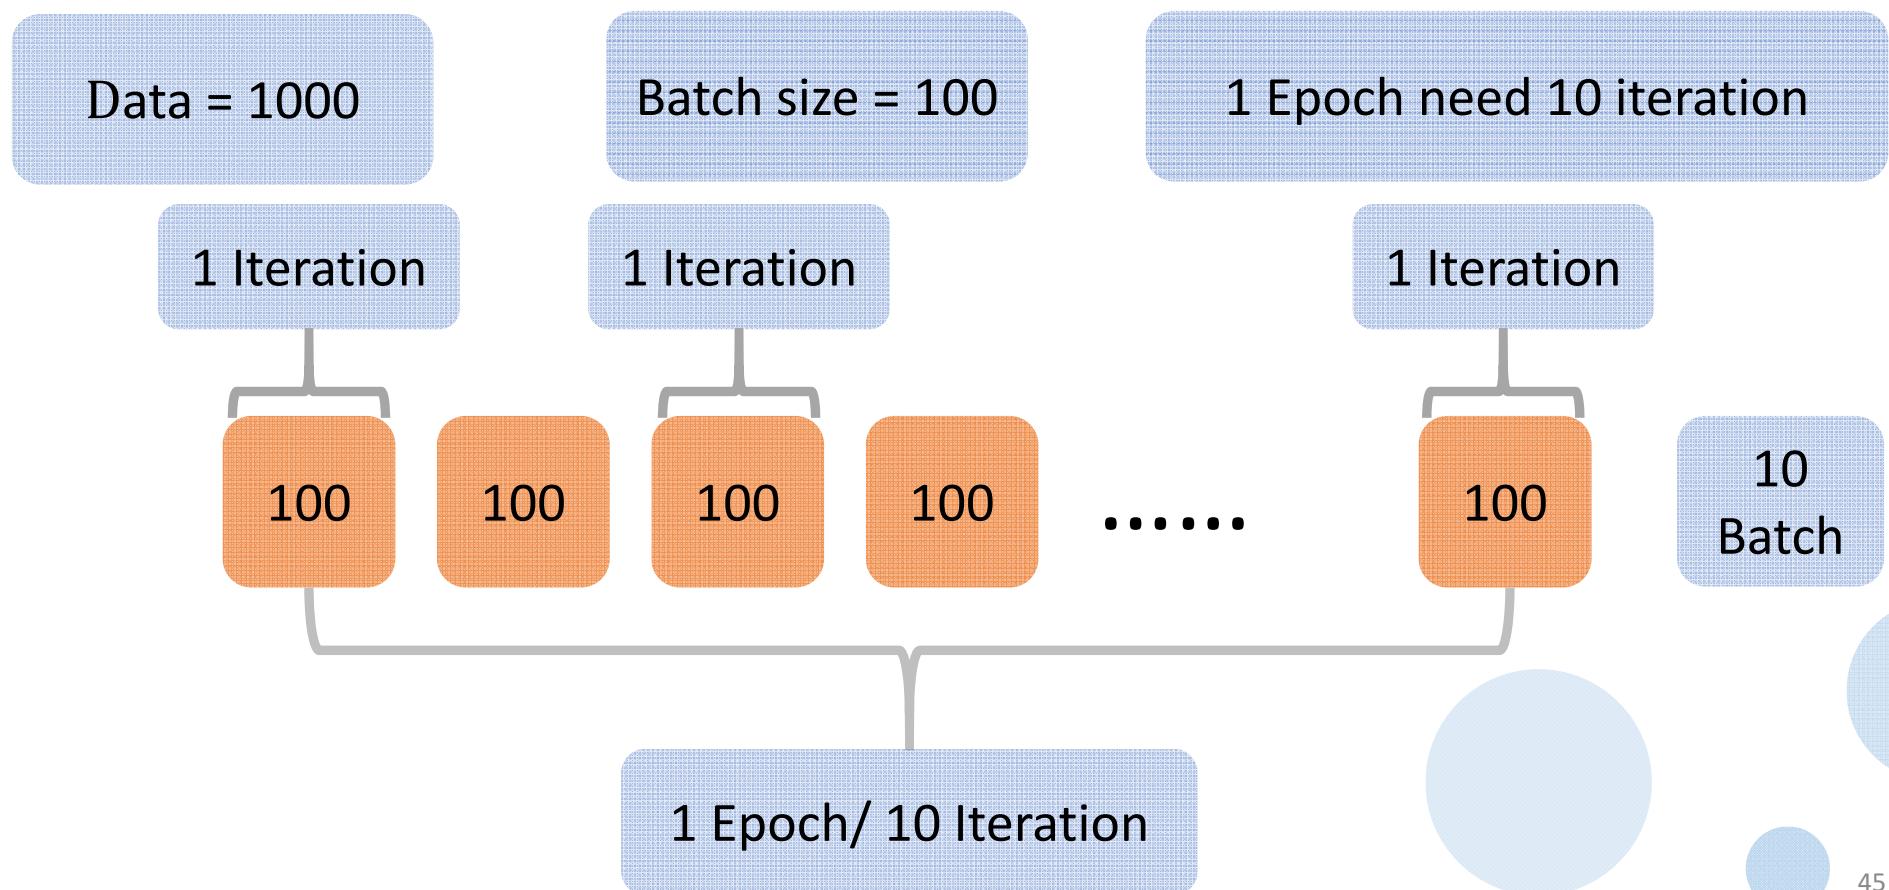
Loss



Batch, Epoch, Iteration

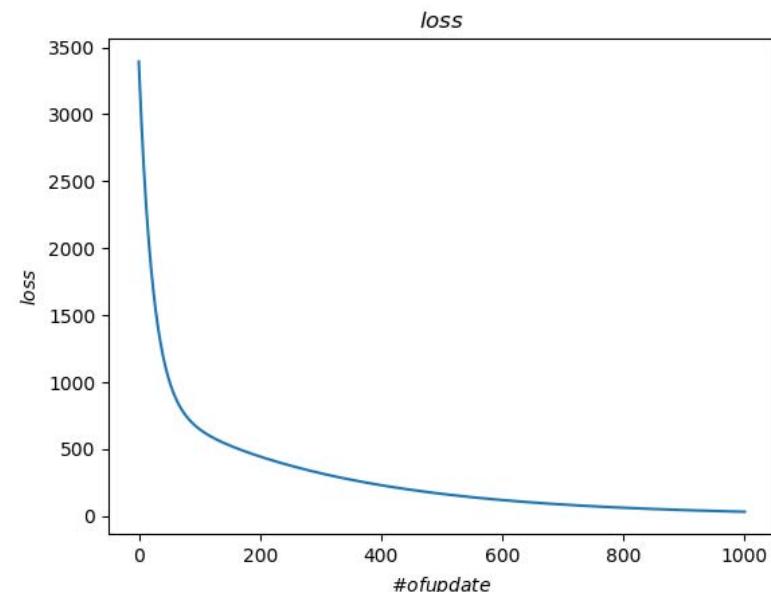
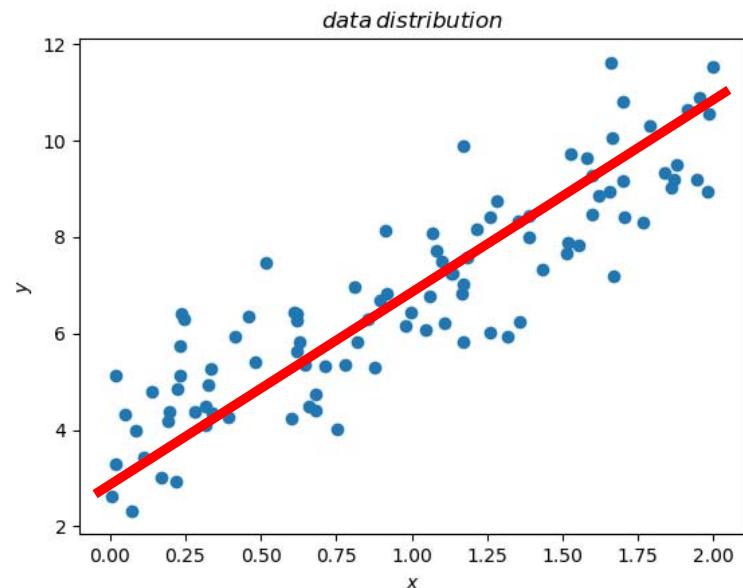


Batch, Epoch, Iteration



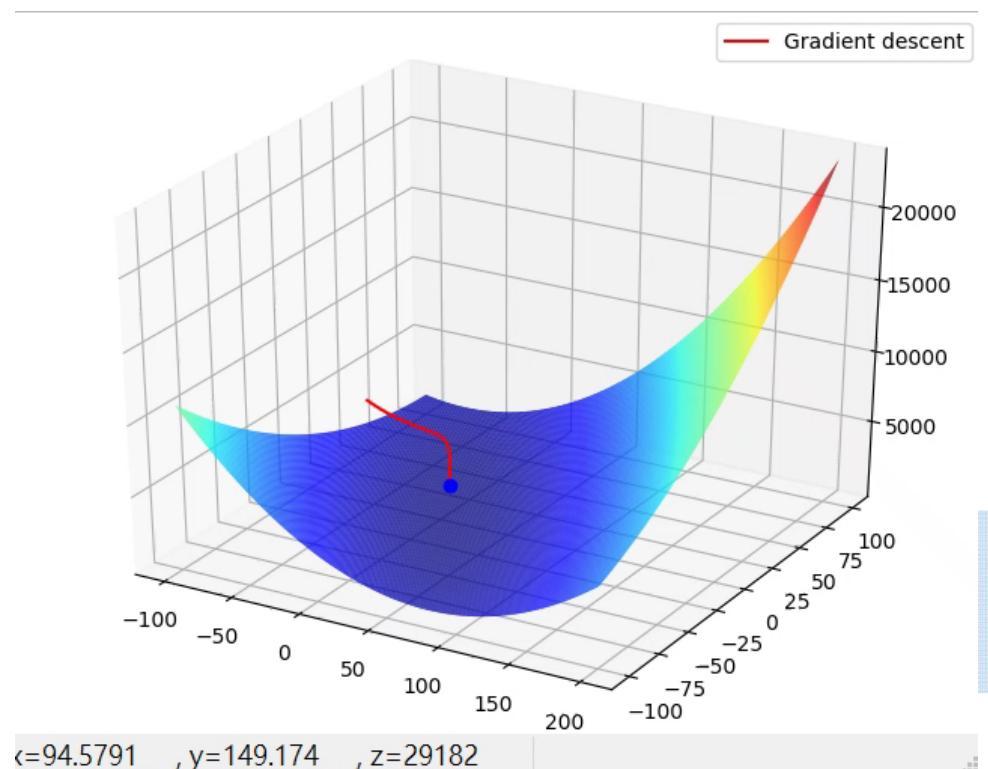
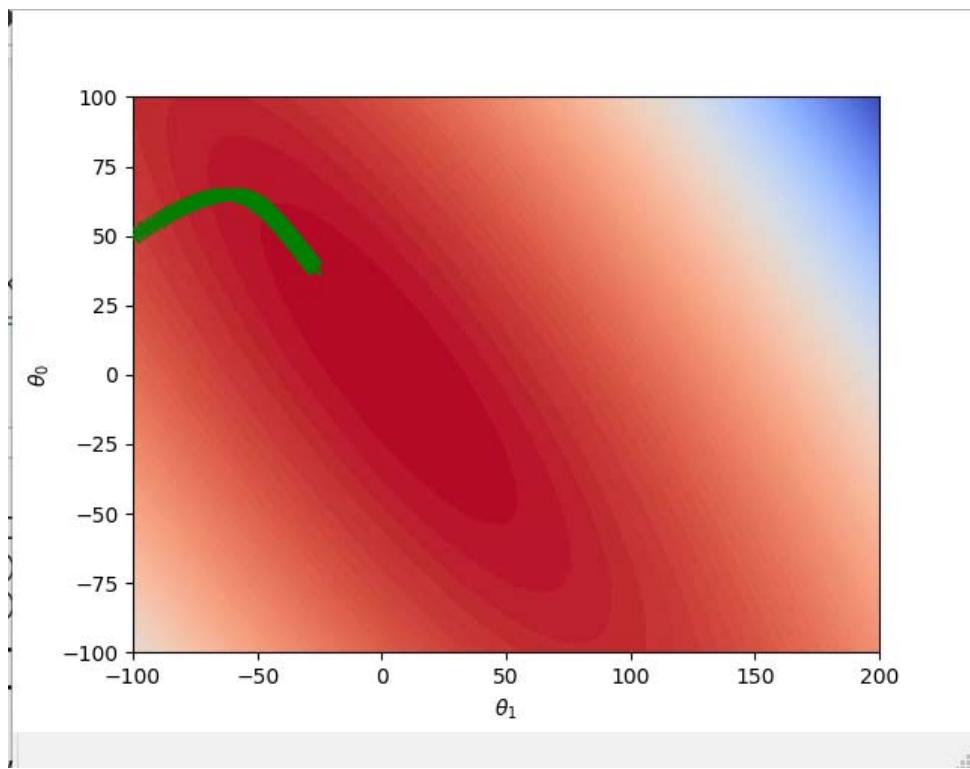
Linear Regression Training process

◆ Data and loss



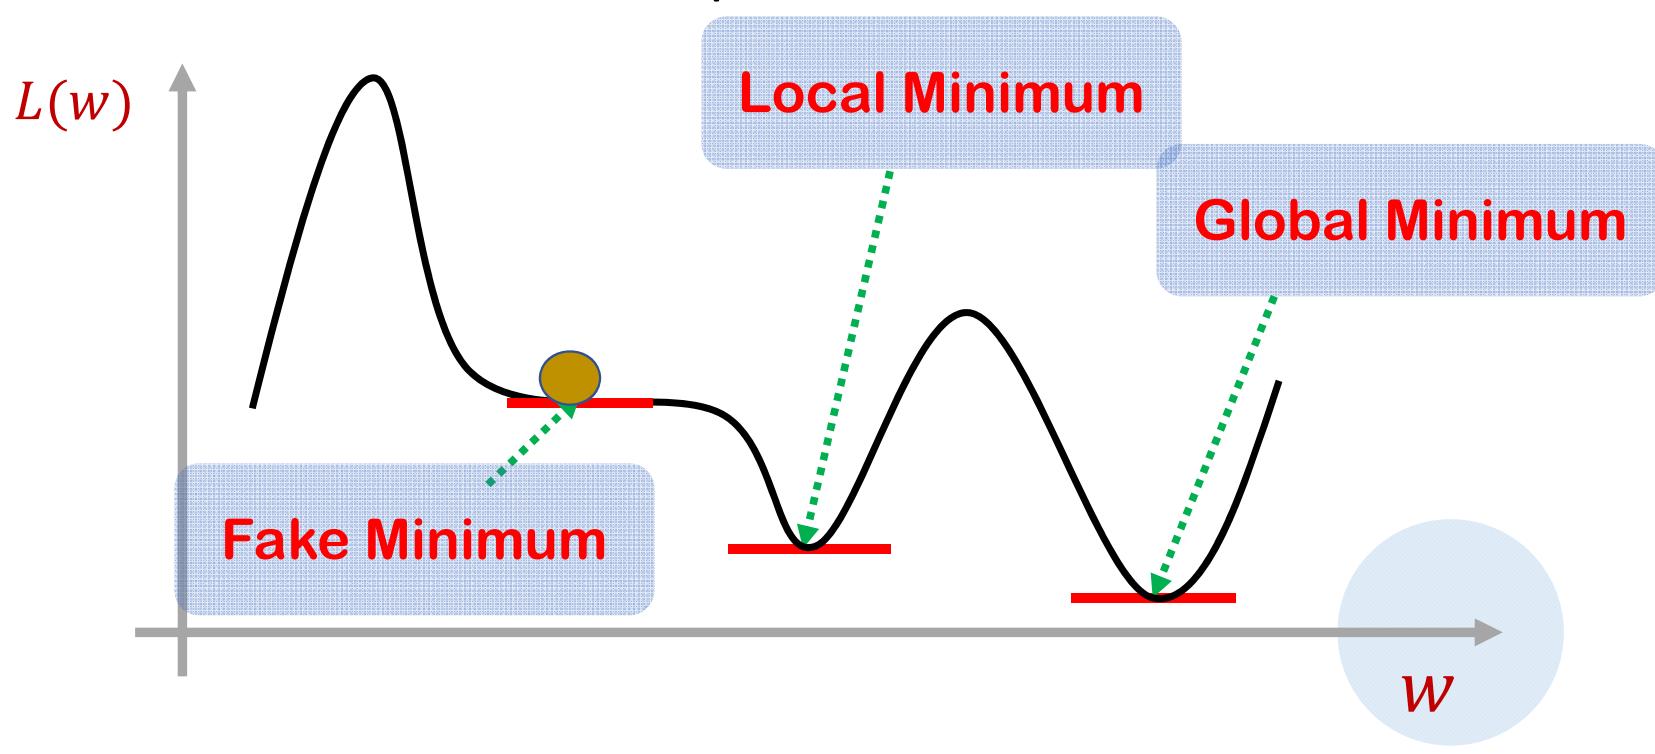
Linear Regression Training process

◆ Training process

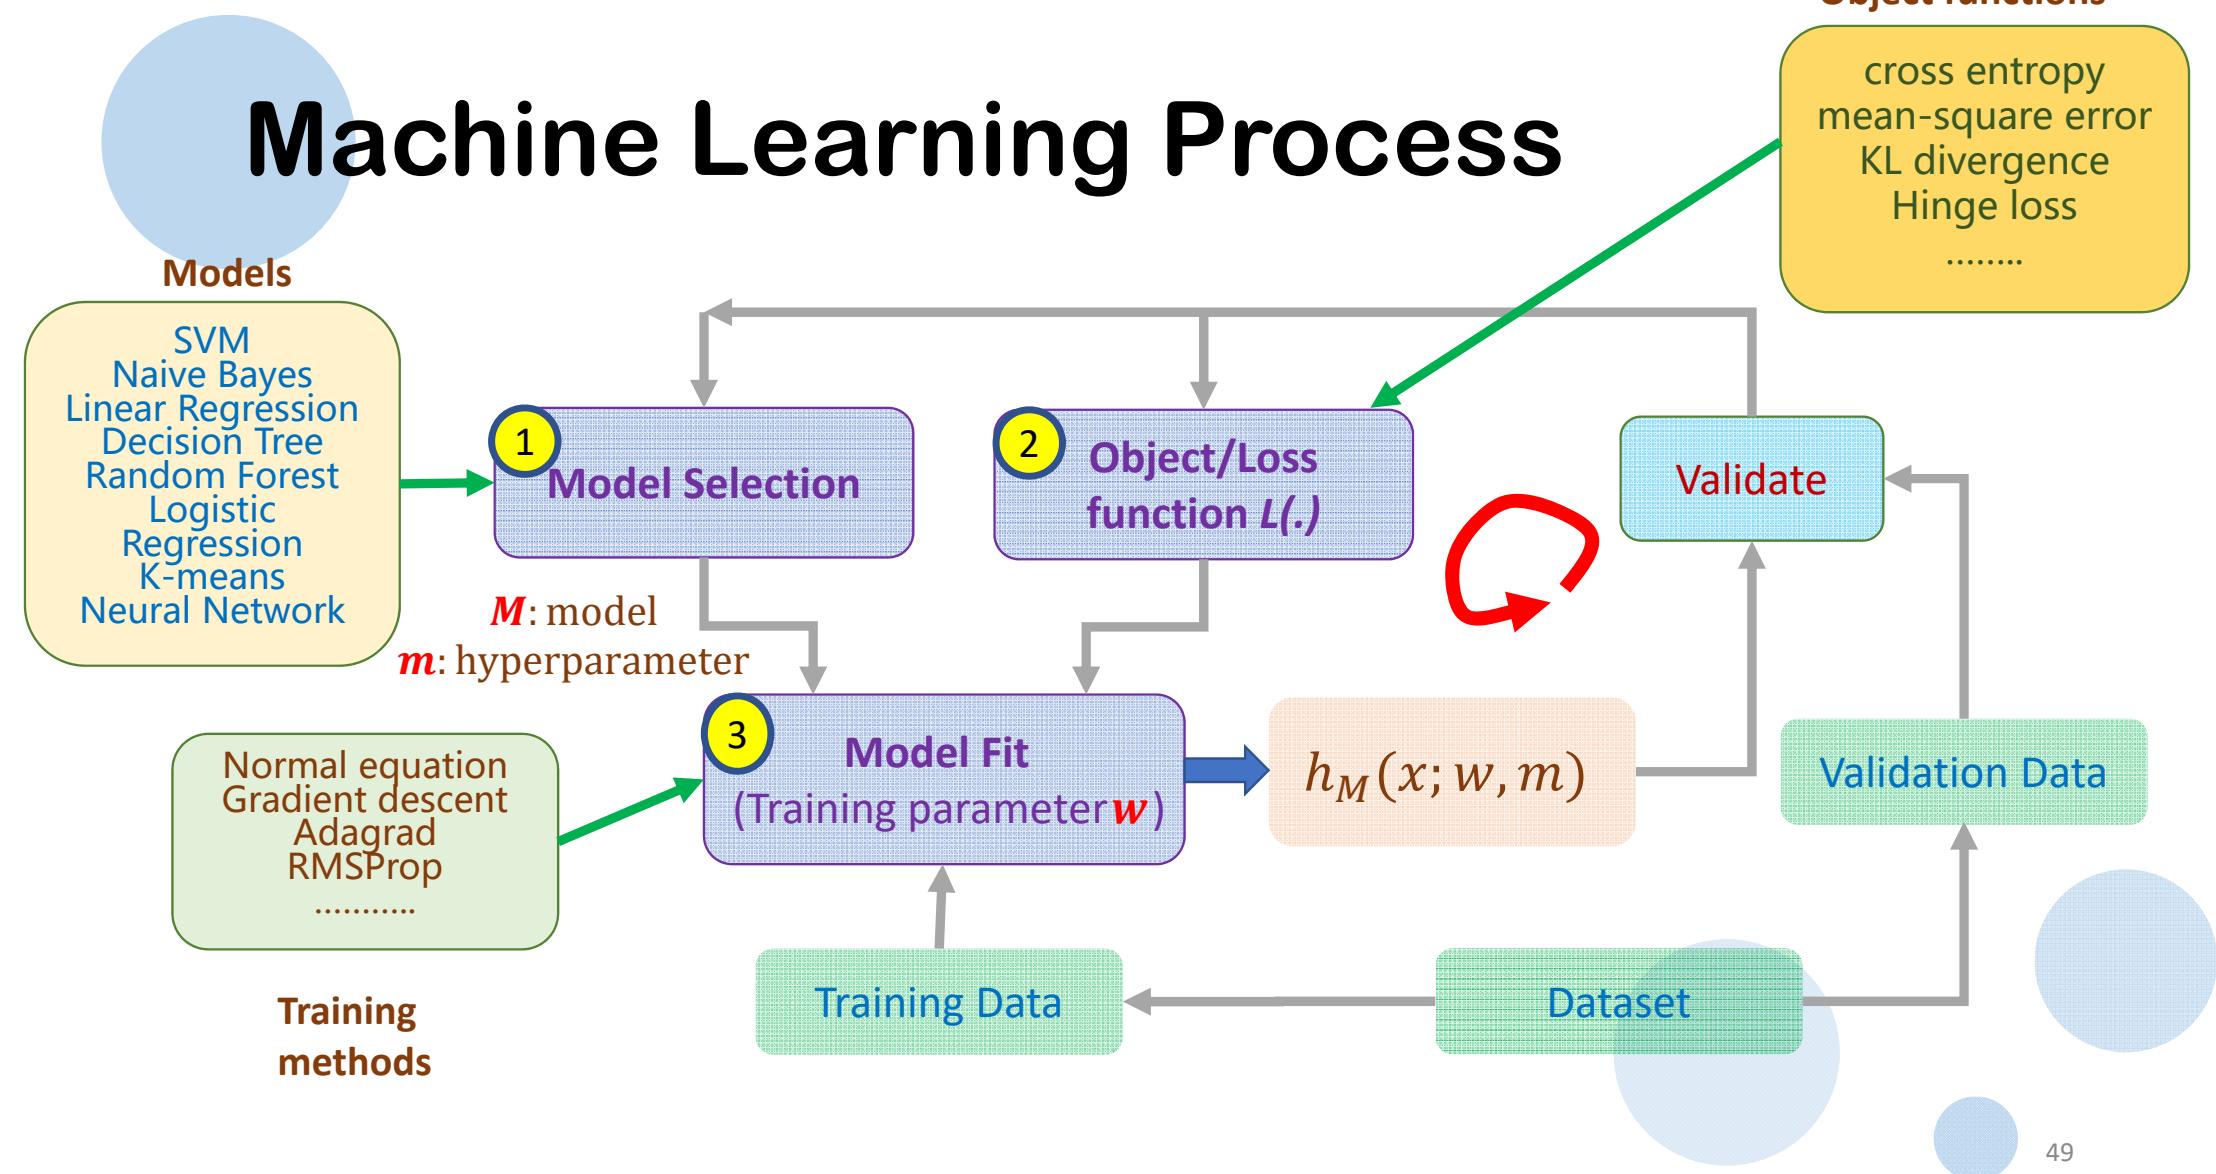


Problems of Gradient Descent

- ❖ Stuck at the local minimum points.
- ❖ Oscillate around the minimum point.
- ❖ Stuck at the saddle points.
- ❖ Very slow update at plateau.



Machine Learning Process



Ex & Hyper (meta) -parameters

❖ Setting of sklearn decision-tree module

sklearn.tree.DecisionTreeClassifier

```
class sklearn.tree.DecisionTreeClassifier (criterion='gini', splitter='best', max_depth=None,  
min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features=None, random_state=None,  
max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, class_weight=None, presort=False)
```

[source]

❖ Setting of sklearn SVM module

sklearn.svm.SVC

```
class sklearn.svm.SVC (C=1.0, kernel='rbf', degree=3, gamma='auto_deprecated', coef0=0.0, shrinking=True,  
probability=False, tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1,  
decision_function_shape='ovr', random_state=None)
```

[source]

Import package

Load dataset

Choose model
Set hyper para

```
from sklearn import svm  
from sklearn import datasets
```

```
#digit dataset from sklearn  
digits = datasets.load_digits()
```

```
#creat the LinearRegression model  
clf = svm.SVC(C=1.0, kernel='linear')
```

```
#set training set  
x, y = digits.data[:-1], digits.target[:-1]
```

#train model

```
clf.fit(x,y)
```

Set label

#predict

```
y_pred = clf.predict([digits.data[-1]])  
y_true = digits.target[-1]
```

Fit model

```
print(y_pred)  
print(y_true)
```

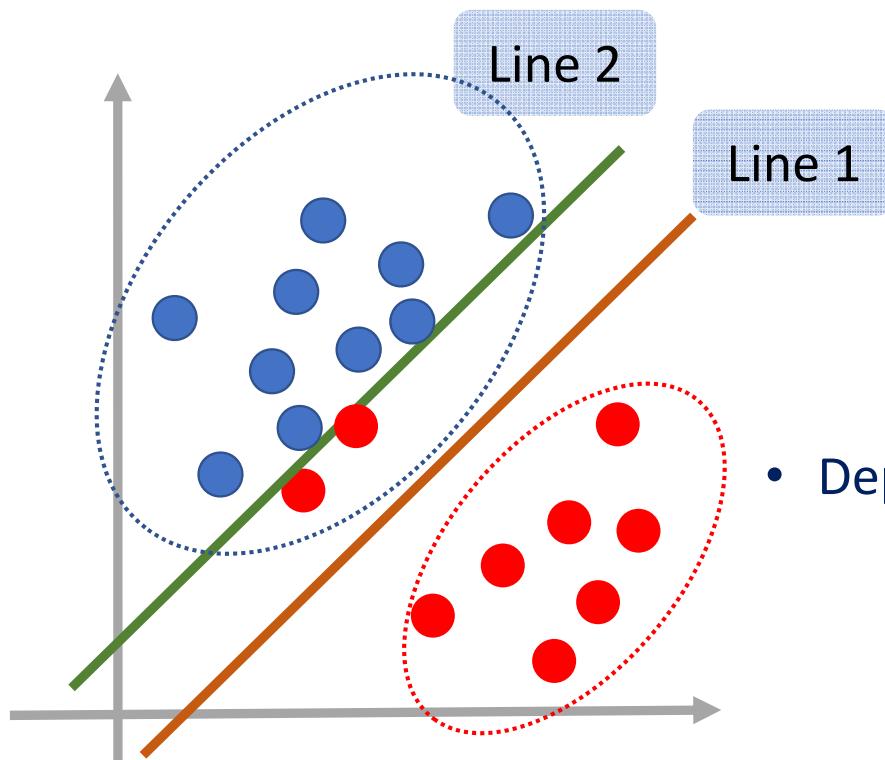
Inference

[8]

8

Display result

Which separation line is better for binary classification ?



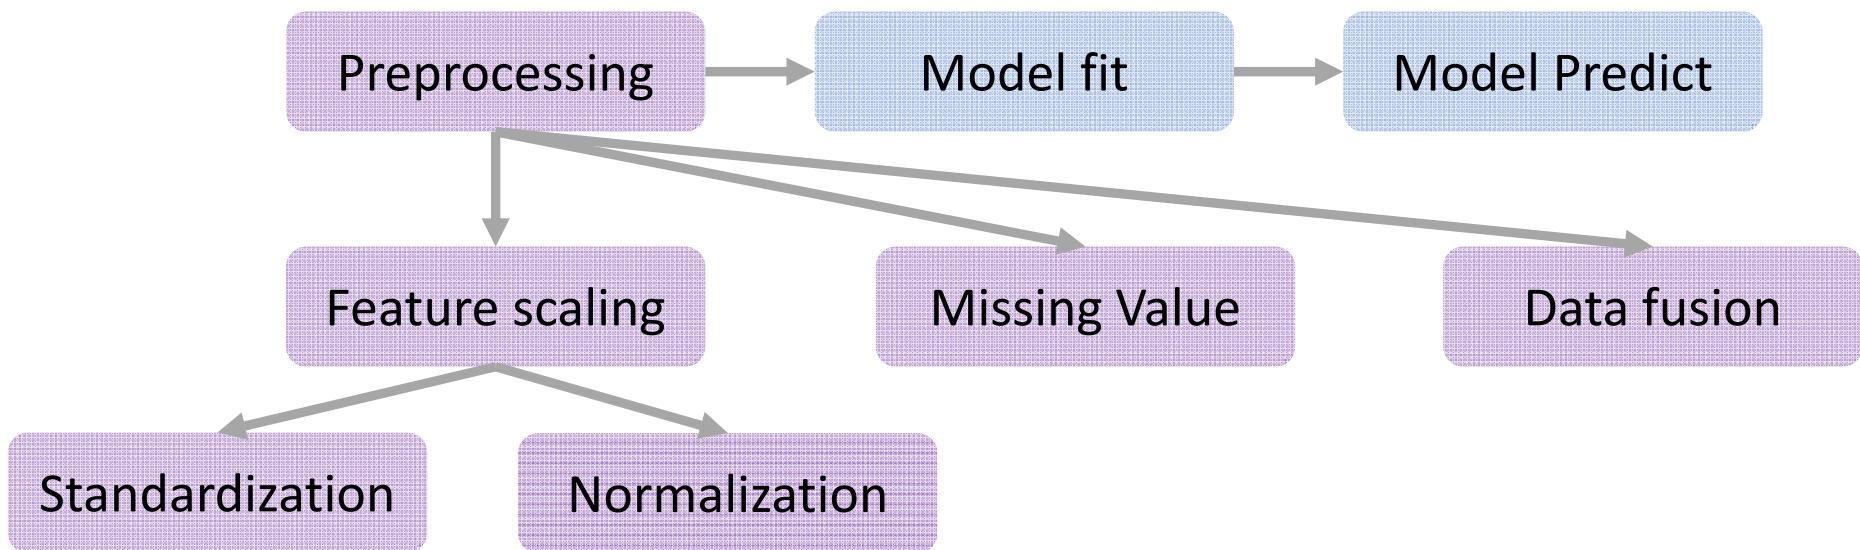
- Think about how the clustering approach (**unsupervised**) classifies.
- Depend on your defined **goodness**



3

Data preprocessing

Preprocessing before training



Preprocessing package



scikit
learn

Home Installation Documentation Examples

Google Custom Search

Previous 5.2 Feature Next 5.4 Up 5. Datasets

5.2. Feature extraction

scikit-learn v0.21.2 Other versions

Please cite us if you use the software.

«

5.3. Preprocessing data

5.3.1. Standardization, or mean removal and variance scaling

- 5.3.1.1. Scaling features to a range
- 5.3.1.2. Scaling sparse data
- 5.3.1.3. Scaling data with outliers
- 5.3.1.4. Centering kernel matrices

5.3.2. Non-linear transformation

- 5.3.2.1. Mapping to a Uniform distribution
- 5.3.2.2. Mapping to a Gaussian distribution

5.3.3. Normalization

5.3.4. Encoding categorical features

5.3.5. Discretization

- 5.3.5.1. K-bins discretization
- 5.3.5.2. Feature binarization

5.3.6. Imputation of missing values

5.3.7. Generating polynomial features

5.3. Preprocessing data

The `sklearn.preprocessing` package provides several common utility functions and transformer classes to change raw feature vectors into a representation that is more suitable for the downstream estimators.

In general, learning algorithms benefit from standardization of the data set. If some outliers are present in the set, robust scalers or transformers are more appropriate. The behaviors of the different scalers, transformers, and normalizers on a dataset containing marginal outliers is highlighted in [Compare the effect of different scalers on data with outliers](#).

5.3.1. Standardization, or mean removal and variance scaling

Standardization of datasets is a **common requirement for many machine learning estimators** implemented in scikit-learn; they might behave badly if the individual features do not more or less look like standard normally distributed data: Gaussian with **zero mean and unit variance**.

In practice we often ignore the shape of the distribution and just transform the data to center it by removing the mean value of each feature, then scale it by dividing non-constant features by their standard deviation.

For instance, many elements used in the objective function of a learning algorithm (such as the RBF kernel of Support Vector Machines or the l1 and l2 regularizers of linear models) assume that all features are centered around zero and have variance in the same order. If a feature has a variance that is orders of magnitude larger than others, it might dominate the objective function and make the estimator unable to learn from other features correctly as expected.

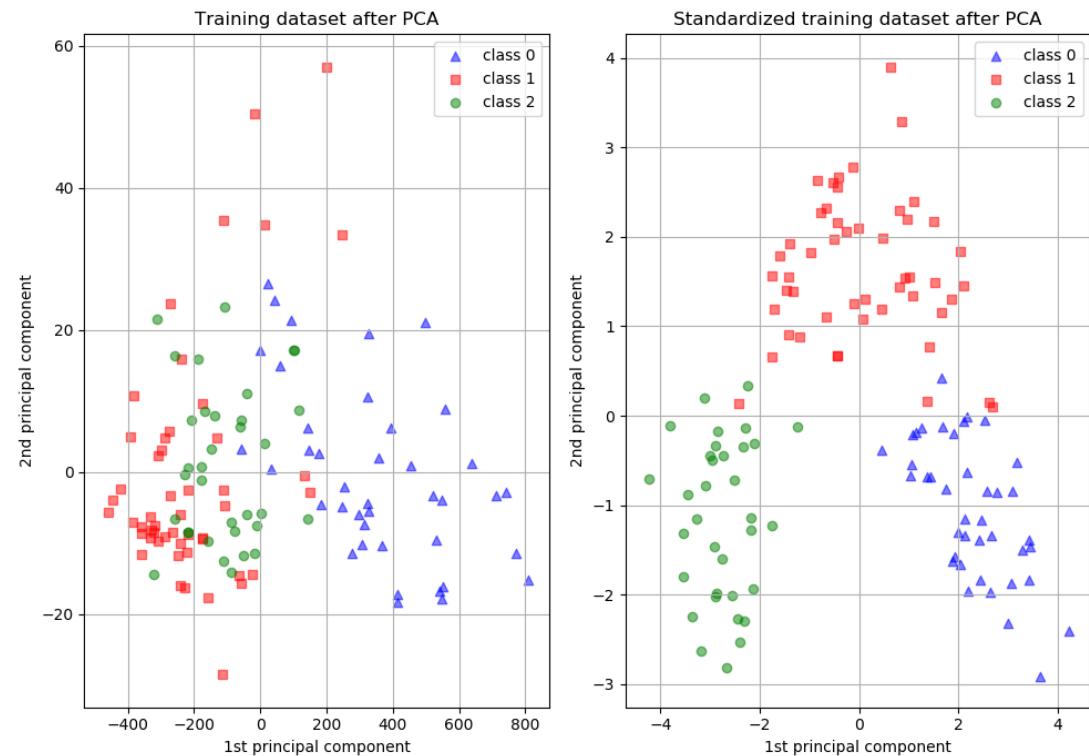
The function `scale` provides a quick and easy way to perform this operation on a single array-like dataset:

```
>>> from sklearn import preprocessing
>>> import numpy as np
```

Preprocessing ?

Standardization

❖ Standardization



Preprocessing for missing value

```
import pandas as pd  
  
data=pd.read_csv('M4_1.csv')  
data
```

	A	B	C	D
0	1.0	3	5.0	8
1	NaN	2	5.0	2
2	4.0	7	NaN	2
3	4.0	2	5.0	1

dropna()
)



```
data.dropna()
```

	A	B	C	D
0	1.0	3	5.0	8
3	4.0	2	5.0	1

dropna(axis=1)



```
data.dropna(axis=1)
```

	B	D
0	3	8
1	2	2
2	7	2
3	2	1

- ◆ Remember to assign dropped result to a new dataframe
- ✓ newdata=data.dropna()

Preprocessing for missing value

```
import pandas as pd  
  
data=pd.read_csv('M4_1.csv')  
data
```

	A	B	C	D
0	1.0	3	5.0	8
1	NaN	2	5.0	2
2	4.0	7	NaN	2
3	4.0	2	5.0	1

```
data.fillna(0)
```

fillna(value)

	A	B	C	D
0	1.0	3	5.0	8
1	0.0	2	5.0	2
2	4.0	7	0.0	2
3	4.0	2	5.0	1

Interpolate()

```
data.interpolate()
```

	A	B	C	D
0	1.0	3	5.0	8
1	2.5	2	5.0	2
2	4.0	7	5.0	2
3	4.0	2	5.0	1

Preprocessing for missing value

```
import pandas as pd  
  
data=pd.read_csv('M4_1.csv')  
data
```

	A	B	C	D
0	1.0	3	5.0	8
1	NaN	2	5.0	2
2	4.0	7	NaN	2
3	4.0	2	5.0	1

fillna(data.mean())



```
data.fillna(data.mean())
```

	A	B	C	D
0	1.0	3	5.0	8
1	3.0	2	5.0	2
2	4.0	7	5.0	2
3	4.0	2	5.0	1

Preprocessing in scikit-learn

Standardization

Standardize features by removing the mean and scaling to unit variance.

```
from sklearn import preprocessing
import numpy as np

xdata=[1, 2, 3, 4]
xscale=preprocessing.scale(xdata)
print(xscale, 'mean=', xscale.mean(), 'variance', xscale.var())
```

```
[-1.34164079 -0.4472136  0.4472136  1.34164079] mean= 0.0 variance 1.0
```

Preprocessing in scikit-learn

Standardization

The standardization process can be stored in a module which can be fit first, and then applied to other datasets later.

```
from sklearn import preprocessing

xdata=[1, 1, 0, 0]
scaler=preprocessing.StandardScaler()
scaler.fit(xdata)

standx=scaler.transform(xdata)
print(standx, 'mean=', standx.mean(), ', variance:', standx.var())

ydata=[1, 2, 3, 4]
ystand=scaler.transform(ydata)
print(ystand)

[ 1.  1. -1. -1.] mean= 0.0 , variance: 1.0
[1.  3.  5.  7.]
```



- ◆ We can use *scaler.fit_transform()* as well.
- ◆ The fit function is to find parameter for standardization
$$\frac{1}{\sigma} \varphi \left(\frac{x-u}{\sigma} \right)$$
- ◆ $\sigma = 0.5, u = 0.5$

Preprocessing in scikit-learn

Normalization

Scale input vectors individually to unit norm (vector length).

```
from sklearn import preprocessing

xdata=[1, 2, 3, 4]
print(preprocessing.normalize(xdata,norm='max'))
print(preprocessing.normalize(xdata,norm='l1'))
print(preprocessing.normalize(xdata,norm='l2'))
```

```
[[0.25 0.5 0.75 1. ]]
[[0.1 0.2 0.3 0.4]]
[[0.18257419 0.36514837 0.54772256 0.73029674]]
```

Preprocessing in scikit-learn

MinMaxScaler

Transforms features by scaling each feature to a given range.

```
from sklearn import preprocessing  
  
xdata=[1, 2, 3, 4]  
scale_x=preprocessing.minmax_scale(feature_range=(0,2),X=xdata)  
print(scale_x)
```

```
[0.          0.66666667 1.33333333 2.          ]
```

Preprocessing in scikit-learn

MinMaxScaler

We can also create an scaling object to do scaling operation.

```
from sklearn import preprocessing

xdata=[1, 2, 3, 4]
minmaxscaler=preprocessing.MinMaxScaler()
scale_x=minmaxscaler.fit_transform(xdata)
print(scale_x)
ydata=[-2, 0, 6, 10]
print(minmaxscaler.transform(ydata))
```

```
[0.          0.33333333 0.66666667 1.          ]
[-1.         -0.33333333 1.66666667 3.          ]
```

Preprocessing in scikit-learn

- ◆ Impute module in preprocessing package can also be used to fill the missing values.

```
import numpy as np
#from sklearn import SimpleImputer
from sklearn.preprocessing import Imputer

imp = Imputer(missing_values=np.nan, strategy='mean', axis=0)
xdata=[[1, 2], [np.nan, 1], [7, 6]]
imp.fit(xdata)
print(imp.transform(xdata))
X = [[np.nan, 2], [9, np.nan], [100, 6],[np.nan, np.nan]]
print(imp.transform(X))
```

Can also specify '**median**' or '**most_frequent**'

```
[[1. 2.]
 [4. 1.]
 [7. 6.]]
[[ 4.  2.]
 [ 9.  3.]
 [100.  6.]
 [ 4.  3.]]
```

1	2
Nan	1
7	6

Preprocessing in scikit-learn

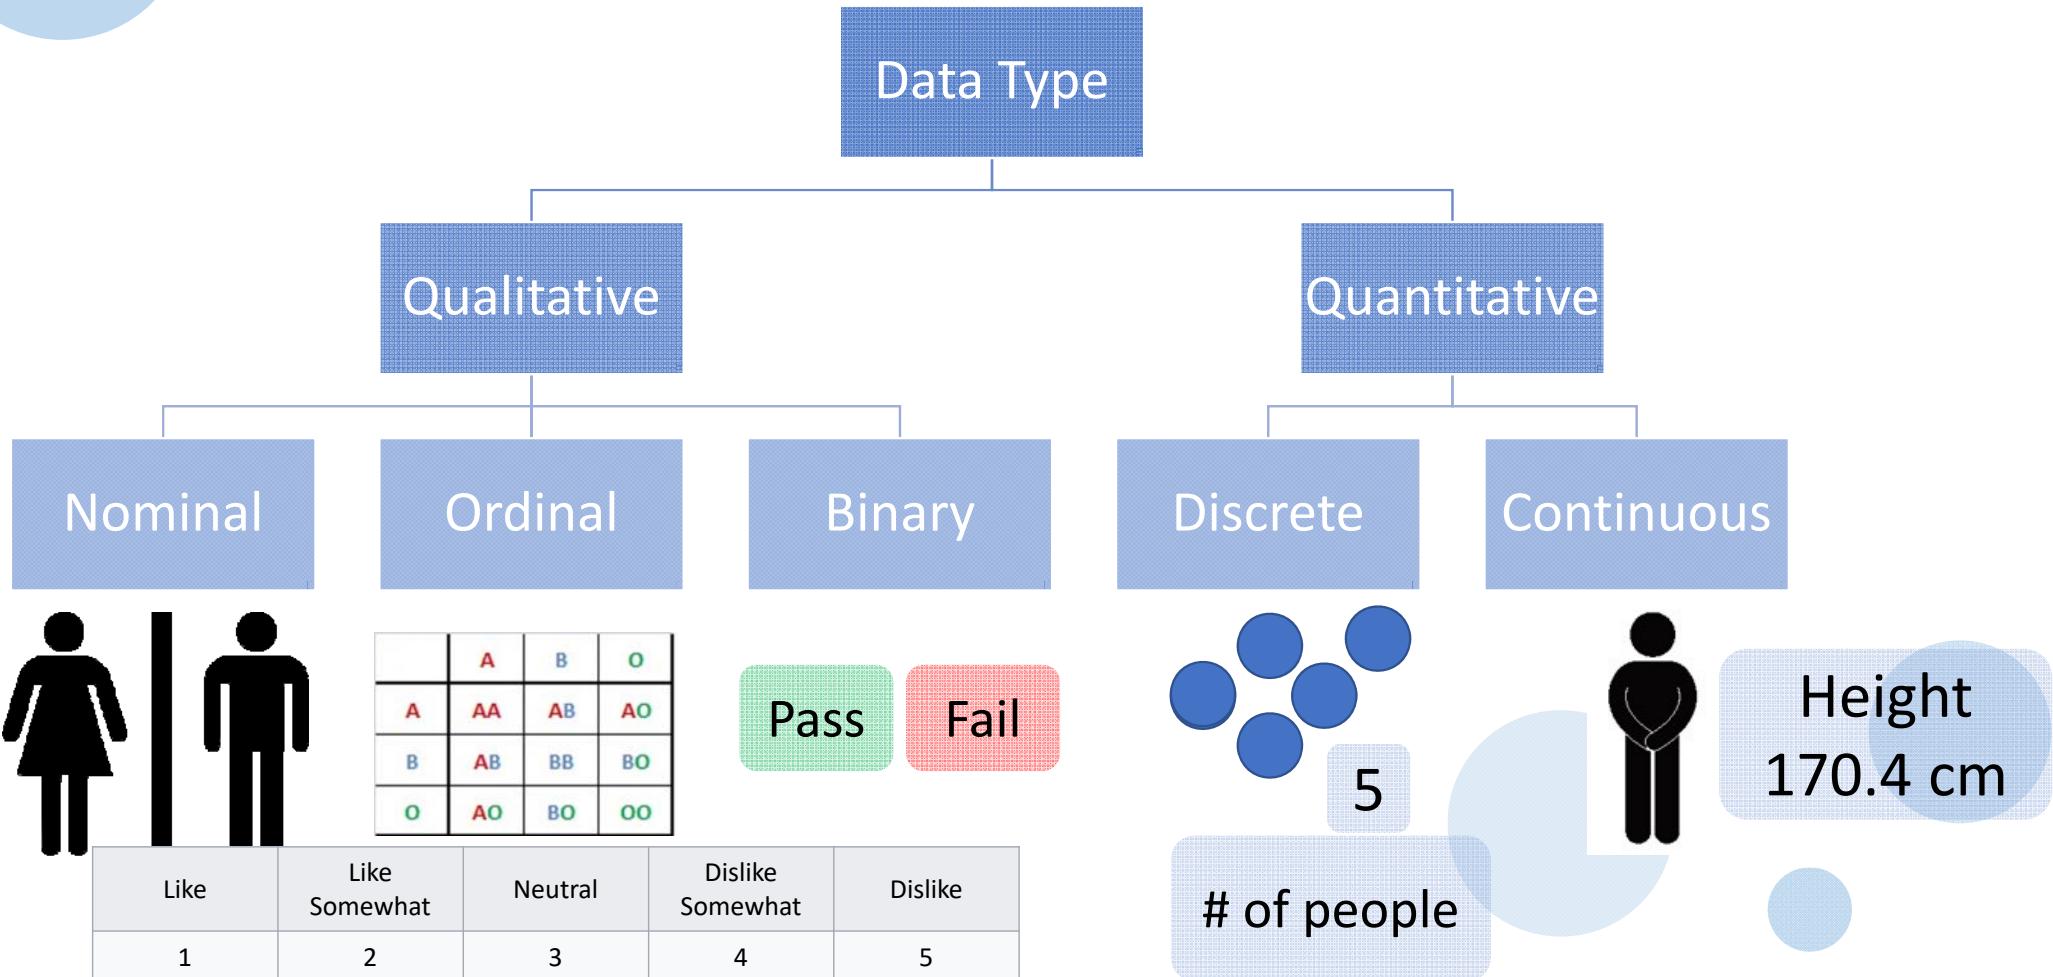
- ❖ Can also fill the missing values using the most frequent data.

```
import numpy as np
#from sklearn import SimpleImputer
from sklearn.preprocessing import Imputer

imp = Imputer(missing_values=np.nan, strategy='most_frequent', axis=0)
xdata=[[1, 2], [np.nan, 1], [7, 6],[7,2]]
imp.fit(xdata)
print(imp.transform(xdata))
X = [[np.nan, 2], [9, np.nan], [100, 6],[np.nan, np.nan]]
print(imp.transform(X))
```

```
[[1. 2.]
 [7. 1.]
 [7. 6.]
 [7. 2.]]
[[ 7.  2.]
 [ 9.  2.]
 [100.  6.]
 [ 7.  2.]]
```

Type of input variables/feature



Encoding of variables/features

Label encoding

Fruit	Label
Apple	1
Orange	2
Banana	3
Grape	4

One-hot encoding

Apple	Orange	Banana	Grape
1	0	0	0
0	1	0	0
0	0	1	0
0	0	0	1

Python package used for coding

```
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
le.fit(['Apple', 'Orange', 'Banana', 'Grape'])
    0      3      1      2
print(le.transform(['Orange', 'Grape', 'Banana', 'Grape', 'Apple']))
print(le.inverse_transform([0, 2, 1, 0]))
```

Import package

Encoder

Fit label

Encode data

```
[3 2 1 2 0]
['Apple' 'Grape' 'Banana' 'Apple']
```

One-hot encoding example

```
from sklearn import preprocessing
lo = preprocessing.OneHotEncoder()
le = preprocessing.LabelEncoder()

le_result=le.fit_transform(['Apple','Orange','Banana','Grape'])
lo_result=lo.fit_transform(le_result.reshape(-1,1)).toarray()

print(lo_result)
```

```
[[1. 0. 0. 0.]
 [0. 0. 0. 1.]
 [0. 1. 0. 0.]
 [0. 0. 1. 0.]]
```

One-hot coding using pandas

- ◆ Pandas provides *get_dummies()* function can also provide similar one-hot encoding result.

```
import pandas as pd

data = pd.DataFrame({"顏色": ["紅色", "紅色", "藍色", "黃色", "黃色"],
                     "Item": [10, 21, 30, 41, 56]})

data
```

	Item	顏色
0	10	紅色
1	21	紅色
2	30	藍色
3	41	黃色
4	56	黃色

```
pd.get_dummies(data['顏色'])
```

	紅色	藍色	黃色
0	1	0	0
1	1	0	0
2	0	1	0
3	0	0	1
4	0	0	1

```
np.array(pd.get_dummies(data['顏色']))

array([[1, 0, 0],
       [1, 0, 0],
       [0, 1, 0],
       [0, 0, 1],
       [0, 0, 1]], dtype=uint8)
```

```
pd.get_dummies(data)
```

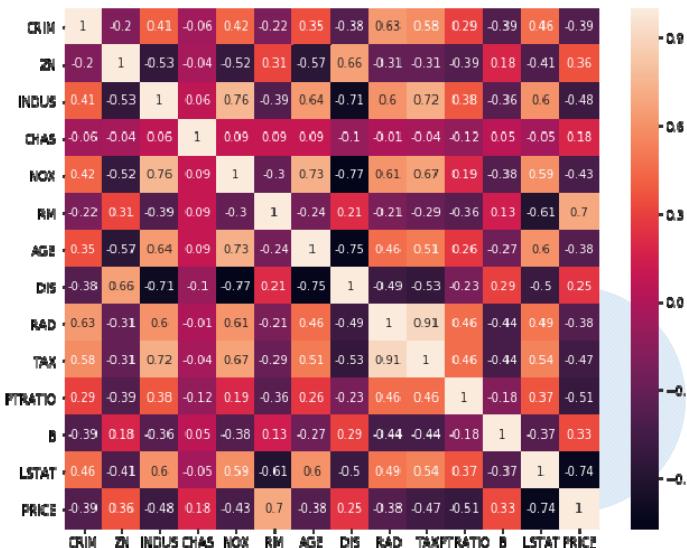
	Item	顏色_紅色	顏色_藍色	顏色_黃色
0	10	1	0	0
1	21	1	0	0
2	30	0	1	0
3	41	0	0	1
4	56	0	0	1

特徵選取

- ◆ 使用pandas套件找出相關係數，seaborn套件視覺化顯示相關係數的矩陣，範例如下：

```
import seaborn as sns  
correlation_matrix = df.corr()  
sns.heatmap(correlation_matrix, annot=True)
```

df.corr() : 產生相關係數的矩陣



相關係數

- ◆ 相關係數為兩個變數之間的**共變異數**和**標準差**的商
- ◆ 可以被轉換成下面幾種等價的形式

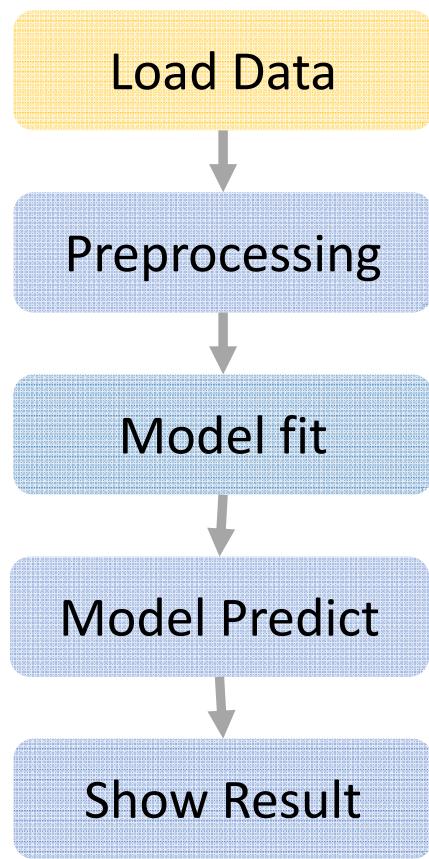
$$\rho_{X,Y} = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y} = \frac{E[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y}$$
$$r = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2} \sqrt{\sum_{i=1}^n (Y_i - \bar{Y})^2}}$$

$$r = \frac{1}{n-1} \sum_{i=1}^n \left(\frac{X_i - \bar{X}}{\sigma_X} \right) \left(\frac{Y_i - \bar{Y}}{\sigma_Y} \right)$$

Using Scikit-Learn for ML



Data Loading



❖ sklearn provides a dataset package which contains some practice data which can be loaded.

```
from sklearn import datasets  
diabetes = datasets.load_diabetes()
```

❖ Very often, the data are stored in csv files

```
import pandas as pd  
df = pd.read_csv('file.csv', header=None, sep=' ')
```

Sklearn dataset

<https://scikit-learn.org/stable/datasets/index.html>

6.2. Toy datasets

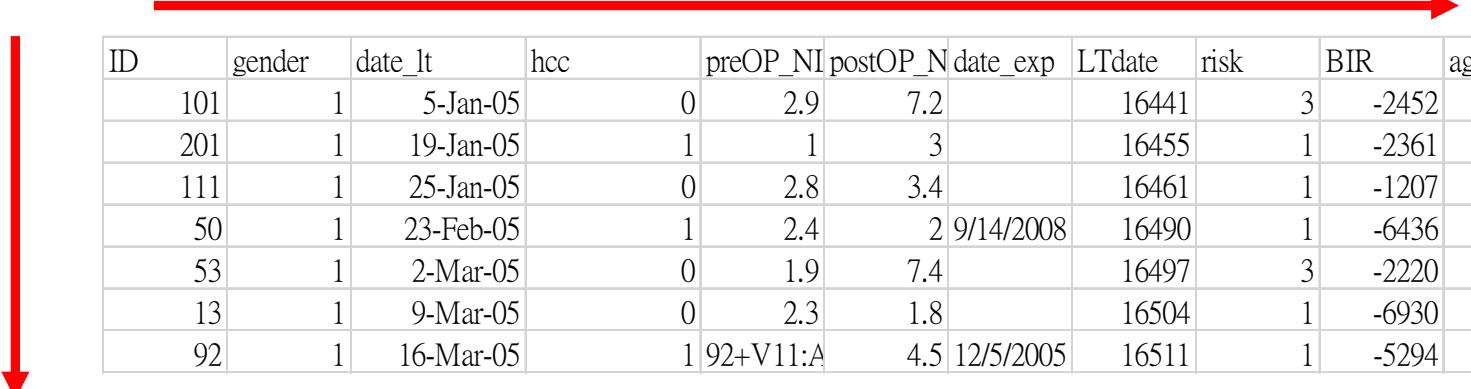
scikit-learn comes with a few small standard datasets that do not require to download any file from some external website.

They can be loaded using the following functions:

<code>load_boston</code> ([return_X_y])	Load and return the boston house-prices dataset (regression).
<code>load_iris</code> ([return_X_y])	Load and return the iris dataset (classification).
<code>load_diabetes</code> ([return_X_y])	Load and return the diabetes dataset (regression).
<code>load_digits</code> ([n_class, return_X_y])	Load and return the digits dataset (classification).
<code>load_linnerud</code> ([return_X_y])	Load and return the linnerud dataset (multivariate regression).
<code>load_wine</code> ([return_X_y])	Load and return the wine dataset (classification).
<code>load_breast_cancer</code> ([return_X_y])	Load and return the breast cancer wisconsin dataset (classification).

csv file of data

❖ You have organize your data in the following format

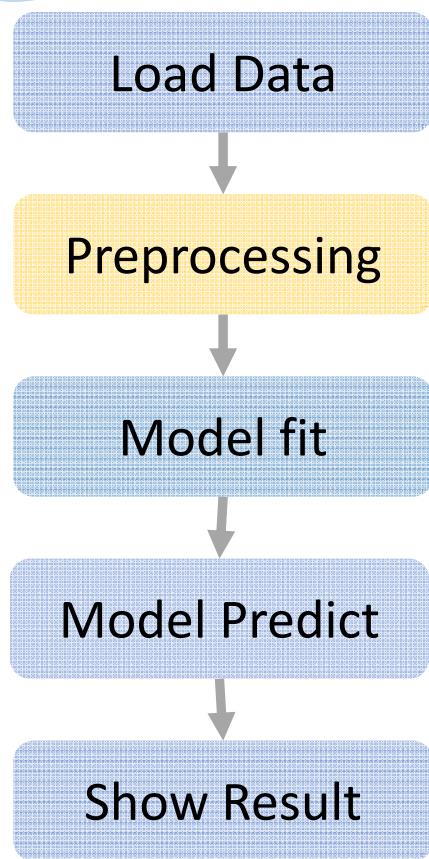


ID	gender	date_lt	hcc	preOP_NI	postOP_N	date_exp	LTdate	risk	BIR	age
101	1	5-Jan-05		0	2.9	7.2		16441	3	-2452
201	1	19-Jan-05		1	1	3		16455	1	-2361
111	1	25-Jan-05		0	2.8	3.4		16461	1	-1207
50	1	23-Feb-05		1	2.4	2	9/14/2008	16490	1	-6436
53	1	2-Mar-05		0	1.9	7.4		16497	3	-2220
13	1	9-Mar-05		0	2.3	1.8		16504	1	-6930
92	1	16-Mar-05		1	92+V11:A	4.5	12/5/2005	16511	1	-5294

cases

features

Data preparation/preprocessing



- ❖ Split dataset into training & testing sets
- ❖ Split each set into target and features

`x_train = x[:-20]`

`x_test = x[-20:]`

`y_train = diabetes.target[:-20]`

`y_test = diabetes.target[-20:]`

- ❖ Can use normalize or other methods

`from sklearn import preprocessing`

`x = preprocessing.scale(x)`

`y = preprocessing.scale(y)`

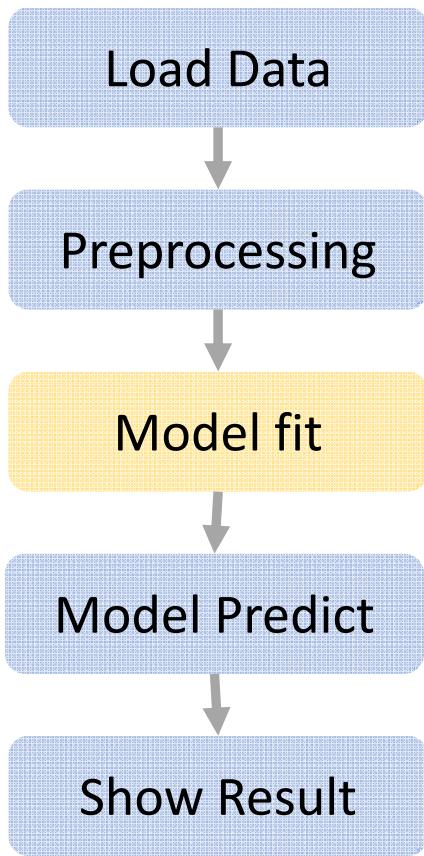
Some quick notes

- ❖ The input features **X** should be in the form of 2D array no matter the number of features being used.
 - ◆ Ex: `X=[[0.5],[1.2],[3.4].....]` instead of [0.5, 1.2, 3.4]
- ❖ The function ***train_test_split()*** can be used to split the data into the training and testing datasets.

```
❖ import numpy as np
from sklearn.model_selection import train_test_split
X, y = np.arange(10).reshape((5, 2)), range(5)
print('Input X:', X)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.33, random_state=42)
print('Training X:', X_train)
print('Training Y:', y_train)
print('Test X:', X_test)
print('Test Y:', y_test)
```

Input X: [[0 1]
[2 3]
[4 5]
[6 7]
[8 9]]
Training X: [[4 5]
[0 1]
[6 7]]
Training Y: [2, 0, 3]
Test X: [[2 3]
[8 9]]
Test Y: [1, 4]

Model object creation and fit



- ❖ Create a model object using the module provided by sklearn packages

```
from sklearn import linear_model  
regr = linear_model.LinearRegression()
```

- ❖ We can set the relevant hyper-parameters

- ❖ Feed the data into the model object to fit/train.

```
regr.fit(x_train, y_train)
```

Linear model provided by sklearn

sklearn.linear_model: Generalized Linear Models

The `sklearn.linear_model` module implements generalized linear models. It includes Ridge regression, Bayesian Regression, Lasso and Elastic Net estimators computed with Least Angle Regression and coordinate descent. It also implements Stochastic Gradient Descent related algorithms.

User guide: See the [Generalized Linear Models](#) section for further details.

<code>linear_model.ARDRegression ([n_iter, tol, ...])</code>	Bayesian ARD regression.
<code>linear_model.BayesianRidge ([n_iter, tol, ...])</code>	Bayesian ridge regression.
<code>linear_model.ElasticNet ([alpha, l1_ratio, ...])</code>	Linear regression with combined L1 and L2 priors as regularizer.
<code>linear_model.ElasticNetCV ([l1_ratio, eps, ...])</code>	Elastic Net model with iterative fitting along a regularization path.
<code>linear_model.HuberRegressor ([epsilon, ...])</code>	Linear regression model that is robust to outliers.
<code>linear_model.Lars ([fit_intercept, verbose, ...])</code>	Least Angle Regression model a.k.a.
<code>linear_model.LarsCV ([fit_intercept, ...])</code>	Cross-validated Least Angle Regression model.
<code>linear_model.Lasso ([alpha, fit_intercept, ...])</code>	Linear Model trained with L1 prior as regularizer (aka the Lasso)
<code>linear_model.LassoCV ([eps, n_alphas, ...])</code>	Lasso linear model with iterative fitting along a regularization path.
<code>linear_model.LassoLars ([alpha, ...])</code>	Lasso model fit with Least Angle Regression a.k.a.
<code>linear_model.LassoLarsCV ([fit_intercept, ...])</code>	Cross-validated Lasso, using the LARS algorithm.
<code>linear_model.LassoLarsIC ([criterion, ...])</code>	Lasso model fit with Lars using BIC or AIC for model selection
<code>linear_model.LinearRegression ([...])</code>	Ordinary least squares Linear Regression.
<code>linear_model.LogisticRegression ([penalty, ...])</code>	Logistic Regression (aka logit, MaxEnt) classifier.
<code>linear_model.LogisticRegressionCV ([Cs, ...])</code>	Logistic Regression CV (aka logit, MaxEnt) classifier.
<code>linear_model.MultiTaskLasso ([alpha, ...])</code>	Multi-task Lasso model trained with L1/L2 mixed-norm as regularizer.
<code>linear_model.MultiTaskElasticNet ([alpha, ...])</code>	Multi-task ElasticNet model trained with L1/L2 mixed-norm as regularizer
<code>linear_model.MultiTaskLassoCV ([eps, ...])</code>	Multi-task Lasso model trained with L1/L2 mixed-norm as regularizer.
<code>linear_model.MultiTaskElasticNetCV ([...])</code>	Multi-task L1/L2 ElasticNet with built-in cross-validation.
<code>linear_model.OrthogonalMatchingPursuit ([...])</code>	Orthogonal Matching Pursuit model (OMP)
<code>linear_model.OrthogonalMatchingPursuitCV ([...])</code>	Cross-validated Orthogonal Matching Pursuit model (OMP).
<code>linear_model.PassiveAggressiveClassifier ([...])</code>	Passive Aggressive Classifier
<code>linear_model.PassiveAggressiveRegressor ([C, ...])</code>	Passive Aggressive Regressor
<code>linear_model.Perceptron ([penalty, alpha, ...])</code>	Read more in the User Guide .
<code>linear_model.RANSACRegressor ([...])</code>	RANSAC (RANdom SAmple Consensus) algorithm.
<code>linear_model.Ridge ([alpha, fit_intercept, ...])</code>	Linear least squares with L2 regularization.
<code>linear_model.RidgeClassifier ([alpha, ...])</code>	Classifier using Ridge regression.

[Previous](#) [Next](#) [Up](#) [API Reference](#)

scikit-learn v0.21.2
Other versions

Please cite us if you use
the software.

`sklearn.linear_model.LinearRe
gression`
Examples using
`sklearn.linear_model.LinearRe`

sklearn.linear_model.LinearRegression

`class sklearn.linear_model. LinearRegression (fit_intercept=True, normalize=False, copy_X=True, n_jobs=None)` [\[source\]](#)

Ordinary least squares Linear Regression.

Parameters: `fit_intercept : boolean, optional, default True`

whether to calculate the intercept for this model. If set to False, no intercept will be used in calculations (e.g. data is expected to be already centered).

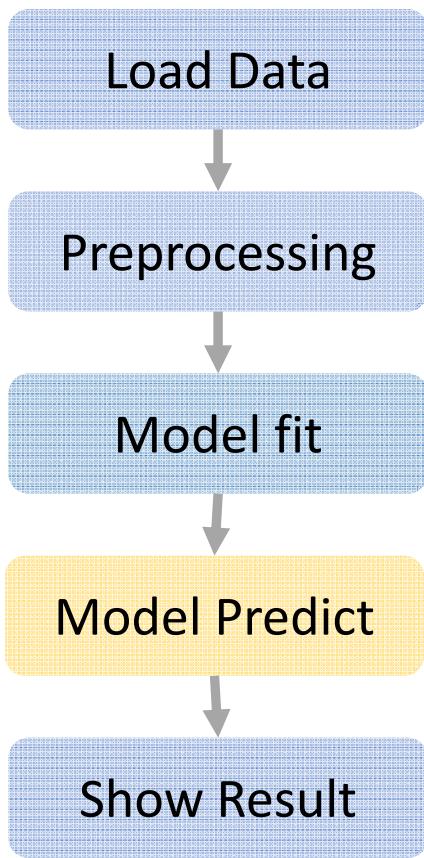
`normalize : boolean, optional, default False`

This parameter is ignored when `fit_intercept` is set to False. If True, the regressors X will be normalized before regression by subtracting the mean and dividing by the l2-norm. If you wish to standardize, please use `sklearn.preprocessing.StandardScaler` before calling `fit` on an estimator with `normalize=False`.

Methods

<code>fit (self, X, y[, sample_weight])</code>	Fit linear model.
<code>get_params (self[, deep])</code>	Get parameters for this estimator.
<code>predict (self, X)</code>	Predict using the linear model
<code>score (self, X, y[, sample_weight])</code>	Returns the coefficient of determination R^2 of the prediction.
<code>set_params (self, **params)</code>	Set the parameters of this estimator.

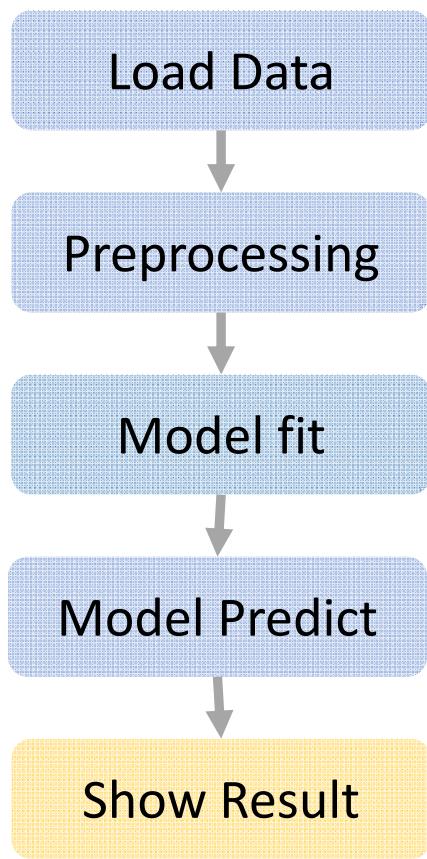
Predict test data using the trained model



- ❖ Call the *predict()* method of the trained model object.

y_pred = regr.predict(x_test)

Show the predicted result



- ❖ Can use graphics to show results

```
import matplotlib.pyplot as plt  
plt.scatter(x_test, y_test, color='blue')  
plt.plot(x_test, y_pred, color='red')  
plt.show()
```

- ❖ Can also use some metrics

- ❖ Can call `score()` method of the model object
 - ❖ Or use metric module

Metric/score for evaluation

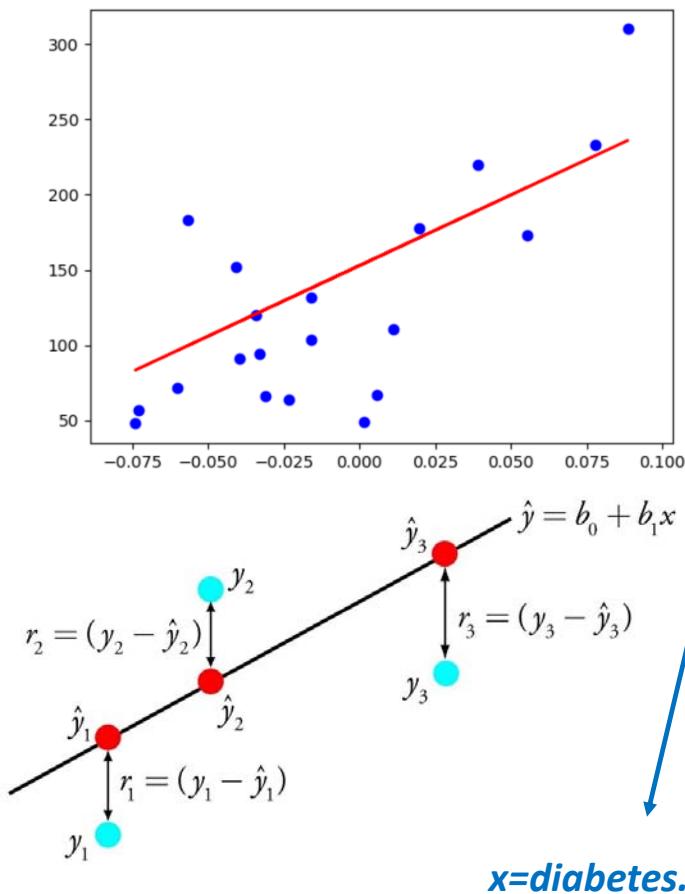
◊ *From sklearn import metrics*

Regression metrics

See the [Regression metrics](#) section of the user guide for further details.

<code>metrics.explained_variance_score</code> (<code>y_true</code> , <code>y_pred</code>)	Explained variance regression score function
<code>metrics.max_error</code> (<code>y_true</code> , <code>y_pred</code>)	max_error metric calculates the maximum residual error.
<code>metrics.mean_absolute_error</code> (<code>y_true</code> , <code>y_pred</code>)	Mean absolute error regression loss
<code>metrics.mean_squared_error</code> (<code>y_true</code> , <code>y_pred</code> [, ...])	Mean squared error regression loss
<code>metrics.mean_squared_log_error</code> (<code>y_true</code> , <code>y_pred</code>)	Mean squared logarithmic error regression loss
<code>metrics.median_absolute_error</code> (<code>y_true</code> , <code>y_pred</code>)	Median absolute error regression loss
<code>metrics.r2_score</code> (<code>y_true</code> , <code>y_pred</code> [, ...])	R^2 (coefficient of determination) regression score function.

Overall Linear Regression Example



```
from sklearn import linear_model, datasets  
import numpy as np  
import matplotlib.pyplot as plt  
  
diabetes = datasets.load_diabetes()  
  
x = diabetes.data[:, np.newaxis, 2]  
  
x_train = x[:-20]  
x_test = x[-20:]  
  
y_train = diabetes.target[:-20]  
y_test = diabetes.target[-20:]  
  
regr = linear_model.LinearRegression()  
regr.fit(x_train, y_train)  
y_pred = regr.predict(x_test)  
  
plt.scatter(x_test, y_test, color='blue')  
plt.plot(x_test, y_pred, color='red')  
  
plt.show()
```

Import package

Load dataset

Split dataset

Select model

Fit model

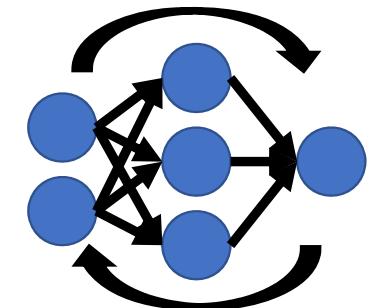
inference

Display result

Model Selection

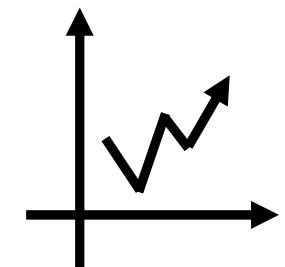
Training set

Training model



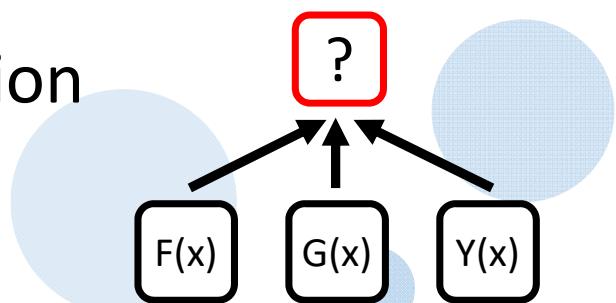
Testing set

Assess the performance



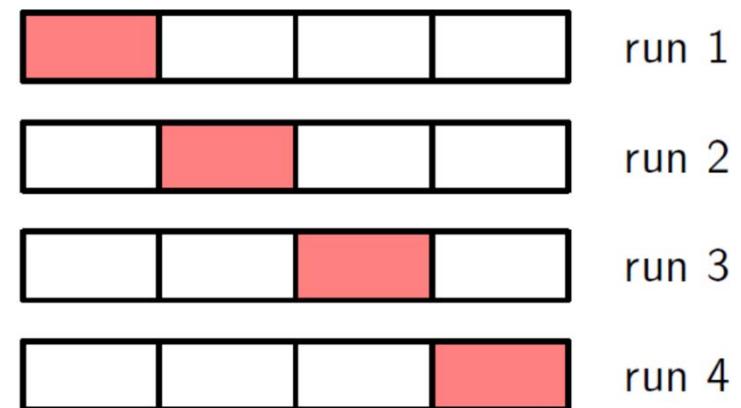
Validation set

Use for model selection

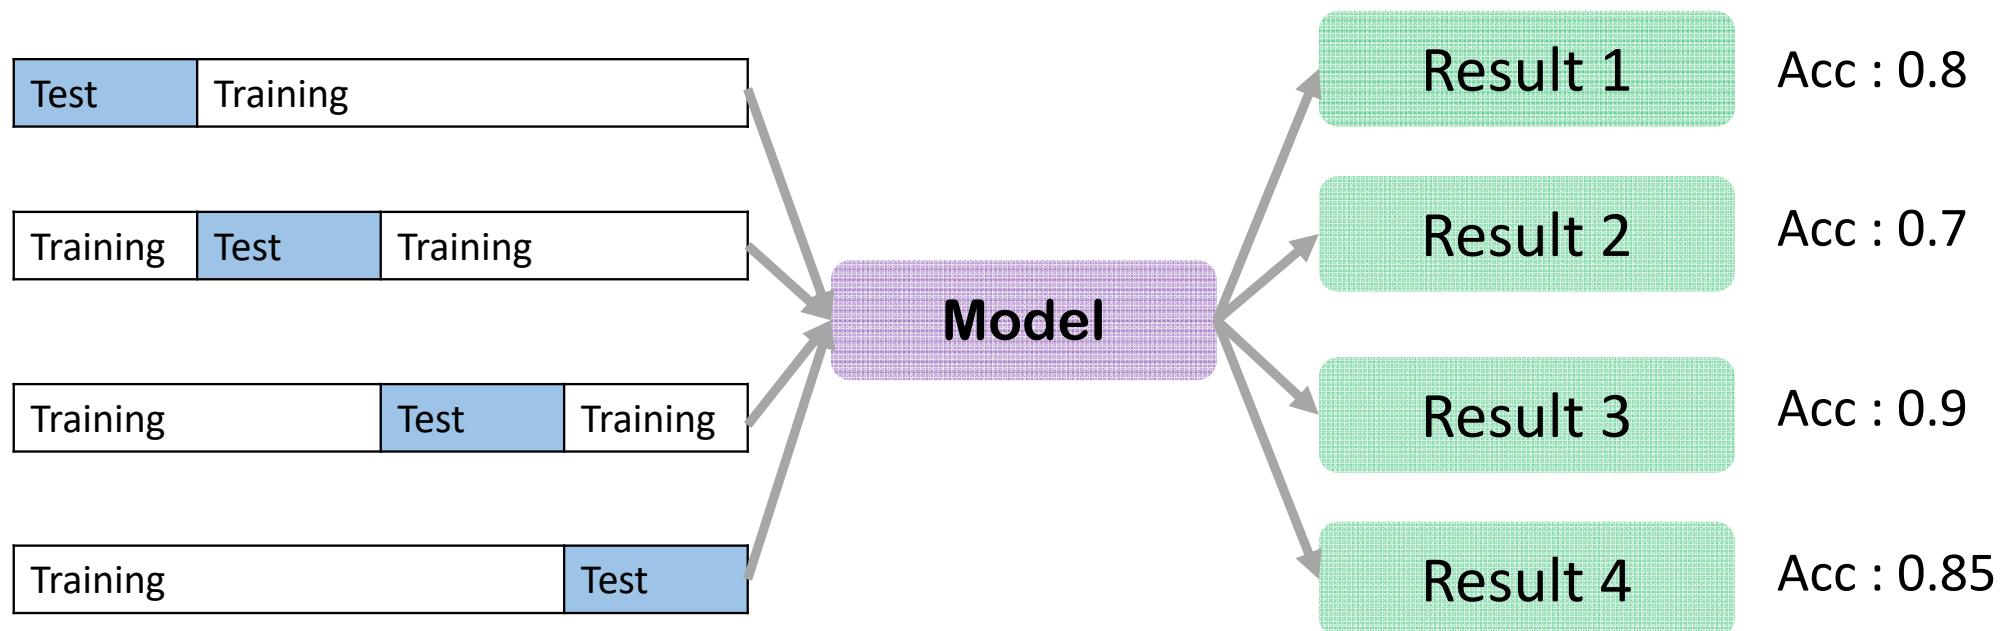


Model Selection

- ❖ **Cross-validation** method can be applied for limited data, which allows a proportion $(S - 1)/S$ of the available data to be used for training.



Cross-validation



$$\text{Mean Acc} : (0.8 + 0.7 + 0.9 + 0.85)/4 = 0.8125$$

Cross-validation

```
# from sklearn import datasets, Linear_model
from sklearn.model_selection import cross_val_score
diabetes = datasets.load_diabetes()
X = diabetes.data[:150]
y = diabetes.target[:150]
lreg = linear_model.LinearRegression()
print(cross_val_score(lreg, X, y, cv=5))
```

```
[0.36324841 0.28239194 0.4211776  0.30071196 0.61240533]
```